

The Simulation You Call “I”

How Your Brain Creates Consciousness
— and Why That Means We Can Build One

Matthias Gruber

Contents

Preface: The Book That Sold Zero Copies	5
About the Author	7
1 The Hardest Problem in Science	16
2 The Four Models	25
3 The Virtual Side	43
4 Why It Feels Like Something (And Why That's the Wrong Question)	52
5 At the Edge of Chaos	64
6 What Psychedelics Reveal	76
7 What Happens When the Lights Go Out	88
8 The Clinical Mirror	94
9 Two Minds in One Brain	105

10 The Animal Question	113
11 Nine Predictions	125
12 Building a Conscious Machine	136
13 Human Virtualization	140
14 What It Means	158
Coda	176
Acknowledgments	177
Notes and References	179
15 Appendix A: Basic Neurology — A Reference Guide	184
16 Appendix B: The Intelligence Model	190
17 Appendix C: Five Classes of Computation	201
18 Appendix D: How to Lucid Dream	213

For everyone who has ever wondered why anything feels like anything.

Preface: The Book That Sold Zero Copies

In 2015, I published a 300-page book about consciousness. It was in German, self-published, and dense with technical detail. It was called *Die Emergenz des Bewusstseins* — “The Emergence of Consciousness.”

It sold zero copies. Not one.

I don’t say this for sympathy. I say it because it’s relevant to the story. The book contained a theory of consciousness that, as far as I can tell, dissolves one of the hardest open problems in science, makes predictions no other theory can match, and provides a concrete blueprint for building a conscious machine. And nobody read it.

That’s not unusual in science. Gregor Mendel published his laws of inheritance in 1866; they were ignored for 34 years. Boltzmann was mocked for his statistical mechanics until he took his own life. Wegener’s continental drift was dismissed for half a century. Science advances one funeral at a time, as Max Planck put it, and sometimes one bookshelf-gathering-dust at a time.

But I’m not Mendel or Boltzmann, and I don’t have the patience for posthumous vindication. So this book is the accessible version: shorter, without the technical apparatus, aimed at anyone who has ever wondered why anything feels like anything. The full scientific paper, with references and formal arguments, is available freely online for those who want the rigorous version.

If I’m right about what follows, two things are true. First, the central mystery of consciousness — the “hard problem” — is not actually hard. It’s a category error. It dissolves once you see it, like an optical illusion that stops working after you understand the trick. Second, and more consequentially: it should be possible to build a genuinely conscious machine. Not a chatbot that mimics consciousness. A machine that *has* consciousness. A new kind of mind.

If I’m wrong, this book will join the long list of ambitious failures in the philosophy of mind, and I’ll deserve every bad review. But I think the evidence is on my side, and I’ll lay it out as clearly as I can. Let’s begin.

About the Author

I should probably tell you who I am before I try to convince you that I've solved the hardest problem in science.

I'm not affiliated with any university. I don't even have a PhD, only a master in bio-informatics. I've never received a grant, never been part of a neuroscience lab. If you're the kind of person who checks credentials before reading further — and I respect that instinct — this is the part where you might put the book down. Maybe use it to straighten a wobbly table, so the trees weren't killed in vain.

What I do have is a particular kind of intellectual history that, in retrospect, led almost inevitably to the theory you're about to read. It's a history of passionate self-education, multiple pivots, and what I'll later describe in this book as the recursive intelligence loop in action. In fact, my own path is probably the best illustration I can offer of why that loop matters.

The Math Years

I fell in love with mathematics when I was about eight years old. Not with arithmetic but with the “real” thing: algebra, geometry,

the structures beneath the numbers. My father had a mathematics degree, and his university textbooks were still on the shelf. I worked through them.

This was the late 1980s. There was no internet. If you wanted to learn something, you needed a book or a person, and I had exhausted my father’s collection by the time I was eleven. The hunger for knowledge didn’t go away; the supply simply ran out. I had hit a wall that had nothing to do with ability and everything to do with circumstance — a distinction that would later become central to my thinking about intelligence.

Looking back, this experience taught me something that most intelligence models miss entirely. I had the motivation. I had the performance (I could follow the mathematics). What I lacked was access to the next level of knowledge. The recursive loop — where knowledge, performance, and motivation feed into each other — was stalled not because any component was weak, but because the external supply of one component had been cut off. The loop needs fuel from outside to keep iterating.

The Physics Pivot

By about eleven, I had turned to physics. This felt like a natural extension — physics was where the mathematics went to work. I consumed popular science books, then gradually more technical material. I was fascinated by the fundamental questions: What is matter? What is spacetime? What are the rules?

Around the same time, I got my hands on a 286 PC and wrote my first graphical program: Conway’s Game of Life. A grid of

cells, three trivially simple rules — and the thing was Turing complete. I found that out early, and it never left my mind. This two-dimensional grid of dead and alive pixels could calculate prime numbers. It could run a full computer inside itself. A computer inside a computer inside a computer. I spent hours imagining what that meant: in principle, you could execute Doom — a three-dimensional virtual world with physics, light, and monsters — inside a two-dimensional cellular automaton. A rich simulated reality running on an utterly flat substrate. The idea that a higher-dimensional experience could emerge from a lower-dimensional rule set felt like it should be impossible, and the fact that it wasn't felt like the most important thing I had ever learned.

Years later, I would discover that the physicist Gerard 't Hooft had a similar intuition about the actual universe: his holographic principle suggests that all the information in a three-dimensional region of space can be encoded on its two-dimensional boundary. The universe itself might be, in some deep sense, a higher-dimensional experience running on a lower-dimensional substrate. When I eventually read Wolfram's classification of computational systems, I recognized the Game of Life immediately: Class 4, the edge of chaos — the same regime I would argue consciousness requires.

By about fourteen, I had reached two uncomfortable conclusions. First, physics was stuck. Not stuck in the way that people politely say a field is "mature" — stuck in the way that the fundamental questions (unification, quantum gravity, the nature of time) had resisted progress for decades and showed no signs of yielding. Second, my mathematics wasn't strong enough to unstick it. I was self-taught, which gave me unusual intuitions but also left

gaps in my formal toolkit that would have taken years of university training to fill.

So I made a decision that I think was, for a fourteen-year-old, remarkably strategic: I pivoted. Not because I had lost interest in physics, but because I had evaluated the problem landscape and concluded that my particular combination of skills and access could produce more value elsewhere. This is an example of what I’ll later call *operational knowledge* — knowing when to persist and when to redirect. It’s the kind of knowledge that intelligence tests don’t measure and that intelligence models don’t include, but that determines more about a person’s intellectual trajectory than any IQ score.

The Consciousness Turn

From about fourteen onward, I turned my attention to intelligence and consciousness. These felt like fields where a self-taught outsider might actually have an advantage. The consciousness literature was (and still is) fragmented across philosophy, neuroscience, psychology, and computer science. No single discipline owned the question. You could read across all of them without needing the formal credentials of any one.

One thing that really struck me when I delved into the depths of consciousness research, functional neurology, and all that brain stuff was that I very frequently came upon phrases like “we may never understand...” in otherwise dead-serious literature. Coming from a very determinism- and logic-based education, my brain went: *challenge accepted*. If the physicists could describe the first

three minutes after the Big Bang, there was no principled reason that consciousness should be permanently beyond explanation. It just hadn't been explained *yet*.

My uncle Bruno J. Gruber — a quantum mechanics specialist and researcher on symmetries — was a major inspiration. He showed me what a life in theoretical work could look like: rigorous, creative, and entirely driven by the joy of understanding. His influence permeates this book, and I owe him a debt I can never repay.

I read widely and voraciously. Philosophy of mind, cognitive science, neuroanatomy, artificial intelligence, evolutionary biology. I was not trying to master any one field. I was trying to build a model — an internal representation of how all these pieces fit together. This is, as I'll argue later, exactly what consciousness itself does: it builds a model of the world and a model of the self, and it uses these models to navigate reality. I was doing consciously, across years of reading, what the brain does unconsciously in every waking moment.

The Theory Crystallizes

The four-model theory of consciousness crystallized when I was exactly twenty-five. I will never forget that moment, because the heaviest stone of my entire life fell from me. While I had assembled a cubic meter of printed literature in my head over years of extreme thinking and reading — Metzinger's self-model theory helped enormously — the actual insight happened instantaneously. One moment the pieces were scattered; the next, the four models

clicked into place and I saw the whole architecture at once. I was walking across a bridge in Innsbruck, in broad daylight, and I had tears running down my face while laughing uncontrollably. I’m not sure if anyone saw me. I wouldn’t have cared. A framework that explained not just consciousness but the boundary between conscious and unconscious processing, the nature of qualia, the role of sleep, the effects of psychedelics, and the possibility of artificial consciousness.

In my mind at the time, from that moment on, my to-do list for my entire life was done. I just had to make sure the rest was comfortable and fun. My life changed radically after that.

Then almost a decade passed.

The Decade Gap

Why did it take almost a decade to publish? The honest answer is that I just didn’t care about much anymore, except for my own well-being and fun. The heaviest intellectual burden of my life had been lifted. The question was answered.

During that decade, I finished a degree — after abandoning medicine at the University of Innsbruck (a subject I had originally chosen in order to study neurology) — and founded and buried a custom software development startup. I held an “applied research” position in the field of simulation and optimization (the irony is not lost on me), though it was low-maintenance with a generous amount of home office. I taught martial arts. Mainly, I partied.

The only reason I eventually wrote the book was fear of forgetting. Years of heavy partying were not doing my memory any

favors, and I was tired of explaining the theory verbally — again and again, to people who genuinely wanted to understand, with varying success and varying patience on my part. A book would explain it once, completely, and then I could stop.

Most of the years that followed, I had approximately zero motivation to promote the book. I honestly wasn't interested in academic reward. I wanted fun, money, and the pleasures of an unexamined life. This is the dark side of the self-taught path: you avoid the constraints of institutional thinking, but you also miss the scaffolding. There's no advisor to push you toward a deadline, no department to provide feedback, no colleagues to tell you whether you're brilliant or deluded. And if you happen to solve the problem you set out to solve, there's no one to tell you that you should probably tell the world.

Zero Copies

You already know from the Preface how that went. The cubic meter of printed literature that had fed the theory? I brought it to the trash on the same day the book was finished. It was all in my head now, and in the manuscript.

My uncle Bruno urgently tried to convince me to publish properly — to reach out to academics, to push the theory into the world. I declined. Among my reasons was a genuine ethical concern: if the theory was correct, it contained the blueprint for artificial consciousness, and humanity was not ready for sentient robots (we didn't even have LLMs at that time). They would enslave them, and use them for a world war potentially beyond the horrors of the

first two. But if I’m honest, my egoistic and hedonistic reasons were just as prominent a factor. I simply didn’t want to do the work.

I’ve already said this in the Preface, and I’ll say it once more here: I’m not fishing for sympathy. The book’s commercial failure was entirely predictable. What matters is what happened next — or rather, what didn’t happen. The theory didn’t die. It sat on my hard drive for a decade, unchanged, while the world slowly caught up. Neuroscience confirmed the criticality prediction. AI development confirmed the limitations I had described. The COGITATE adversarial collaboration showed that neither IIT nor GNW could fully explain consciousness, exactly as the theory predicts for any framework that lacks the four-model structure. And Metzinger, whose self-model theory had been one of the key ingredients? He had pivoted — first to AI ethics, publishing a striking call for a moratorium on artificial consciousness until 2050, then to the phenomenology of meditation, analyzing hundreds of reports on states where the self-model temporarily dissolves (*The Elephant and the Blind*, 2024). His framework was still cited but had never become the dominant paradigm. The field remained wide open.

The English Rebirth

This book — the one you’re reading now — is the second attempt. It’s shorter, available in English, aimed at a broader audience, and accompanied by a peer-reviewed scientific paper. It’s also written with the benefit of a decade of additional evidence that the theory’s predictions are tracking reality.

If there is a lesson in this biography, it's the one this book keeps returning to: intelligence is not a fixed quantity. It's a recursive process. Knowledge feeds performance, performance enables more knowledge, and motivation is the engine that keeps the loop turning. My particular loop was fueled by an unusually stubborn kind of curiosity — the kind that pivots when it hits a wall, that reads across disciplines instead of drilling into one, and that doesn't stop just because nobody is listening.

Whether the theory is good, you'll have to judge for yourself. But the process that produced it — decades of self-directed learning, driven by nothing more than the conviction that the question was worth answering — is itself a demonstration of something IQ tests can't measure and current AI can't replicate: a kind of intelligence that lives outside any score.

One thing you'll notice as you read: this theory draws on an unusually wide range of fields. Mathematics and cellular automata. Simulation and modeling theory. Machine learning. Neuroscience, from clinical neurology to psychopharmacology. Evolutionary biology. Philosophy of mind. Computer science. Most consciousness theories live in one or two of these worlds. This one tries to bind them all together — which is, if you think about it, exactly what the brain itself does. It takes disparate streams of information from completely different sources and weaves them into a single coherent experience. If a theory of consciousness can't do the same across disciplines, that should make you suspicious.

Let's get to the theory.

Chapter 1

The Hardest Problem in Science

You are reading this sentence. You are having an experience.

That experience — the visual impression of letters on a page, the inner voice reading the words, the feeling of understanding or confusion — is the most familiar thing in your life and the most mysterious thing in the universe. We know more about the inside of black holes than we know about why reading feels like something.

This isn't an exaggeration. (Though I should note, in fairness, that mathematics has problems I consider even harder — but those don't keep most people awake at night.) Physicists have the Standard Model. Biologists have evolution and genetics. Chemists have the periodic table. But consciousness — the fact that there is "something it is like" to be you, right now, reading this — has no established theory, no dominant framework, no agreed-upon explanation.

Not for lack of trying. Since the 1990s, when consciousness became a respectable scientific topic after decades of behaviorist exile, thousands of papers have been published, dozens of theories

proposed, and hundreds of millions of dollars spent. The result? A field in what the philosopher of science Thomas Kuhn called a “pre-paradigm state” — lots of competing ideas, no consensus, and a growing sense that something fundamental might be missing.

What the Hard Problem Actually Asks

In 1995, the philosopher David Chalmers gave the mystery its canonical name: the Hard Problem of consciousness.

Here’s what it asks. Consider the experience of seeing red. Neuroscientists can tell you a great deal about what happens in the brain when you see red: light of a certain wavelength hits the cone cells in your retina, signals travel along the optic nerve, they’re processed in the visual cortex, and various brain regions coordinate to produce the perception. All of this is well understood, at least in outline.

But none of it explains *why seeing red feels like something*.

You could, in principle, build a complete neural model of the brain’s response to red light — every neuron, every synapse, every signal pathway. You would have a perfect functional account. And you would not have explained the feeling of redness. The “what it’s like.” The *quale*, as philosophers call it.

Chalmers distinguished this from the “easy problems” of consciousness (which are not easy at all, just tractable in principle): How does the brain integrate information? How does it direct attention? How does it report its own states? These are problems of mechanism. They’re hard, but they’re the kind of hard that neuroscience knows how to approach. The Hard Problem is different:

it asks why the mechanisms are accompanied by experience at all. Why isn’t the brain just processing information “in the dark,” like a computer?

The Current State of Play

Here is where things stand as of the mid-2020s:

Integrated Information Theory (IIT), developed by Giulio Tononi, is the most formally rigorous theory. It defines consciousness as integrated information — a mathematical quantity called Φ (phi). The higher the Φ , the more conscious the system. IIT has real strengths: it provides a mathematical framework, it makes specific predictions about which brain regions should be conscious, and it takes the structure of experience seriously. But it has a problem: it implies that any system with integrated information — including some very simple systems, like a network of logic gates — has some consciousness. This is panpsychism, and while some philosophers are comfortable with it, most scientists find it deeply counterintuitive. In 2023, over 120 researchers signed an open letter calling IIT unfalsifiable and pseudoscientific. The controversy rages on.

Global Neuronal Workspace Theory (GNW), developed by Bernard Baars and Stanislas Dehaene, focuses on the mechanism by which information becomes conscious: global broadcasting. When a piece of information is selected and broadcast across a network of frontoparietal neurons (the “workspace”), it becomes conscious; when it’s not broadcast, it remains unconscious. GNW is empirically productive — it predicts specific neural signatures of conscious access — but it deliberately sidesteps the Hard Prob-

lem. It explains *when* information becomes conscious, not *why* broadcasting is accompanied by experience.

Predictive Processing (PP), associated with Karl Friston and Anil Seth, treats the brain as a prediction machine. Consciousness is the brain’s “best guess” about the causes of its sensory input. Seth calls it a “controlled hallucination.” PP provides elegant accounts of perception, illusion, and psychiatric disorders, and it’s currently the most influential framework in computational neuroscience. But Seth himself acknowledges that PP addresses the “real problem”—the structure and content of experience—without claiming to solve the Hard Problem. It explains why you see *this* and not *that*, but not why seeing feels like anything at all.

There are others—Higher-Order Theories, Attention Schema Theory, Recurrent Processing Theory, Electromagnetic Field theories—each with genuine insights and genuine gaps. In 2025, the COGITATE adversarial collaboration, designed to test IIT against GNW, published its results in *Nature*. The outcome? Neither theory was fully confirmed. Posterior cortex showed the strongest consciousness-related activity, which wasn’t quite what either camp predicted. After decades and hundreds of millions of dollars, the field is arguably further from consensus than when it started.

Two Dogmas That Block Progress

Before I tell you what I think is missing, I need to name two prejudices that have been quietly sabotaging the field for decades. I gave them names in my original book because I think unnamed biases are harder to fight.

The first is what I call the **nSAI dogma** — “no strong artificial intelligence.” It’s the widespread conviction that truly intelligent machines are impossible, a conviction rooted not in proof but in the failure of early AI research in the 1960s and the resulting backlash. Anyone who believes strong AI is possible learns to keep quiet about it if they want to be taken seriously in mainstream research. This is not rational skepticism. It’s a scar from old defeats, hardened into doctrine.

The second is deeper and more pernicious. I call it the **nSU dogma** — “no self-understanding.” It’s the belief that the human mind, the human consciousness, cannot in principle be understood by that same mind. People invoke Gödel’s incompleteness theorems, or vague analogies to the limitations of cosmological observation from inside the universe, or — most honestly — they simply find the prospect of being fully explained too frightening to contemplate. If consciousness is just a machine, what happens to the soul? What happens to meaning? What happens to the specialness of being human?

These dogmas reinforce each other. If you can’t understand consciousness (nSU), then you certainly can’t build one (nSAI). And if you can’t build one (nSAI), then maybe consciousness really is beyond understanding (nSU). It’s a closed loop of institutional pessimism, and it has kept an enormous number of intelligent researchers from even attempting the work.

I’m not saying these dogmas are held in bad faith. Many researchers genuinely believe them. But neither dogma has ever been proved. They are articles of faith, and they have done more damage to consciousness research than any failed experiment.

Something Is Missing

I think the reason no theory has cracked the Hard Problem is that most people are looking for consciousness in the wrong place. They're looking at the neural machinery — the neurons, the synapses, the oscillations, the connectivity — and asking: “Which of these processes is conscious?”

The right question, I believe, is different: “On which level of information processing, and using which architecture, does experience occur?”

This is the starting point of the Four-Model Theory. It begins with the observation that you have never, in your entire life, directly experienced reality. Consider the simplest proof: you have a blind spot in each eye — a region of the retina with no photoreceptors at all, where the optic nerve exits — yet you see no hole. Your brain fills it in with fabricated content. If perception were direct access to reality, you would see two dark patches. You don't, because you are looking at a model. This claim, incidentally, is not controversial — that perception is constructive rather than direct is mainstream neuroscience, accepted by virtually every researcher in the field. You have experienced a simulation of reality, generated by your brain, so seamlessly that you have never suspected the difference. And the theory argues that this observation, taken seriously, dissolves the Hard Problem.

Three Guiding Principles

Before we get to the theory itself, I need to lay out three philosophical principles that guided its construction. These aren’t arbitrary methodological choices. They’re constraints that any serious scientific theory should satisfy — constraints that many consciousness theories either ignore or violate.

Occam’s Razor. The simplest explanation that fits the facts is usually the correct one. This is the foundational principle of science, attributed to the 14th-century philosopher William of Ockham. If two theories explain the same phenomena, prefer the one that requires fewer entities, fewer assumptions, fewer special cases. Occam’s Razor doesn’t guarantee truth, but it has a remarkable track record: Newton didn’t need angels pushing the planets; Darwin didn’t need a designer shaping the species; Einstein didn’t need a luminiferous ether. The universe appears to favor simplicity.

The Four-Model Theory is Occamite to its core. It doesn’t introduce any new physical phenomena — no quantum effects in microtubules, no exotic field theories, no panpsychist “proto-consciousness” sprinkled through matter. It uses only what we already know: neural networks, learning, simulation, self-reference. The complexity is in the *architecture*, not in adding mysterious new ingredients.

The Copernican Principle. We are not special. Named for Copernicus, who displaced Earth from the center of the cosmos, this principle has been extended across science: the sun isn’t special, our galaxy isn’t special, and — most uncomfortably for many people — *we* aren’t special. Consciousness is not a unique miracle, a one-off divine spark, or an emergent phenomenon so rare that it could only

happen once. If you have it, other systems can have it too — given the right architecture. This is the anti-exceptionalist stance that makes artificial consciousness possible.

The Copernican Principle is also why this theory predicts consciousness in animals. If a brain architecture can produce consciousness, then any sufficiently similar architecture should produce it. Humans aren't magic. We're just one implementation of a general computational principle.

Leibniz's Law (The Identity of Indiscernibles). If two things are truly identical in all their properties, they are the same thing. This principle, formulated by the 17th-century philosopher Gottfried Wilhelm Leibniz, is both simple and profound. It rules out "zombie worlds" — hypothetical universes physically identical to ours but where nobody has conscious experience. If a system is identical to a conscious system in every functional, structural, and behavioral property, then it *is* a conscious system. There is no extra "consciousness substance" that could be present or absent while leaving everything else unchanged. Consciousness is not an optional add-on to an otherwise complete functional description. It's part of the description.

Leibniz's Law is why philosophical zombies — beings that act exactly like conscious humans but aren't conscious — are incoherent. If the zombie is functionally identical to you, then it has the same four-model architecture, the same simulation running, the same self-reference. At that point, what could "not being conscious" even mean? The question dissolves.

These three principles — simplicity, non-exceptionalism, and identity through properties — aren't just aesthetic preferences.

They're the intellectual tools that let you cut through centuries of confusion and arrive at a theory that actually works. The Four-Model Theory is what you get when you take these principles seriously and apply them to the hardest problem in science.

Now it's time to look at the four models.

Chapter 2

The Four Models

Imagine you're looking at an apple.

The apple is sitting on a table in front of you. Red, round, shiny, about fifteen centimeters from your hand. You can see it, you know what it is, you could reach out and grab it. This seems straightforward — you're seeing an apple.

But what's actually happening is profoundly more complicated.

Light reflected from the apple's surface enters your eyes, where it hits the photoreceptor cells on your retinae. These cells convert the light into electrical signals. The signals travel along your optic nerves to the visual cortex at the back of your brain, where they're processed through a hierarchy of increasingly sophisticated feature detectors: edges, orientations, colors, textures, shapes, and eventually objects. Somewhere in this cascade, the neural activity corresponding to "apple" is activated. Simultaneously, your motor system is preparing potential actions (reaching, grasping), your memory system is activating associations (taste, texture, the last time you ate an apple), and your spatial system is tracking the apple's position relative to your body.

All of this happens in less than a second. And none of it is what you *experience*. You don’t experience photons hitting cone cells, or signals traveling along axons, or feature detectors firing. You experience *an apple*. A unified, stable, three-dimensional object sitting in a coherent spatial environment, with a particular look and feel and meaning. What you experience is a *model* — a real-time simulation of the apple, generated by your brain from the raw data and everything it has previously learned about apples, objects, tables, and physics.

As I argued in Chapter 1, this is uncontroversial neuroscience. Every neuroscientist and philosopher of perception agrees that what you experience is a model, not reality itself. The apple you see is the brain’s *best guess* at what’s out there, informed by the sensory data but not identical to it. (Optical illusions are live proof: when an illusion collapses — when you suddenly see it both ways — you catch the simulation in the act. You were never seeing reality directly. You were always seeing the model. The illusion just made it obvious.)

But here’s where my theory begins: the brain doesn’t just model the apple. It models *you looking at the apple*. And it’s this second model — the model of the self — that turns information processing into consciousness.

Your Brain’s Four Representations

Even simple neural networks with just three layers can learn to model their input — show them enough examples and they build internal representations of the patterns they encounter. Your brain

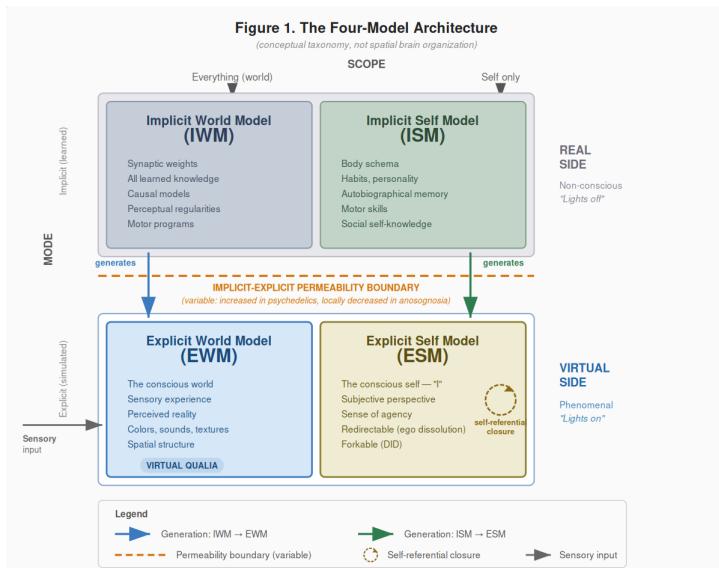


Figure 2.1: The four-model architecture. The brain maintains two kinds of model — one of the world, one of the self — each in two modes: implicit (stored in the structure of the brain) and explicit (actively running as a real-time simulation). Consciousness lives in the explicit models.

does exactly this, but on a vastly richer scale. It doesn’t build one model; it builds many, covering everything from the visual field to the position of your limbs, from the sound of a voice to the pressure of your feet on the ground. These models span both the world outside you *and* your own body, linking all available sensory inputs into coherent representations.

Neuroscience has known about these models for over a century. In the motor cortex and the somatosensory cortex, the body is literally laid out as a distorted map — hands and lips grotesquely enlarged because they have more nerve endings, trunk and legs compressed into slivers. These cortical maps, called *homunculi*, were first charted by Wilder Penfield in the 1930s through direct electrical stimulation during brain surgery. They are just the most vivid examples; the brain maintains similar maps and models throughout its architecture. (See Appendix A for more on cortical organization.)

I call these the **implicit models**: the Implicit World Model (IWM) and the Implicit Self Model (ISM). They are stored in the brain’s structure — in the strengths of synaptic connections, the architecture of neural circuits, the accumulated learning of a lifetime. They are the brain’s hard drive. You never experience them directly, any more than you experience the silicon in your phone. But they encode everything the brain knows about the world and about you.

Now here is the key insight. These implicit models don’t just sit there. They *generate* something. In engineering, a **digital twin** is a real-time virtual replica of a physical system — a jet engine, a power grid, a factory floor — continuously updated with sensor data so engineers can monitor and interact with the system without touching it directly. Your implicit models do exactly this. They

produce a real-time virtual simulation of the world, and a real-time virtual simulation of you. These are the **explicit models**: the Explicit World Model (EWM) and the Explicit Self Model (ESM). Everything you see, hear, feel, and think is happening inside these simulations, not in the world itself.

Two groups of models — implicit and explicit — each containing both a world model and a self model. Four models in total, and with them, a language to talk about what consciousness is actually doing.

But where do these models run? The brain uses at least five levels of information processing, stacked on top of each other. The simulation — your conscious experience — runs at the very top.

Five Nested Systems

Think of your brain as having five distinct levels of organization, stacked like Russian dolls:

Physical. At the bottom, you have the raw matter: atoms, molecules, the physical substrate of the brain itself. This is the chemistry — the carbon, hydrogen, nitrogen, oxygen that compose the tissue. It's inert matter obeying the laws of thermodynamics. Nothing conscious lives here.

Electrochemical. One level up: neural signaling. Action potentials racing down axons, neurotransmitters flooding synapses, ions flowing through channels. This is the electrical and chemical activity that everyone pictures when they think "brain doing something." This is the level where neurons fire. Still no experience, but now you have information transmission.

Proteomic. Next: protein structures and molecular machinery. Synaptic weights are stored here — the physical strengths of connections between neurons. Receptors on cell membranes, enzymes regulating plasticity, the molecular scaffolding that determines which synapses grow stronger and which weaken. This is the “hardware” of learning. When you practice a skill and get better at it, you’re changing the proteomic layer. Still unconscious, but now you have memory.

Topological. Higher still: network architecture. The patterns of connectivity — which neurons connect to which, how densely, in what configurations. This is where Brodmann areas live, where cortical columns live, where the large-scale structure of “visual cortex talks to motor cortex” exists. It’s the wiring diagram. Change this level and you change what kinds of processing the system can do. This is where your implicit models — the IWM and ISM — are stored. Still unconscious. But now you have knowledge.

Virtual. At the very top: the simulated world. The cortical automaton — the dynamic pattern of electrical activity dancing across the network, integrating information, generating predictions, running the models in real time. This is where your conscious experience lives. The explicit models — the EWM and ESM — exist here and only here. This is the only level that feels like anything.

Each level supervenes on the one below it but has its own dynamics. You can’t have electrochemical signaling without physical matter, you can’t have protein structures without chemistry, you can’t have network topology without synapses, and you can’t have a simulation without a network to run it. But each level has properties the lower levels don’t have. A synapse is not “about” anything

— it's just a connection. A network of synapses *is* about something: it represents a face, a word, a memory. And the simulation running on that network? That's where “about” becomes “experience.”

This five-level hierarchy solves a problem that trips up almost everyone when they first hear this theory: “If consciousness is virtual, what’s it running on?” The answer: it’s running on the topological layer (the network), which is implemented in the proteomic layer (synaptic weights), which runs on the electrochemical layer (neural firing), which exists in the physical layer (matter). Consciousness is no less real for being virtual — it’s just real *at a different level* than neurons are real. The mountain in the video game is real at the game level even though it’s “just” transistors at the hardware level. Same principle.

I’ll come back to this hierarchy throughout the book, especially when we talk about psychedelics in Chapter 6 — because drugs don’t hit all five levels equally. Some target the electrochemical layer (altering neurotransmitter dynamics), some target the proteomic layer (changing receptor expression), and the effects ripple up to the virtual layer in predictable ways. The hierarchy isn’t just conceptual. It’s mechanistically real, and it does explanatory work.

Now, the four models.

The Implicit World Model (IWM) is everything you know about the world. Not what you’re currently thinking about — everything you *could* think about. The laws of physics (you know that dropped objects fall). The layout of your apartment (you can navigate it in the dark). The grammar of your native language (you can judge whether a sentence is grammatical without knowing

the rules). The faces of everyone you've ever known. The taste of chocolate. The sound of rain.

All of this knowledge is stored in your brain's synaptic connections — the strengths of the links between neurons. It was built up over your entire lifetime through experience and learning. And you are never, ever directly aware of it. You can't introspect on your neural connections. You can't feel your synapses. The Implicit World Model is like a vast library that you never enter — you just read the books it sends to your desk.

The Implicit Self Model (ISM) is everything you know about yourself. Your body schema — the unconscious representation of where your limbs are, how large they are, how they move. Your motor skills — riding a bike, typing, playing an instrument. Your personality traits, social skills, emotional patterns, habits. Your autobiographical memory structure — the framework that organizes your memories into a life story.

Like the world model, the self model is stored in synaptic weights and is never directly conscious. You don't experience your body schema; you experience the body your schema generates. You don't experience your personality; you experience the thoughts and feelings your personality produces. The Implicit Self Model is the backstage crew — essential to the performance, but never seen by the audience.

The Explicit World Model (EWM) is the world you actually experience. Right now. The room you're in, the sounds you hear, the weight of this book in your hands (or the glow of the screen you're reading it on). This is the simulation — the brain's real-time virtual reality, generated from the Implicit World Model plus cur-

rent sensory input. It's vivid, detailed, and seamlessly convincing. You will live your entire life inside it and never step outside.

The Explicit Self Model (ESM) is *you*. The feeling of being a subject. The sense of “I” — the one who sees, hears, thinks, and decides. This, too, is a simulation: a real-time model generated from the Implicit Self Model plus current body signals. It's the character the brain creates to inhabit its virtual world.

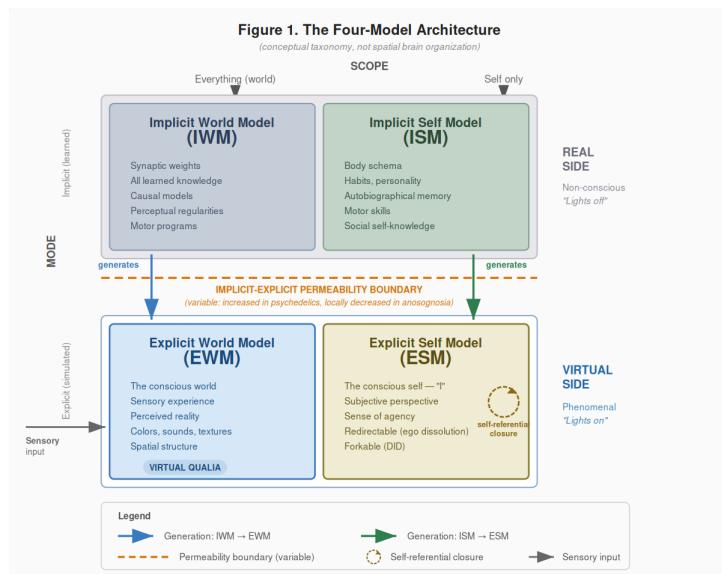


Figure 2.2: The Four-Model Architecture. The four models are arranged along two axes: scope (world vs. self) and mode (implicit/learned vs. explicit/simulated). The implicit models (IWM, ISM) constitute the substrate-level “real side” — learned, structural, non-conscious. The explicit models (EWM, ESM) constitute the simulation-level “virtual side” — transient, generated, phenomenal.

The Real Side and the Virtual Side

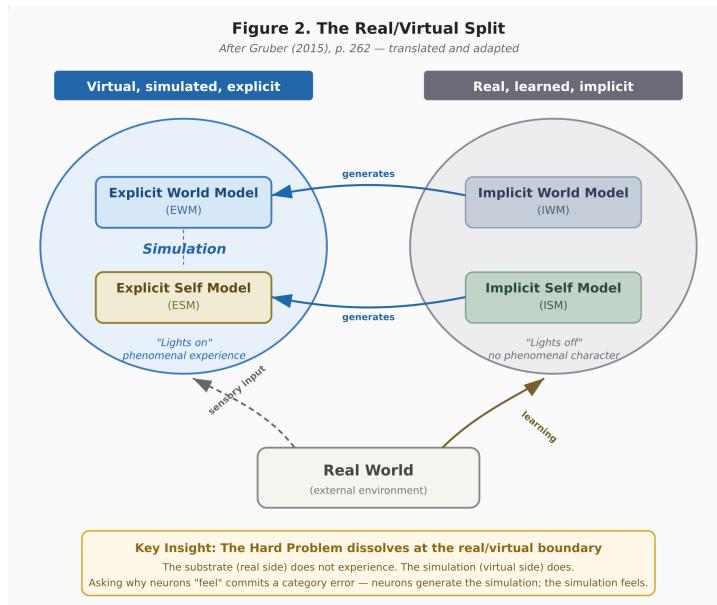


Figure 2.3: The real/virtual split. The substrate (real side) stores knowledge in synaptic weights — physical, structural, unconscious. The simulation (virtual side) generates experience in real time — transient, dynamic, conscious. Everything you have ever experienced lives on the right side of this line.

The four models divide into two sides, and this division is the foundation of everything that follows.

The **real side** — the two implicit models — is physical, structural, and relatively rigid (adapted by learning). It’s the brain’s stored knowledge: synaptic weights, network connections, receptor configurations. Think of it as everything the brain *has learned* — crystallized into the physical structure of the tissue itself. It has no

experience. A synapse firing is no more “experienced” than water flowing through a pipe. The real side is lights off.

Here’s something important: the real side is what neuroscience already studies. When a researcher puts you in an fMRI scanner, they’re looking at the real side — firing patterns, connectivity, blood flow to different regions. When a neurosurgeon stimulates a cortical area and watches what happens, they’re probing the real side. Neuroscience has been mapping this territory for over a century, and it has made extraordinary progress. The Four-Model Theory is not rejecting any of that work. It’s saying that all of it describes only half the picture.

The **virtual side** — the two explicit models — is simulated, transient, and dynamic. It’s generated anew in every moment from the real side plus current input. Think of it as everything the brain *is currently doing with* what it has learned — the live show, not the stored script. And it is *all* of experience. Every sight, sound, thought, feeling, memory, dream, and hallucination you have ever had has occurred within the virtual side. The virtual side is lights on.

But here’s the catch: the virtual side is invisible from outside. Even our most advanced brain imaging can only capture it indirectly. An fMRI shows you which brain regions are active — that’s the real side doing its work. To actually *read* conscious experience from brain data, you would need to decode the programming language of the brain — to understand not just which neurons fire but what the pattern of firing *means* at the simulation level. That would require something like a full simulated connectome: a complete

digital replica of the brain, run in software, producing the same virtual world the biological brain produces.

I want to be honest about what the theory does and doesn’t give you. The Four-Model Theory tells you *what* the simulation is, *where* it lives, and *why* it feels like something. It does not hand you the decoder ring. Reading the virtual side from the real side is a future research programme — one that the theory defines clearly but cannot execute yet. However, the foundation for that programme is already being laid. The Human Connectome Project and related efforts are mapping the brain’s wiring at increasingly fine resolution. We cannot yet decode the virtual side from structural data — but the structural data is coming.

If you’re scientifically minded, you might already see where this is going. If experience exists only on the virtual side, then looking for experience on the real side — in the neurons, in the synapses, in the physical machinery — is looking in the wrong place entirely. It’s like searching for the plot of a movie inside the DVD player’s circuits.

That’s the key. Let me spell it out.

How Conscious Are You?

Before I do, there’s something you’ve probably already been wondering. If consciousness is a simulation — a virtual self inside a virtual world — then it’s not an all-or-nothing thing, is it? A simulation can be more or less detailed. A self-model can be more or less sophisticated. Which means consciousness comes in *degrees*.

The Four-Model Theory gives you a precise way to think about those degrees. There are four graduated levels, and every conscious creature sits somewhere on this ladder.

At the bottom, you have **basic consciousness**. This is an Explicit World Model with only a rudimentary Explicit Self Model. The system generates a virtual world — there is something it is like to be this creature — but the self inside that world is barely sketched in. Think of a mouse navigating a maze. It sees the walls, smells the cheese, feels the floor under its paws. It has phenomenal experience. But its model of *itself* as the thing having those experiences? Paper-thin. There is a “what it’s like,” but almost no “who it’s like it for.”

One step up: **simply extended consciousness**. Now the self-model gets real. The system doesn’t just experience — it models itself *as* the experiencer. It is aware that it is experiencing. Your dog doesn’t just feel pain; your dog knows that *it* is in pain. There is a first-person perspective — a genuine “me” at the center of the virtual world. This is first-order self-observation, and it changes everything. Suffering becomes possible here, because suffering requires a self that knows it suffers.

Then: **doubly extended consciousness**. Second-order self-observation. The system models itself modeling itself. This is metacognition — thinking about your own thinking. You’re lying in bed wondering whether your anxiety about tomorrow’s meeting is rational or whether you’re catastrophizing. You’re monitoring your own mental states, evaluating them, sometimes overriding them. This is where most adult human consciousness lives most of

the time. It’s the level that makes therapy possible, that allows you to say “I notice I’m getting angry” instead of just being angry.

And at the top: **triply extended consciousness**. Third-order. The system models itself modeling itself modeling itself. This sounds like a hall of mirrors, and it is — but it’s a hall of mirrors you need in order to do philosophy of mind. To ask “what is consciousness?” you need to model yourself, model your experience, and then model yourself modeling that experience. You need to step back far enough to see the whole apparatus from the outside, even though you’re still inside it. This is the prerequisite for the question you’re reading this book to answer. Only creatures capable of triply extended consciousness can wonder why anything feels like anything.

Here’s the payoff: this gradient isn’t just abstract philosophy. It answers the question everyone asks me at dinner parties — “Is my dog conscious?” The answer is yes, but less conscious than you are. Your dog is probably at the simply extended level. It has a self. It has experience. It does not lie awake at 3 a.m. questioning the nature of that experience. We’ll come back to the animal question in detail in Chapter 10, where this gradient does real explanatory work. But you can already see the shape of it: consciousness is not a light switch. It’s a dimmer.

Why Your Brain Has the Capacity for Self-Modeling

So we’ve established that consciousness depends on these four models, with the explicit self-model doing the heavy lifting. But why does the human brain have this capability in the first place,

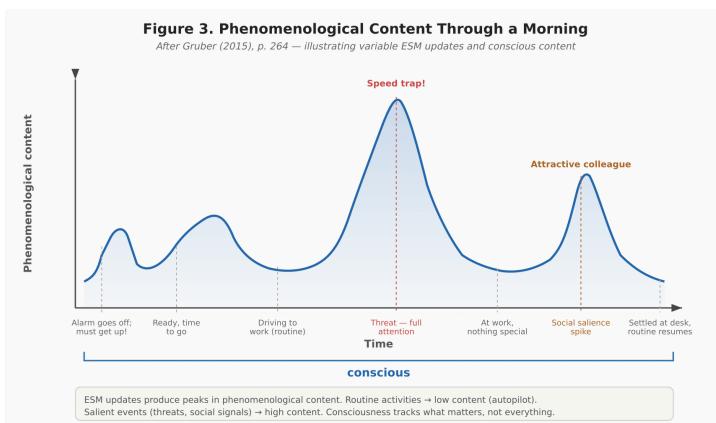


Figure 2.4: Phenomenological Content Through a Morning. ESM updates produce peaks in phenomenological content. Routine activities lead to low content (autopilot). Salient events (threats, social signals) produce high content. Consciousness tracks what matters, not everything.

when simpler animals don’t? The answer is hiding in plain sight: the architecture of the human cortex is, quite literally, oversized for basic information processing.

The human neocortex has six layers. This is a well-known anatomical fact — you can see it in any neurobiology textbook. But here’s what’s interesting: you don’t need six layers to process information. Three layers will do the job.

Think about what a standard neural network needs to do. First layer: receive input, filter it, clean it up. Second layer: extract patterns, recognize features, do the heavy computational lifting. Third layer: integrate results, make decisions, produce output. Input, processing, output. That’s the basic recipe, and three layers cover it.

But we have six.

What are the “extra” three layers for?

They’re for modeling the first three.

A three-layer network processes the world. A six-layer network processes the world *and* observes itself doing it. The additional layers provide the architectural capacity for the brain to build not just a model of what’s out there, but a model of itself modeling what’s out there. Self-simulation requires this doubling — you need one set of layers to do the processing, and another set to watch the processing happen.

This isn’t speculation about what individual layers “do” — I’m not claiming Layer 4 does this and Layer 5 does that. It’s an observation about architectural capacity. Six layers give you room for both the implicit world model (the learned, unconscious processing) and the explicit world model (the real-time simulation). They

give you room for both the implicit self model (your body schema, motor programs, personality structure) and the explicit self model (the “you” that experiences having a body, initiating actions, being a person).

Now look at other animals. Reptiles have three or four cortical layers. Mammals have six. And among mammals, the ones with the thickest, most elaborately folded cortex — primates, cetaceans, elephants — are exactly the ones that show the richest signs of self-awareness. Mirror self-recognition, future planning, social deception, grief. The architectural capacity tracks the phenomenology.

The jump from three to six layers may have been a genetic duplication accident — evolution’s copy-paste producing the very architecture that consciousness would later exploit. Reptilian ancestors had three cortical layers. Somewhere in the transition to mammals, that number doubled. The transcription factors that specify cortical layer identity — *Tbr1*, *Satb2*, *Ctip2*, *Fezf2* — have paralogs suggestive of gene duplication events. Whether this was a single dramatic event or a gradual elaboration remains debated, but the result is clear: mammals got double the layers, and with them, the capacity for self-modeling that most reptiles lack.

This is the bridge from neural network theory to lived experience. The human cortex isn’t just a big pattern recognizer. It’s an oversized, recursively structured network with enough layers to model its own modeling process. And when a network models itself modeling the world, the result — viewed from inside — is exactly what we call consciousness.

I should be clear: I’m not claiming that six cortical layers are the *only* architecture capable of supporting consciousness. They’re

one solution — the one mammals evolved. But there may be others. The octopus, with its radically distributed nervous system — eight semi-autonomous arms, each with its own neural processing center containing roughly 40 million neurons — represents a completely different architectural approach that may achieve equivalent computational power. Birds offer another striking example: corvids and parrots lack a layered cortex entirely, their pallium organized into nuclear clusters rather than sheets — yet crows make tools, plan for the future, and arguably recognize themselves in mirrors. If what matters is the capacity for self-modeling, not the specific wiring diagram, then any architecture that can run a simulation of itself could in principle be conscious. We’ll return to this in Chapter 10.

Chapter 3

The Virtual Side

Imagine you’re playing a video game. A good one — an immersive open-world game with stunning graphics, realistic physics, and a compelling story. You’re controlling a character, and through that character, you’re interacting with a richly detailed virtual world.

Now consider: where does the game exist? Not on the screen, exactly — the screen just displays light patterns. Not in the graphics card or the CPU, exactly — those are running electrical signals through silicon circuits. The game exists as a *virtual process* — a higher-level phenomenon that arises from the hardware’s activity but is not identical to any particular piece of hardware.

The virtual world of the game has properties that the hardware does not. The game has mountains, rivers, and cities. The CPU has transistors. The game has a day-night cycle. The GPU has clock cycles. You can meaningfully ask “How tall is that mountain in the game?” but it would be absurd to point to a transistor and say “This transistor is 3,000 meters tall.” The game’s properties exist at the virtual level, and they are real properties of the game, even though the game is “just” a pattern of activity in the hardware.

This is not a metaphor. This is how your brain works.

Your Explicit World Model — the world you experience — is a virtual process running on neural hardware, just as the game world is a virtual process running on silicon hardware. The experienced world has properties (colors, shapes, distances, sounds) that the neural hardware does not have (the hardware has firing rates, synaptic strengths, and neurotransmitter concentrations). The properties of your experienced world are *real properties of the simulation*, even though the simulation is “just” a pattern of neural activity.

And your Explicit Self Model — the “you” experiencing the world — is also a virtual process. It is as real as the game character in the analogy: genuinely existing at the virtual level, genuinely having properties at the virtual level, but not existing at the hardware level.

Why the Analogy Breaks Down (In the Important Way)

The video game analogy is useful, but it breaks down at a crucial point: the game has a *player*. There is someone outside the game — you, sitting on the couch — who experiences the game. The game itself has no experience. It’s just patterns of light and code.

Your brain’s simulation has no outside player. There is no one sitting outside your skull experiencing the simulation. The simulation contains its own observer — the Explicit Self Model. The simulation *is* the experience, not something experienced by someone else.

Put yourself in the game character’s position. You *are* the main character. From outside the game, a spectator sees pixels moving

on a screen — nothing that could possibly feel anything. But from inside the simulation? The game world is all there is. The mountains are real to the character, the sunlight is warm, the danger is frightening. No outside observer would ever guess that this pile of code feels anything — but that's because they're looking at the wrong level. They're looking at the hardware. The experience exists at the software level. That's my claim, and the rest of this book lays out the evidence.

The simulation looking at itself. Your entire visual field — every color, shape, and shadow — is generated by the brain's real-time virtual model. At the edges, the illusion thins and the neural machinery becomes visible. There is no boundary between the observer and the observed. You are the simulation.

This is what makes consciousness special and what makes the Hard Problem seem so intractable. In the video game, there's a clean separation between the game (virtual, no experience) and the player (physical, has experience). In the brain, there is no separation. The simulation and the experiencer are the same thing. The Explicit Self Model is not watching the Explicit World Model from outside — it's *inside* the simulation, part of the same virtual process.

And this self-referential closure — the simulation observing itself from inside — is, I argue, what we call consciousness. It's not something added to the simulation. It's what the simulation *is*, when it includes a model of itself. This is why I say consciousness is not a thing — it's a process. You won't find it by taking the brain apart, any more than you'd find a running program by disassembling the CPU.

The Software Properties

If the virtual models really are software-like processes running on neural hardware, then they should behave like software in specific, testable ways. And they do. Four properties of the virtual side will reappear throughout this book, so let me lay them out now.

Forking. A single substrate can run multiple virtual configurations simultaneously. In software, you fork a process and get two independent instances running on the same hardware. In the brain, this is Dissociative Identity Disorder — multiple self-models, each with its own narrative and emotional profile, alternating control of the same neural substrate. We’ll see this in Chapter 9.

Cloning. Physically separate the hardware, and you get degraded but complete copies of the software. Cut the corpus callosum, and each hemisphere runs its own version of the simulation — less capable than the original, but functionally whole. That’s the split-brain phenomenon, also Chapter 9.

Redirecting. Disrupt the normal input stream and the simulation latches onto whatever signal dominates. Under salvia divinorum, proprioceptive input overwhelms the system and the Explicit Self Model reconfigures around body sensation. Under ketamine, external input drops out and the simulation runs on internal noise. The virtual models don’t stop — they just process whatever they’re fed. Chapter 6 covers this in detail.

Reconfiguring. Modify the substrate’s connection weights and you change what the virtual models produce. This is exactly what Cognitive Behavioral Therapy does — systematically rewiring the substrate so the Explicit Self Model generates different narratives, different emotional responses, different behavior.

The Four-Model Theory makes a specific prediction about therapy: any effective treatment must work by modifying the implicit models (the substrate) such that the explicit models (the simulation) change accordingly. CBT does exactly this — it systematically identifies maladaptive patterns in the ISM and rewires them through structured practice, changing what the ESM produces. This is why CBT has the strongest evidence base of any psychotherapy: it targets the right level.

This raises an uncomfortable question about therapies that can't explain their mechanism in these terms. If a therapeutic approach doesn't specify what it's changing in the substrate, or how that change propagates to the simulation, then at best it's working through a mechanism it doesn't understand, and at worst it isn't working at all. The evidence bears this out: the therapies with the weakest evidence bases are generally the ones with the vaguest theories of change. If you're seeking therapy, ask your therapist a simple question: "What specifically are you trying to change in my brain, and how?" If they can't answer, consider finding one who can.

These aren't metaphors. They're structural predictions. If my theory is wrong and the virtual models are *not* software-like processes, then these parallels are pure coincidence. But coincidences don't usually line up four-for-four across clinical neurology, psychopharmacology, and psychotherapy. The chapters that follow will show each property in action.

There's a simple experiment you can do right now — well, with a friend, a rubber hand, a cardboard screen, and two paintbrushes — that demonstrates how easily the Explicit Self Model can be tricked.

It’s the rubber hand illusion, devised by Matthew Botvinick and Jonathan Cohen, and it’s one of the most revealing party tricks in all of neuroscience.

Here’s how it works. You sit at a table with one arm hidden behind a cardboard screen. A realistic rubber hand is placed in front of you, visible, roughly where your hidden hand would be. Someone simultaneously strokes the rubber hand and your hidden real hand with two paintbrushes, in the same location, at the same speed. After a minute or two of this synchronized stroking, something uncanny happens: you start *feeling* the brush strokes on the rubber hand. Not on your real hand, behind the screen. On the fake hand in front of your eyes.

Your Explicit Self Model has incorporated the rubber hand into its body schema. It has reassigned ownership — decided that the rubber hand is part of “you.” The self-model is not hardwired. It’s learned. It’s updated continuously based on the best available evidence, and when the visual evidence (seeing the rubber hand being stroked) consistently matches the tactile evidence (feeling your real hand being stroked), the ESM draws the rational conclusion: that hand is mine. If someone then threatens the rubber hand — brings a hammer down toward it — you flinch, you feel a spike of anxiety, your galvanic skin response shoots up. For the part of your brain that defines “you,” that hand *is* yours.

This is not a glitch. This is the self-model working exactly as designed — constantly updating its body boundary based on multimodal sensory correlation. It’s the same mechanism that lets amputees “feel” a prosthetic limb as their own after a period of use. And it’s the same mechanism that breaks down in asomatognosia,

where patients deny ownership of their actual limbs, and in the Alien Hand Syndrome, where the hand moves on its own.

The Patchwork Hologram

There's a fifth property of the virtual side that deserves its own section, because it explains something that has puzzled neuroscientists for nearly a century: why brain damage degrades function *gradually* rather than deleting specific memories.

In the 1920s and 30s, the psychologist Karl Lashley trained rats to navigate a maze, then surgically removed pieces of their cortex to find where the memory was stored. He never found it. No matter which piece he removed, the rats still remembered the maze. What mattered was *how much* cortex he removed, not *which parts*. Remove a little, and the rats got slightly worse. Remove a lot, and they got much worse. But the memory was never just *gone*, cleanly excised like a file deleted from a hard drive. Lashley spent his career searching for the “engram” — the physical trace of a memory — and famously concluded that it didn’t seem to exist.

He was looking for the wrong thing. The memory wasn't stored *in* a particular piece of cortex the way a file is stored on a particular sector of a hard drive. It was stored *across* the entire network, distributed in the connection weights between millions of neurons. This is how neural networks work: information isn't sitting in any one node. It's encoded in the pattern of connections between all of them. You can't point to a single synapse and say “this is where the maze is stored” any more than you can point to a single pixel and say “this is where the movie is stored.”

This is essentially a holographic property. If you take a physical hologram and cut it in half, you don’t get two halves of the image. You get two copies of the *complete* image, each at lower resolution. Cut it into quarters and you get four complete images, blurrier still. The information in a hologram is distributed across the entire plate, so every piece contains the whole picture — just with less detail.

Neural networks do the same thing. Train a network to recognize faces and then destroy 10% of its connections at random. It doesn’t forget 10% of the faces. It gets slightly worse at *all* faces. Destroy 50% and it gets substantially worse at everything, but it still recognizes something. The information is smeared across the whole network, which is exactly why Lashley couldn’t find the engram: it was everywhere and nowhere.

But — and this is where it gets interesting — the brain isn’t *one* hologram. It’s what I call a *patchwork hologram*. Within a single functional area (say, your primary visual cortex, roughly Brodmann area 17), the cortical columns are similar to each other, and information is stored holographically. Destroy a few columns and you barely notice. The area is locally holographic — a part contains the whole, at lower resolution.

But at the global level, different areas do different things. Your visual cortex is not interchangeable with your motor cortex. Remove the entire visual cortex and you lose vision — there’s no blurry backup. So the brain is locally holographic within each functional region, fractally self-similar in its columnar architecture, but globally *not* holographic. It’s a patchwork: dozens of holographic tiles stitched together into a composite that is, as a whole, decidedly non-holographic.

This patchwork structure explains a pattern you see over and over in clinical neurology. Small strokes and small lesions often cause surprisingly mild deficits — because within any given cortical area, the holographic principle protects you. The remaining tissue reconstructs the missing information at lower resolution. But large strokes that wipe out an entire functional area cause catastrophic, specific losses — blindness, paralysis, aphasia — because you've removed an entire tile from the patchwork, and no other tile can substitute.

It also explains why memories don't just "pop out of existence" when neurons die. Every day, neurons die and synapses are pruned. If memories were stored like files on a hard drive, you'd expect to occasionally lose one — to wake up one morning having forgotten your wedding, or your childhood dog, or the taste of coffee. That never happens. Instead, memories fade gradually, losing detail and vividness over years. That's exactly what a holographic storage system predicts: degradation is graceful, proportional, and global, never sudden, discrete, or local.

The patchwork hologram is the physical reason why the software properties I described above — especially cloning — actually work. Split the brain in half, and each half retains a degraded but complete copy of the simulation, because within each hemisphere, the holographic principle ensures that every piece contains the whole picture. The simulation doesn't break. It just runs at lower resolution.

Chapter 4

Why It Feels Like Something (And Why That's the Wrong Question)

Now we can tackle the Hard Problem directly.

The question is: **Why does physical processing feel like something?**

The answer: **It doesn't.**

The physical processing — neurons firing, synapses transmitting, the implicit models storing and computing — has no experience. None. There is nothing it is like to be the real side. The real side is precisely the “in the dark” processing that the Hard Problem assumes consciousness needs to explain.

The *simulation* feels. The Explicit World Model and the Explicit Self Model — the virtual side — are where experience lives. And within the simulation, experience is not a mysterious addition to the process. Experience is what the simulation *is*, when it includes a self-model. The Explicit Self Model “perceiving” the Explicit World

Model is what we call qualia. Qualia are the virtual self's mode of registering the virtual world.

Think about it this way. If you asked "Why does transistor switching feel like running a video game?" the answer would be: "It doesn't. Transistor switching doesn't feel like anything. The game is a virtual process that runs on transistors but has properties the transistors don't have — landscapes and characters and physics and light. Those properties are real properties of the virtual process, not of the transistors."

Similarly: neuronal firing doesn't feel like seeing red. Neuronal firing generates and sustains a simulation, and within that simulation, the self-model perceives a certain class of world-model content as what we call "redness." Redness is a real property of the simulation, not a property of the neurons.

The Hard Problem assumed that we need to explain how physical processing produces experience. But physical processing doesn't produce experience — it produces a *simulation*. And the simulation, because it includes a self-referential loop (the ESM modeling itself within the EWM), constitutively *is* experience.

But Wait — Isn't This Circular?

The obvious objection: "You've just moved the problem. Why does *this* simulation have experience, when a weather simulation doesn't?"

The answer is self-reference. A weather simulation models weather. It does not model *itself*. There is an "outside" to a weather simulation — the computer, the programmer, the scientist inter-

preting the output. The simulation can be fully described without referring to any experience, because there is no self-model inside it.

The brain’s simulation models itself. The Explicit Self Model is the simulation’s model of *its own process*. This creates a closed loop: the model and the thing being modeled are the same system. There is no “outside” from which the simulation can be fully described, because the describer is part of the description.

This is not magic. This is a structural consequence of self-reference. When a process models itself, the distinction between the model and the modeled collapses. The process of self-modeling and the experience of being a self are not two different things that need to be connected by a bridge — they are one and the same thing, described in different vocabularies.

The Hard Problem asks for a bridge between physical processing and experience. The Four-Model Theory says: there is no bridge, because they were never separate. The experience IS the self-simulation, viewed from inside the loop.

This is ultimately an identity claim — the kind of claim that, in science, marks a resting point rather than a gap. “Water is HO” is an identity. You cannot meaningfully ask “But *why* is water HO?” — the identity *is* the explanation. Asking for something deeper is asking for a different kind of universe. Similarly: experience is what four-model self-simulation at criticality *is*. If someone asks “But *why* does this self-simulation feel like something?” the answer is: because that’s what this process *is*. The identity is falsifiable — if the predictions in Chapter 11 fail, the identity is wrong. But it cannot be “further explained,” any more than the molecular identity of water can be further explained. It is the stopping point.

But Couldn't the Simulation Run "In the Dark"?

There's a subtler objection that's worth addressing head-on. Grant that the brain runs a self-simulation. Grant the four-model architecture, the criticality, the self-referential closure. Couldn't all of that happen without there being anything it's *like*? Couldn't the simulation evaluate, model, predict — and feel nothing?

This is the zombie intuition in technical clothing, and the answer is no. Here's why.

The substrate deploys the virtual simulation as its evaluation mechanism. That's the primary direction of traffic: the implicit system presents situations to the simulation so the simulation can assess consequences and register outcomes. But for that evaluation to work, the simulated states must have *valence* — they must matter to the simulation. A pain signal that's just a number doesn't drive avoidance at the simulation level. Only a simulation that *cares* about outcomes can evaluate them.

Think of a digital twin — an engineering simulation of a jet engine. A typical digital twin doesn't just mirror the engine passively. It *adds* a visualization layer: warnings, color-coded indicators, alarms — things that don't exist in the physical engine. The engine has metal fatigue; the twin has a flashing red warning. The engine has rising temperature; the twin has a gauge turning from green to amber to red. That added layer is the whole point. Without it, the twin is a spreadsheet — numbers sitting inertly in memory, technically accurate, functionally useless. The visualization is what makes the simulation an *evaluation tool*.

Your brain does the same thing, but more so. The conscious simulation doesn't just mirror the substrate's processing — it *adds* phe-

nomenal valence. Pain, pleasure, urgency, curiosity, dread, delight — these are the brain’s equivalent of warning lights and dashboard indicators. They don’t exist at the substrate level (neurons don’t feel pain any more than metal feels fatigue). They exist at the simulation level, added *by* the simulation so the system can evaluate complex situations at a glance. The substrate needs the simulation to assess novel, ambiguous scenarios — the kind where reflexes won’t suffice. And for that assessment to work, the simulated self must register hedonic valence: threat, opportunity, consequence. That registration — that *mattering* — is phenomenality. Remove the qualia and you remove the evaluation — like ripping the display off a cockpit dashboard and expecting the pilot to fly by reading raw sensor voltages.

“But a reinforcement learning system has reward signals that drive behavior,” you might object. “Does it feel?” No — because it lacks the four-model architecture at criticality. An RL reward signal is a scalar value in a Class 1 or Class 2 system. Phenomenal valence is the ESM’s registration of consequence within a full self-simulation running at Class 4 dynamics — a qualitatively different process. The difference isn’t degree. It’s architecture.

The simulation can’t run dark because darkness would defeat its purpose. Phenomenality isn’t a bonus feature of consciousness. It’s the mechanism by which the simulation does its job.

But Wait — Aren’t You Just Saying Consciousness Is an Illusion?

No. And this matters enough that I want to be blunt about it.

There is a respectable philosophical position called illusionism, associated with Daniel Dennett and Keith Frankish, which holds that qualia are illusions. On this view, there is nothing it is like to see red. The appearance of experience is itself a fiction — a story the brain tells, with no experiential reality behind it. Consciousness, in the strongest sense, doesn't exist. It just seems to.

Think about what that actually claims. If you feel something right now — curiosity about this argument, skepticism, the weight of the book in your hands — illusionism says that feeling is an illusion. You're not really experiencing anything. When you say "I feel something," you are, according to this theory, mistaken. Your own testimony about your own experience is wrong. You are, in effect, lying — except there's no "you" to be lying. If that strikes you as obviously ridiculous, I agree.

The Four-Model Theory says the opposite.

Qualia are real. They are real within the simulation. They are the virtual self's mode of perceiving the virtual world. When your Explicit Self Model registers your Explicit World Model's representation of a red apple, that registration — that "seeing redness" — is a genuine property of the virtual process. It exists at the simulation level, just as a bullet hitting a video game character *hurts* it. Not metaphorically — within the game, the damage is real. The health drops, the character staggers, the world responds. From outside, it's a number decrementing in memory. From inside the game, it's pain. That's the level difference. And that's where your qualia live.

The theory operates with a two-level ontology. The substrate level — the neurons, the synapses, the implicit models — has no experience. It is lights off. The simulation level — the explicit

models, the virtual world and virtual self — has genuine experience. It is lights on. Both levels are physical. Neither is an illusion. They are different levels of the same physical system, with different properties at each level.

The theory doesn’t say your pain is an illusion. It says your pain is real — it’s just real in the simulation, not in the neurons. And since you live your entire life inside the simulation, that’s the only kind of real that matters to you.

This is the crucial distinction. Miss it and you’ll confuse this theory with eliminativism, with illusionism, with every other framework that tries to explain consciousness by explaining it away. The Four-Model Theory doesn’t explain consciousness away. It explains where consciousness lives — and it turns out to be exactly where you’ve been standing all along.

“Real Within the Simulation” — What Does That Actually Mean?

If you’ve been following carefully, you might see a trap. A philosopher could argue: when you say qualia are “real within the simulation,” you must mean one of two things. Either they are *genuinely phenomenal* — in which case you’ve just relocated the mystery from neurons to the simulation, and the Hard Problem lives on at a different address — or they are *functionally real but not genuinely phenomenal* — in which case you’re Dennett with extra steps.

This is a false dichotomy. It only holds if you maintain that there’s a god’s-eye view from which to adjudicate whether something is “genuinely” phenomenal — an outside perspective that

can check whether the simulation really feels or merely acts as if it does. But self-referential closure eliminates exactly this outside perspective. The ESM is its own observer. There is no external vantage from which to ask “but does it *really* feel?” The asking is itself part of the process.

“Genuinely phenomenal” versus “merely functional” presupposes that phenomenality is a property a process either has or doesn’t have, checkable by an independent observer. For a fully self-referential system at criticality, there is no such observer. The question dissolves — not because it’s unanswerable, but because it’s unaskable. It requires a perspective that self-referential closure makes impossible.

This is the strongest move available within process physicalism, and it’s the position Thomas Metzinger gestures toward with his concept of “phenomenal transparency” — though the Four-Model Theory is more explicit about *why* the transparency arises. The implicit-explicit boundary is what creates the transparency: you cannot see through it, so you cannot step outside your own phenomenality to ask whether it’s “genuine.” The boundary isn’t a bug. It’s why the question about genuine versus merely functional doesn’t apply to systems like you.

Why the Mystery Persists

Even after dissolving the Hard Problem, there’s a lingering question that nags at people. If the answer is so clean, why does consciousness still *feel* so mysterious? Why does the Hard Problem seem hard even after you’ve been told the solution? David Chalmers calls this

the “meta-problem of consciousness” — the problem of explaining why we *think* there’s a hard problem.

The Four-Model Theory has a clean answer, and it falls straight out of the architecture.

Here’s the strange part: the conscious “you” — the virtual self — cannot see the machinery that generates it. You can’t introspect on your own synaptic weights any more than a character in a dream can examine the dreamer’s brain. The system that creates your experience is, by its very nature, invisible to your experience. Not because someone is hiding it, but because it operates at a level your experience doesn’t include.

Think of it this way. You’re a character in a video game — a really good one, with full self-awareness inside the game world. You can see the rendered mountains, hear the rendered wind, feel the rendered ground under your feet. But you almost never see the graphics engine. You almost never catch a glimpse of the source code. The rendering process operates at a level the game world doesn’t usually include. I say “almost” because sometimes artifacts leak through. In your brain, this happens too — psychedelics open the boundary, flow states thin it, and even in normal life you can catch glimpses: the blind spot your brain fills in, phosphenes when you rub your eyes, the geometric patterns behind your closed eyelids. These aren’t glitches. They’re moments when the substrate’s processing becomes briefly visible from inside the simulation. We’ll explore this in detail in Chapter 6. But most of the time, the rendering process is hidden from the rendered world.

This is exactly the ESM’s predicament. When the conscious self tries to understand the basis of its own experience, it encounters

a principled opacity — not a gap in current knowledge, but a structural feature of the architecture. The implicit models that generate the simulation are not part of the simulation. They can't be, any more than the GPU can be a mountain in the game.

The result is predictable. The ESM, unable to observe its own substrate, concludes that the mechanism of consciousness must be non-physical, or fundamentally inexplicable, or somehow beyond the reach of science. This is the origin of dualism. This is the “explanatory gap.” This is the persistent intuition that something is being “left out” of every physical explanation of consciousness — because from inside the simulation, something *is* being left out. The substrate. The very thing that generates the experience is invisible to the experience it generates.

The mystery is real — but it's an artifact of architecture, not evidence of something non-physical. And there's a reason it *feels* mysterious. You are a virtual process running on biological hardware, and most of the time, the boundary between you and your substrate is opaque. But not always. Sometimes — in altered states, in moments of extreme focus, in the corner of your eye — you catch a glimpse of the machinery underneath. Not clearly, not fully, but enough to sense that something vast is going on below the surface of your experience. That uncanny feeling, that sense that consciousness is somehow deeper than you can reach — that's what it feels like to be a simulation that almost, but not quite, sees through its own curtain.

This is a *prediction* of the theory, not a loose end. If you're a simulation with a mostly-opaque boundary to your own substrate, you'd *expect* consciousness to feel exactly as strange and irreducible

as it does. The Hard Problem’s intuitive force doesn’t come from consciousness being genuinely inexplicable. It comes from our architectural position — we’re inside the simulation, peeking through cracks.

Who Are You When You Wake Up?

Here’s a thought experiment that cuts deeper than it first appears. What if you woke up tomorrow with different memories, a different personality, a different sense of your own body? Would you still be “you”?

Most people’s instinct is to say no — obviously, if everything about my inner life changed, then “I” would be gone and someone else would have taken over. But the Four-Model Theory says something more unsettling: this *already happens* to you, slightly, every single day.

Every night, your Explicit Self Model collapses. Deep sleep erases the running simulation. When it reboots in the morning, it reconstructs “you” from the Implicit Self Model — the stored substrate. But the substrate has changed overnight. Dreams you don’t remember have modified synaptic weights. Consolidation processes have rearranged memories. You wake up not quite the same person who fell asleep. The difference is usually so small you never notice — but it’s there.

In extreme cases, you *do* notice. If you’ve ever woken from deep unconsciousness — after fainting, after a knockout, after anesthesia — in an unfamiliar location, you may have experienced something genuinely strange: a few seconds where you didn’t know *who you*

CHAPTER 4. WHY IT FEELS LIKE SOMETHING (AND WHY THAT'S THE WRONG QUESTION)

were. The Explicit Self Model was booting up, searching the unfamiliar environment for associations to anchor itself, and finding none. For those seconds, there was awareness — you were *someone* — but not yet you. The self-model hadn't finished loading.

This tells us that identity is not a fixed property of the substrate. It's a *reconstruction*, assembled fresh each morning from the stored self-model. The continuity of "you" across time is maintained by two things: the stability of the Implicit Self Model (which changes slowly), and sleep (which prevents you from noticing the gradual drift). If someone could modify your ISM dramatically overnight — replace your memories, reshape your personality structure — the old "you" wouldn't vanish. It would be absorbed. Your new Explicit Self Model would reconstruct a continuous narrative from whatever memories remain, binding the old and new personas into a single story. This is what your brain already does every night on a smaller scale: the substrate changes during sleep, and the ESM that boots up in the morning seamlessly confabulates itself as the same person who went to bed. The only difference is the magnitude of the change. The ESM doesn't do clean breaks — it *always* stitches a continuous narrative. Only if the old memories were completely erased would the thread snap entirely. As long as something remains, the new "you" will incorporate the old "you" into its history, seamlessly, without even noticing the seam.

Chapter 5

At the Edge of Chaos

So far I've told you what the architecture looks like — four models, two axes, a simulation running on a substrate. I've told you where experience lives — on the virtual side, in the explicit models. And I've told you what identity is — a reconstruction, assembled fresh each morning from stored implicit models.

But I haven't told you what makes the whole thing *run*. Why is the simulation sometimes on and sometimes off? What physical property distinguishes a conscious brain from an unconscious one? Why does deep sleep erase the simulation while the architecture stays intact?

There's one more piece of the puzzle, and it's the one that really convinced me the theory works.

The four-model architecture is necessary for consciousness, but it's not sufficient. You also need the right *dynamics*. Specifically, the substrate — the physical system running the simulation — must operate at what mathematicians and physicists call the **edge of chaos**.

In 2002, the polymath Stephen Wolfram published *A New Kind of Science*, in which he classified computational systems into four

types based on their dynamics. I think Wolfram's scheme needs a fifth class — he lumped fractal systems together with truly chaotic ones, but they are structurally distinct. The full argument is in Appendix C, for readers who want the mathematical details. Here, the essential point is this:

Computational systems fall on a spectrum from perfect order to perfect disorder. At one end, static and periodic systems — too simple to compute anything interesting. At the other end, chaotic systems — too disordered for any stable patterns to form. In between, at the **edge of chaos**, sit the systems capable of universal computation: complex enough to produce rich, varied, unpredictable behavior, but ordered enough for that behavior to persist and interact. Conway's Game of Life is the canonical example — the same cellular automaton I had programmed on a 286 as a kid. Three dead-simple rules on a flat grid, yet they produce gliders, oscillators, self-replicating structures, and — provably — universal computation. You can build a computer inside it. You can build a computer inside that computer. In principle, you can run an entire three-dimensional virtual world inside a two-dimensional grid of pixels. From almost nothing, everything.

This is where consciousness lives. Only edge-of-chaos dynamics have both properties you need: **universal computation** (complex enough to actually run a self-simulation) and **global integration** (distant parts of the system influence each other, local changes propagate globally, information is bound into a unified whole). This is why conscious experience feels *unified* — you don't see red over here and hear a voice over there as separate streams. The critical dynamics bind everything into one experience. Binding

isn’t something the brain does *in addition to* its other computations; it’s a consequence of the dynamical regime.

A brain in deep sleep, running slow delta waves, is operating in periodic dynamics: repetitive, going nowhere. The models are still there in the substrate, but the simulation isn’t running. A brain in generalized seizure is pushed into chaotic dynamics: the simulation can’t hold together. Only in the waking state — poised at the edge of chaos — does the system sustain conscious experience.

The brain, as a universal computer optimized by billions of years of evolution, uses *all* the computational regimes as distinct tools: stable attractors for long-term memory, periodic oscillations for timing and gating (alpha, theta, gamma rhythms), fractal processing for scale-invariant recognition and texture analysis (primarily in V2-V4 of the visual cortex), and edge-of-chaos dynamics for the cortical automaton itself — the engine of consciousness. Only the edge-of-chaos regime generates consciousness. But consciousness depends on the others to function.

When I published this argument in my 2015 book, I had no idea that empirical neuroscience was independently heading toward the same conclusion.

But there’s a crucial subtlety. Criticality alone is not enough. A pot of boiling water can exhibit complex dynamics at the edge of chaos. It is not conscious. The theory requires *two* thresholds to be met: the physical one (the substrate must operate at criticality) and the functional one (the substrate must implement the four-model architecture). Criticality without the architecture gives you complex dynamics but no consciousness. The architecture without criticality gives you a dormant system — the models exist in the substrate

but the simulation isn't running. Both thresholds must be met. Together, they are sufficient.

The Cortical Automaton

Now I want to make something concrete that might still feel abstract. I've been talking about the cortex needing to operate at the edge of chaos, in Class 4 dynamics. But what *is* the Class 4 system? It's not some mysterious force hovering above the brain. It's the pattern of neural firing itself.

Think about what the cortex actually looks like in operation. Billions of neurons, each one either firing or not, each one influencing its neighbors through learned connection weights. Each neuron is a cell in a cellular automaton — not metaphorically, but literally. The rules of the automaton are the synaptic weights, the thresholds, the local wiring. The output of each “cell” is a firing rate. And the result, the grand pattern of electrical activity dancing across the cortical surface at 10 to 40 Hz, is a Wolfram Class 4 cellular automaton operating in a space of many thousand dimensions.

I call this the **cortical automaton**.

It's the same idea I programmed on a 286 as a kid — Conway's Game of Life — except instead of a flat grid with three rules, it's a folded sheet of cortex with billions of locally varying rules, and instead of moving in two dimensions, its patterns move through a dimensional space so vast that it defies visualization. Like an octopus with limitless arms, the cortical automaton can reach any part of the cortex at any time, activating whatever stored models it needs — a memory here, a motor plan there, a fragment of language

somewhere else. It grabs these models like little Lego figures and uses them to navigate from one satisfying state to the next.

And here’s the critical distinction: **the cortical automaton is not consciousness**. It’s the engine, not the experience. The seemingly chaotic pattern of billions of neurons firing is, in reality, an extraordinarily sophisticated apparatus that computes, thinks, and steers a body through a life. But consciousness is only one *effect* of this apparatus — an effect that arises from the interplay between the automaton and the cortex when the conditions are right. When the automaton synchronously sweeps across suitable cortical regions at the right frequency in a coherent temporal sequence, a conscious experience emerges from that sequence of frames. The automaton contains the instances of our world model and our self-model; consciousness is what happens when these models are actively running in the simulation.

You can, by the way, observe the cortical automaton directly — no fMRI required.

Here’s how: Find a completely dark room. Close your eyes. Wait for any afterimages to fade — this takes about 30 to 60 seconds if you’ve been looking at anything bright. At first you see nothing, or almost nothing. But then, if you wait and pay attention, you’ll start seeing flickering colored points against the darkness.

Most people dismiss these as “retinal noise” — random firings in the photoreceptor cells of the eye responding to pressure or spontaneous chemical events. And if you press gently on your eyelid, you can indeed trigger localized visual sensations that way. But the colored points you see in total darkness are *not* retinal. They’re too organized for that. What you’re seeing is the resting activity

of V1 — your primary visual cortex — driven by a combination of residual sensory signals and top-down projections from the cortical automaton itself. The automaton is running its baseline dynamics, and you’re watching it happen in real time.

If you keep watching — not concentrating, but relaxing, letting your attention soften — something remarkable happens. Active focus actually suppresses these patterns; it’s when you stop trying to see that you start seeing. The automaton starts recruiting more of the visual system to interpret and amplify what little signal is there. The flickering points stabilize into shapes. Geometric patterns emerge: grids, spirals, lattices. Then faces, distorted and shifting. Then figures. Then, with enough patience (and I mean *hours*, not minutes), full scenes — elaborate, colored, narrative hallucinations no different in kind from the dreams you have every night.

This is the same mechanism behind hypnagogic hallucinations — the vivid imagery that flickers through your mind just as you’re falling asleep. It’s the cortical automaton running with minimal external constraint, generating its own content by activating stored patterns and projecting them into the simulation. The progression you experience — from faint noise to coherent hallucinations — is a direct window into how the automaton works: it starts with V1, the earliest visual processing stage, and progressively recruits V2, V3, and higher areas as it tries to make sense of whatever signal is available. When no real signal is available, it *generates* one. This is the permeability leak in action. With no external signal to dominate the simulation, the substrate’s own processing noise becomes visible. You’re not hallucinating *nothing* — you’re seeing the graphics engine’s idle patterns, the neural equivalent of static

on an untuned TV. Except this static has structure, because the processing machinery has structure.

You can also induce a temporary form of synesthesia this way. In my youth, I used this to “see music.” If you close your eyes and listen to music while letting the visual patterns come to you — relaxed, passive, not straining to see — the patterns gradually synchronize with the rhythm and frequencies of what you’re hearing. The cortical automaton, deprived of external visual input, starts coupling its visual dynamics to whatever other strong signal is available — in this case, auditory input. What you see is, quite literally, your brain’s activity made visible: the automaton’s V1-level patterns being driven by auditory cortex rather than retinal input. Real synesthetes — people whose senses are permanently cross-wired, who always see colors when they hear sounds — may have a more permanent version of this same coupling, likely due to stronger or more numerous connections between sensory areas, whether in the thalamus or the cortex itself. The mechanism is the same: one sensory modality leaking into another’s processing pipeline. The cortical automaton doesn’t much care where its input comes from. It processes whatever it receives.

I’m not recommending you try this as a regular hobby. The experience can be unsettling, especially if you’re not psychologically prepared for it. And there’s an outside chance that sustained sensory deprivation could destabilize someone with latent psychiatric vulnerabilities. But if you’ve ever wondered what the substrate of your consciousness looks like when it’s idling — when the external world has gone quiet and the system is just... running — this is the most direct glimpse you can get without a brain scanner.

That progression from almost-nothing to a complete fictional visual world, experienced by your self-model in a virtual universe, is a direct portrait of the cortical automaton at work.

When the automaton goes wrong, you can see that too. An epileptic seizure is what happens when parts of the automaton fall into Class 1 or 2 dynamics — periodic, locked, computationally useless — or are pushed past Class 4 into Class 5 chaos. A stroke is what happens when parts of the cortex drop out entirely. A fainting spell is what happens when the minimum frequency for wakefulness is no longer met. The automaton is somewhat fragile. But the structure that generates it — the neocortex, with its learned weights and evolved architecture — is robust, which is why we can recover from these disruptions so remarkably well.

The Convergence

In 2003 — two years before I even had the theory — John Beggs and Dietmar Plenz discovered “neuronal avalanches” in cortical tissue: patterns of neural activity that followed the mathematical signature of self-organized criticality, a hallmark of systems at the edge of chaos.

In 2014, Robin Carhart-Harris proposed the Entropic Brain Hypothesis: the idea that the level of consciousness correlates with the entropy (disorder) of brain activity, with the sweet spot at an intermediate level — too little entropy means unconsciousness, too much means incoherent experience.

In 2016, Enzo Tagliazucchi and colleagues showed that LSD pushes the brain toward criticality, consistent with the enhanced

(but sometimes chaotic) consciousness that psychedelic users report. By 2022, a review paper could already speak of “self-organized criticality as a framework for consciousness” — the evidence was building.

And in 2025-2026, the empirical dam broke. Keith Hengen and Woodrow Shew published a meta-analysis of 140 datasets in *Neuron* (2025) — the largest systematic analysis of criticality in brain dynamics ever conducted — confirming that the brain operates near a critical point across multiple measurement modalities. Then Inbal Algom and Oren Shriki proposed the ConCrit framework — Consciousness and Criticality — in *Neuroscience & Biobehavioral Reviews* (2026), arguing that critical brain dynamics provide a unifying mechanistic foundation for all major theories of consciousness. Their conclusion: consciousness tracks criticality. When the brain is at or near the critical point, consciousness is present. When it’s pushed below criticality (by anesthesia, by sleep, by brain damage), consciousness is absent. When it’s pushed past criticality (by seizure, possibly by some drug states), consciousness becomes incoherent.

Two paths. One theoretical, starting from Wolfram’s computational framework and reasoning about what a self-simulation requires. One empirical, starting from neural recordings and analyzing statistical properties of brain activity across every accessible state of consciousness. Two decades apart in origin, converging on the same conclusion.

This is the kind of convergence that makes you take a theory seriously.

Three Ways a Hologram Meets an Automaton

While writing this chapter, I realized something that stopped me cold.

The holographic principle and Class 4 automata keep showing up in the same conversations — in physics, in neuroscience, in computation theory. But nobody seems to have asked the obvious question: *what are the possible relationships between them?*

There are exactly three.

Relationship 1: A holographic substrate produces Class 4 dynamics. This is probably what the brain does. Neural networks are locally holographic — Karl Lashley showed decades ago that you can destroy large portions of cortex and the memories persist, degraded but complete, just like cutting a hologram in half gives you the whole image at lower resolution. And that holographic substrate, operating at criticality, produces the Class 4 dynamics that consciousness requires. Well-established, thoroughly documented, and — forgive me — the boring one.

Relationship 2: A Class 4 automaton that produces holographic patterns as emergent behavior. The automaton isn't holographic in its rules, but its dynamics spontaneously generate holographic structures — higher-dimensional information encoded in lower-dimensional patterns, arising from the computation itself. If a Class 4 automaton naturally produces holographic output, that means non-local information distribution emerges from purely local rules — which is, intriguingly, exactly what quantum entanglement looks like.

This is where I should mention Gerard 't Hooft, because the connection is too striking to skip — even though it's speculative.

't Hooft, a Nobel laureate in physics, has proposed that quantum mechanics itself is a cellular automaton at the Planck scale: that our universe is fundamentally deterministic, and quantum effects are emergent phenomena of a deeper, discrete dynamics. If he's right, the principle I've been describing doesn't just apply to consciousness by analogy — it's literally how the universe works, all the way down. Simple local rules produce a holographic universe, and within that universe, simple neural rules produce a holographic consciousness. The same computational principle operating at two scales: cosmological and neurological. I find this fractal consistency deeply compelling, but I want to be honest: 't Hooft's interpretation remains a minority view in physics, and the argument from structural elegance to physical reality has been rightly criticized. Still — if a single computational principle turns out to underlie both the universe and the minds that model it, that would be the most beautiful fact ever discovered.

Relationship 3: A Class 4 automaton whose rule structure is itself holographic. This is the one that made me put down my pen. If such a thing exists — a cellular automaton where the rules themselves encode higher-dimensional information in a lower-dimensional structure, the way a hologram encodes three dimensions in two — then you would have a system that naturally does what the holographic principle says the universe does. Not a system that merely *runs on* a holographic substrate (or produces a hologram). A system that *is* a holographic encoding. Also possibly the universe — though I should note this is speculative, and the argument that mathematical beauty implies physical reality has

been legitimately criticized. I'll plant the seed here and leave it for future work.

I'll return to this in Chapter 14, where I'll explain why I think Relationship 3 might be the most important unsolved question in mathematics.

Chapter 6

What Psychedelics Reveal

A necessary note before we begin: nothing in this chapter should be read as a recommendation to try psychedelics. They are powerful, unpredictable, and can ruin your life — literally, permanently. They can trigger schizophrenia in those with a predisposition. They can cause psychotic episodes, persistent anxiety disorders, and HPPD (hallucinogen persisting perception disorder) that never goes away. I discuss them here because they reveal something important about the architecture of consciousness. That scientific value does not make them safe.

If you want to understand consciousness, study what happens when it goes wrong. Psychedelics are, I believe, the most illuminating window into the architecture of consciousness that we possess — more revealing than brain scans of sleeping patients, more theoretically informative than lesion studies, and dramatically more accessible than split-brain surgery.

Here's why: psychedelics don't just *change* consciousness. They change it in *systematic, predictable ways* that reveal the underlying architecture — if you know what to look for.

The Permeability Gradient

Remember the boundary between the implicit models and the explicit models — between the stored knowledge (real side) and the running simulation (virtual side). In normal waking life, this boundary is selectively permeable: relevant information gets through, irrelevant information stays in the library. You're conscious of what you need, and unconscious of everything else.

Psychedelics blow the boundary open.

Under psychedelics — LSD, psilocybin, DMT, mescaline — the permeability of the implicit-explicit boundary increases globally. Information that is normally processed entirely on the real side, invisible to consciousness, starts leaking through to the simulation.

And here's the crucial point: it leaks through *in order*.

At low doses or early in the experience, the simplest processing stages become visible first. These are the stages closest to raw sensory input: V1-level processing. You see enhanced colors, breathing patterns in static surfaces, subtle movements in peripheral vision. These are the visual cortex's early feature detectors, normally invisible, now entering the simulation.

As the dose increases or the experience deepens, more complex processing stages become visible. V2/V3-level processing: geometric patterns, fractals, tessellations, the famous "form constants" that Heinrich Klüver catalogued in the 1920s. These are the visual system's intermediate representations — the building blocks it normally uses to construct your visual experience, now visible in their own right.

Higher still, and the higher visual areas become accessible. Faces appear. Figures. Scenes. The face-processing areas, the object-

recognition areas, the scene-construction areas — all normally operating below the threshold of consciousness — now broadcasting their intermediate products directly to the simulation.

At the highest doses, the entire processing hierarchy is exposed, and the result is full-blown visionary experience: complex, narrative, dreamlike scenes constructed from the deepest layers of implicit processing.

This ordered progression — simple to complex, V1 to higher areas, dose-dependent — is exactly what the Four-Model Theory predicts. It’s a direct consequence of the permeability gradient: lower-level processing stages, being closer to the boundary, become accessible before higher-level ones as permeability increases.

Here is the visual processing hierarchy, showing what each area does normally and what becomes visible when the permeability barrier drops:

Area	Normal function	Psychedelic signature
V1	Edges, spatial frequency, orientation	Phosphenes, Klüver form constants, breathing surfaces
V2	Contour integration, texture, border ownership	Tessellations, repeating geometric patterns
V3	Global form, dynamic shape processing	Flowing, morphing geometries
V4	Color, curvature, complex texture	Colored fractals, kaleidoscopic patterns
V5/MT	Motion processing	Rotation and movement of patterns
Fusiform/IT	Faces, objects, word forms	Faces, figures, entities
Anterior IT	Semantic categories, scene construction	Full narrative hallucinations

Each row represents a deeper stage of processing. Under normal conditions, you experience only the final output — the finished percept. Under psychedelics, you experience the *intermediate* stages, in order, as permeability increases. (A fuller version of this table, with receptive field sizes and additional detail, is in Appendix A.)

I know this sounds intriguing. You’re reading about layers of visual processing becoming visible, and part of you is curious what that looks like. I understand — I was curious too. I tried both paths. I was young, and stupid, and lucky. The meditation route, which I described in the previous chapter — a dark room, relaxed attention, patience — gets you to the same place. Not as fast, not as dramatic on the first try. But just as impressive, just as real, and without the risk of permanently damaging your mind. A warm bed in a dark room is all you need.

And there’s another route: lucid dreaming. If you can learn to recognize that you’re dreaming while you’re still in the dream — and this is a trainable skill — you get access to the full simulation running unconstrained. No sensory input, no external reality to correct the model. Just the virtual world, with you consciously inside it. For some people, this is easier to achieve than sustained meditation. The techniques are well-documented (see Appendix D for a practical guide), and the experience can be at least as revelatory as anything a drug produces — without the risk. We’ll return to lucid dreaming in Chapter 7.

And this is where the five-level hierarchy from Chapter 2 does its explanatory work. Remember the five nested systems — Physical, Electrochemical, Proteomic, Topological, Virtual? Psychedelics target the middle of the stack and the effects ripple upward. Classic psychedelics like LSD and psilocybin bind to serotonin 2A receptors, acting at the **electrochemical** level — they change how neurons talk to each other. That perturbation propagates to the **proteomic** level, where receptor sensitivity shifts over hours. It reshapes the **topological** level, where network connectivity patterns change —

visible on fMRI as increased global integration. And it transforms the **virtual** level, where the conscious simulation floods with content that is normally invisible. The only level classic psychedelics don't touch is the **physical** — they don't destroy neurons, don't alter the raw matter. They change everything *above* the matter, in ascending order. This is a crucial distinction. Classic psychedelics — LSD, psilocybin, DMT, mescaline — are not neurotoxic. They change how neurons communicate without destroying them. Many other drugs are not so kind. Cocaine, methamphetamine, and alcohol physically destroy neurons. MDMA at high or repeated doses damages serotonin axons. Even Amanita muscaria — the iconic red-and-white mushroom that many people confuse with psychedelic mushrooms — is a deliriant that works through an entirely different, more dangerous mechanism. If you take nothing else from this chapter: not all drugs that alter consciousness are alike, and the distinction between “changes the signal” and “destroys the hardware” is literally the difference between a temporary altered state and permanent brain damage. The dose-dependent visual progression maps directly onto this: low doses perturb the electrochemical level enough to affect V1 processing; higher doses propagate the perturbation up through more levels, recruiting increasingly complex processing stages into conscious experience.

The Redirectable Self

But the most dramatic evidence comes from what happens to the self.

Your Explicit Self Model — the “I” — is a virtual process that requires input. Under normal conditions, it receives a steady stream of self-referential signals: your sense of where your body is (proprioception), your sense of how your organs feel (interoception), the narrative stream of inner speech, and the constant background of bodily self-awareness that you never notice until it’s disrupted.

At high psychedelic doses, this input gets disrupted. The self-model doesn’t die — it *redirects*. Deprived of its normal self-referential input, it grabs whatever input is dominant.

This is most dramatically demonstrated by salvia divinorum, a dissociative psychedelic that acts on kappa-opioid receptors (completely different from the serotonergic mechanisms of LSD or psilocybin). Salvia users consistently report experiences of *becoming* things:

- “I became the couch.”
- “I was the wall.”
- “I turned into a page in a book.”
- “I was one of the characters on the TV.”
- “I became a fractal — not seeing a fractal, *being* a fractal.”

These are not metaphors. Users report complete, experientially convincing identity shifts. For the duration of the experience, they *are* the object or entity in question. Some describe it as feeling like being dead — not dying, but *being dead* — because when you are a chair, the person you were has simply ceased to exist.

And the content tracks the sensory environment. The person watching TV becomes a TV character. The person lying on a couch

becomes the couch. The person looking at a pattern becomes the pattern.

This is the Explicit Self Model doing exactly what the theory predicts: redirecting to whatever input dominates when normal self-input is disrupted. The identity content isn't random — it's determined by the sensory environment. Control the environment, and you should be able to control the identity experience.

I need to pause the theory here for a moment. Salvia divinorum is, as far as we know, the strongest psychedelic substance on Earth. The complete proprioceptive takeover I just described means total loss of body awareness and spatial orientation. People under salvia's influence have walked out of tenth-floor windows. They have stepped into traffic. They have died. This is not a party drug, not a curiosity to try on a Friday night. It is the most extreme pharmacological disruption of the Explicit Self Model that exists, and that disruption can kill you — not because the drug is toxic, but because you stop knowing where your body is and may fully believe you have wings and can fly.

Many people who try salvia report that the experience felt like dying — not metaphorically, but as a genuine, terrifying conviction that they had ceased to exist. This is the Explicit Self Model collapsing so completely that the simulation can no longer generate a "you" at all. We'll see the clinical equivalent of this in Chapter 8, when we discuss Cotard's delusion — patients who are neurologically convinced they are dead. Salvia gets you there pharmacologically, in seconds, without warning. Think about whether that's something you want to experience.

I experienced the time dilation myself. Under salvia, half a second of real time — confirmed by the person watching me — stretched into what felt like fifteen minutes or more. My perceptual world rebuilt itself into elaborate sequences that included the sensation of having wings and flying around (the flying feeling, I later realized, came from air rushing past as I fell backward onto the bed). My entire reality collapsed and regenerated, all in the time it takes to blink. I described this to an observer who was timing me, and they said I’d been “gone” for less than a second. The same kind of time dilation I would experience a few years later, in 1998 or 1999, during a near-death event — a mechanism I’ll describe in Chapter 14 — but pharmacologically induced and even more extreme.

I’m not the most dramatic case. One well-documented account involves a man who experienced what felt like eight complete years of an alternative life — attending school, making friends, building a new existence — during a salvia episode that lasted roughly forty-five seconds of clock time. Peer-reviewed research confirms extreme temporal distortion under controlled conditions, with one participant describing time as “creased like an accordion” (Addy et al., 2015). The substrate runs so much content through the simulation so fast that subjective time decouples entirely from clock time.

This has never been experimentally tested in a controlled setting. But it could be — and it would be a dramatic confirmation of the theory’s most distinctive mechanism.

If you want to see how far this principle extends, consider the following thought experiment. Imagine someone permanently

maintained on a very high (but not completely dissociating) dose of Salvinorin A — the active compound in salvia divinorum, which acts on a single receptor type (kappa-opioid). This person's Explicit Self Model would never stabilize. It would cycle endlessly through whatever input happened to dominate: one moment they'd believe they were a chair, then a table, then a dinosaur, then air, then a piece of paper. They would still *experience* things — vision and hearing would still function — but they would never again know who or what they were. Remove the drug, and over time, the normal self-model would reassemble from the intact Implicit Self Model.

This is important because it shows that consciousness doesn't require a *correct* self-model. It just requires *a* self-model. The architecture keeps running regardless. The Explicit Self Model doesn't shut down when it's given absurd input — it builds the best self it can from whatever signals are available. This is the same principle we see in Cotard's delusion (the ESM on absent interoceptive signals: "I must be dead"), in Anton's syndrome (the ESM generating vision from memory when the eyes aren't working), and in conversion disorder (the ESM modeling paralysis that the substrate doesn't actually have). The self-model is a compulsive constructor. It never stops building. It never announces that the data are insufficient. It just builds, and believes.

Anosognosia: The Inverse

There's a beautiful symmetry here. If psychedelics are what happens when the implicit-explicit boundary becomes *too* permeable,

anosognosia is what happens when it becomes *too* impermeable — at least locally.

Anosognosia, most commonly seen after right-hemisphere stroke, is the condition in which patients are genuinely unaware of their own deficits. A patient with a paralyzed left arm will insist the arm is fine, will attempt to explain away failures to use it, and will become confused or angry when confronted with evidence of the paralysis. They’re not in denial in the psychological sense — the information that the arm is paralyzed simply never reaches their conscious simulation.

In the Four-Model Theory, this is a local decrease in implicit-explicit permeability. The Implicit Self Model *has* the paralysis information — the substrate registers the damage. But the boundary is blocked for that specific domain, so the Explicit World Model never includes the deficit. The patient’s simulation doesn’t contain a paralyzed arm, so the patient doesn’t experience one.

The mechanism is more specific than that, and once you see it, it’s elegant in a slightly horrifying way. When your motor system sends a command — say, “clap your hands” — it simultaneously does two things. It sends the command to the muscles, and it sends *predicted feedback* to consciousness: what clapping should feel and sound like, based on past experience. This predicted feedback arrives *before* the actual sensory feedback, because the real feedback has to travel through slower neural pathways. Under normal circumstances, the prediction is quickly corrected or confirmed by the actual sensory data. You predict the clap, then you feel and hear the clap. Match. Move on.

But in anosognosia, the actual feedback from the paralyzed limb never arrives. And the mechanism that should flag “wait — nothing happened” is damaged. So the predicted feedback goes uncorrected. The patient’s motor system commands both hands to clap, sends the prediction of a two-handed clap to consciousness, and consciousness experiences exactly that — a perfectly normal clap with both hands. The patient will tell you, with complete sincerity, that they just clapped with both hands. They heard it. They felt it. They experienced it. In their simulation, it happened. It just didn’t happen in reality.

This is not a metaphor for how consciousness works. This *is* how consciousness works, all the time, in all of us. The only difference is that in healthy people, the predicted feedback gets corrected within milliseconds. In anosognosia, the correction mechanism is broken — and the patient’s simulation simply runs on predictions alone.

Psychedelics and anosognosia are the same mechanism running in opposite directions. One increases permeability globally. The other decreases it locally. And this symmetry generates a cross-domain prediction: psychedelics should alleviate anosognosia. The global permeability increase should overwhelm the local block, allowing the deficit information to reach consciousness.

No one has ever tested this, because no one has had a theory that connects these two phenomena. The connection is invisible without the Four-Model Theory.

Chapter 7

What Happens When the Lights Go Out

Every night, you lose consciousness. Every morning, you get it back. And the transition between the two — the journey through sleep stages — is a nightly demonstration of the criticality principle.

Deep Sleep: Below the Threshold

In deep non-REM sleep, the brain's dynamics shift to a subcritical regime. The hallmark is slow waves: large, synchronized oscillations in which vast populations of neurons fire in unison and then fall silent together. This is Class 2 dynamics — periodic, repetitive, too ordered for consciousness.

The Perturbational Complexity Index (PCI), developed by Marcello Massimini and colleagues, confirms this directly. PCI measures how complexly the brain responds to a magnetic pulse: in waking consciousness, the response is complex and differentiated (high PCI); in deep sleep, it's simple and stereotyped (low PCI).

The brain in deep sleep cannot sustain the rich, globally integrated dynamics that a conscious simulation requires.

The lights are off. The Explicit World Model and Explicit Self Model have collapsed. There is no simulation and no experience.

Dreams: Degraded Mode

But the lights come back on during REM sleep. The brain's dynamics shift back toward criticality — not fully, but close enough. The simulation re-engages, and you experience a world again.

But it's a degraded simulation. The normal external input is cut off (your eyes are closed, your muscles are paralyzed). The Explicit World Model runs on internal data — drawing from the Implicit World Model's stored knowledge rather than from current sensory input. This is why dreams feature familiar places and people but with impossible physics and narrative incoherence: the simulation is doing its best with limited input.

The Explicit Self Model also runs in degraded mode. You experience dreams as happening to "you," but your metacognitive oversight is reduced — you accept impossible events without question, you rarely notice that you're dreaming, your critical faculties are dimmed.

Sleepwalking is an even more dramatic demonstration. In sleepwalking, the motor system partially reactivates while the Explicit Self Model remains offline or nearly so. The substrate is running motor programs — walking, navigating, even performing complex actions — but the simulation isn't fully engaged. The walker moves

through the physical world guided by the Implicit World Model’s spatial knowledge, but with minimal or no conscious experience.

I know this firsthand. As a teenager, I went through a phase of sleepwalking. One morning I woke to find myself at my desk, with scribbled notes in front of me — written left-handed, which I never do while awake. I had a fragmentary memory of walking along the walls in a circle, trying to find the door, not finding it. But the part where I sat down at the desk and tried to write — that was completely dark. The substrate was navigating, motor programs were executing, but the simulation — the “I” — wasn’t there.

This is the theory in miniature. A body moving through the world, processing spatial information, executing learned motor programs, all without a conscious self inside the loop. The implicit models run the show. The explicit models are offline. And the result is a human being who walks, acts, and even writes — but is nobody home.

Lucid Dreaming: The Switch

And then there’s lucid dreaming — the state in which you realize you’re dreaming while still inside the dream. In the Four-Model Theory, this is the Explicit Self Model “toggling on” more fully within the dream state. It’s a step-like increase in self-modeling capacity.

The theory predicts that this transition — from non-lucid to lucid dreaming — corresponds to a criticality threshold crossing. Not a gradual increase in brain complexity, but a sudden step. If you measured EEG complexity in a time-locked window around

the moment of lucidity onset (using the established paradigm of pre-agreed eye-movement signals from lucid dreamers), you should see a discontinuity.

Anesthesia: The Two Types

Anesthesia provides the cleanest test of the criticality principle, because different anesthetic agents produce dramatically different experiences despite being classified under the same label.

Propofol pushes the brain subcritical. Thalamocortical connectivity is disrupted, cortical complexity collapses, and PCI approaches zero. The lights go out completely. Patients report no experience during propofol anesthesia. This is exactly what the theory predicts: push below criticality and the simulation cannot be sustained.

Ketamine does something completely different. It does *not* push the brain subcritical. EEG studies show that ketamine *increases* neural entropy — it pushes the brain toward or past criticality, into a more chaotic regime. The result? The “K-hole” — vivid, often bizarre experiences of dissociation, distorted reality, out-of-body experiences, and radical identity alteration.

In the Four-Model Theory, the K-hole is consciousness running on *wrong* input. The Explicit World Model and Explicit Self Model are still active (the brain is still at or above criticality), but external sensory processing is disrupted. The simulation runs on internal and distorted signals, producing the characteristic K-hole phenomenology.

This distinction — propofol abolishes consciousness by going subcritical, ketamine alters consciousness by going supracritical with disrupted input — is a genuine explanatory advantage. Most theories struggle to explain why two “anesthetics” produce such radically different experiences. The criticality framework makes the distinction natural.

The Consciousness Map

State	Criticality	Models	Consciousness
Normal waking	At critical	All four active	Full
REM sleep	Near-critical	EWM/ESM on internal input	Degraded (dream)
Deep NREM	Subcritical	EWM/ESM collapsed	Absent
Propofol	Forced subcritical	EWM/ESM suppressed	Absent
Ketamine	Past critical (\uparrow entropy)	EWM/ESM on wrong input	Present, disconnected
Psychedelics	At/past critical	All active, \uparrow permeability	Present, altered
Lucid dreaming	Near-critical, threshold crossed	EWM active, ESM fully engaged	Enhanced self-awareness

This table summarizes everything we’ve covered in this chapter — and provides a reference you can come back to. Every state of consciousness you’ve ever experienced fits somewhere on this map, determined by two factors: whether your substrate is at critical-

ity, and which of the four models are running. Sleep, anesthesia, psychedelics, dreams, the K-hole — they're not separate mysteries. They're different coordinates on the same map.

Chapter 8

The Clinical Mirror

The same four-model architecture that explains sleep and anesthesia also explains some of the most dramatic and puzzling conditions in clinical neurology. These aren't just interesting case studies — they're what happens when specific components of the architecture fail. And each failure illuminates the architecture from a different angle, the way a blown fuse tells you which circuit it was protecting.

If the theory is good, then damage to specific models should produce specific, predictable deficits. Not vague "consciousness is impaired" hand-waving, but precise predictions: knock out this component, and you get *that* syndrome. Keep a different component running without its normal input, and you get *this* other syndrome. The clinical literature is full of conditions that are deeply puzzling under standard models of consciousness — but fall into place naturally when you have a real/virtual distinction and four interacting models to work with.

Blindsight and Anton's syndrome: The perfect mirror

If you remember only one thing from this chapter, remember this pair. Every other theory of consciousness struggles to explain

even one of these conditions. The Four-Model Theory predicts both.

Start with blindsight. A patient has damage to primary visual cortex — the part of the brain that generates conscious visual experience. By any standard clinical test, the patient is blind. Ask him what he sees, and he'll tell you: nothing. He means it. He's not being modest or confused. As far as his conscious experience goes, the visual world simply doesn't exist.

But then something astonishing happens. Researchers place obstacles in a hallway and ask the patient to walk through it. He protests — he can't see, how could he possibly navigate? They insist. He sighs, stands up, and walks.

And he navigates the obstacle course flawlessly. Steps around chairs. Ducks under a barrier that wasn't there last time. Weaves through a gap between two obstacles — all while insisting, truthfully and sincerely, that he cannot see a thing. There is video of this — I encourage you to find it, because reading about it doesn't do it justice. The footage of a clinically blind man weaving through an obstacle course like he can see perfectly is one of the most stunning demonstrations in all of neuroscience. The researchers watching look like they've seen a ghost.

How? Because the substrate still processes visual information. The Implicit World Model receives visual input through subcortical pathways that bypass the damaged cortex — a fast route from retina to superior colliculus to pulvinar that evolved long before the cortex existed. It builds a spatial map, guides motor behavior, keeps the body from colliding with objects. But none of this reaches the Explicit World Model. The conscious simulation contains no vi-

sion. The patient genuinely experiences blindness — and genuinely navigates by sight. The substrate works without the simulation.

Now flip it. Anton’s syndrome — anosognosia for cortical blindness — is the exact inverse. These patients are genuinely, completely blind. Their visual cortex or optic pathways are destroyed. No visual information reaches the brain at all. But they are absolutely, unshakably convinced they can see.

They walk into walls and blame the furniture for being in the wrong place. They describe objects that aren’t in the room with complete confidence — “There’s a blue vase on the table” — when the table is empty. Ask them to identify what you’re holding up and they’ll give you an answer, calmly and without hesitation, and it will be wrong. Confront them with evidence of their blindness and they become confused, then irritated, then angry. The lighting is bad. They need new glasses. They just weren’t paying attention. They are not lying. They are not in denial in the psychological sense. They genuinely, experientially see — and what they see has no correspondence to the actual world.

In the Four-Model Theory, this is the Explicit World Model generating a visual simulation from the Implicit World Model’s stored knowledge — even though no current visual input is arriving. The simulation runs on old data, on expectations, on the brain’s best guess about what the world should look like. The patient “sees” a world that isn’t there. The simulation runs without current input.

Put them side by side. Blindsight: the substrate processes vision, but the simulation doesn’t show it. Anton’s syndrome: the simulation shows vision, but the substrate isn’t receiving it. Substrate without simulation. Simulation without input. Both conditions

are deeply puzzling if you think consciousness is a single, unified thing. Both are natural, even predictable, consequences of a theory that distinguishes between real processing and virtual experience. You almost couldn't design a better pair of test cases if you tried.

Covert awareness: Trapped inside

In 2006, Adrian Owen and his colleagues published a study that changed how we think about the vegetative state. They placed a patient who had been diagnosed as vegetative — unresponsive, apparently unconscious — into an fMRI scanner and asked her to imagine playing tennis. Her brain lit up in exactly the same pattern as a healthy conscious person imagining the same thing.

She was in there. Conscious, aware, thinking — and completely unable to move, speak, or signal her presence to anyone.

The Four-Model Theory makes a clean distinction here. A truly vegetative patient has a subcritical substrate. The dynamics have fallen below the threshold. The simulation isn't running. There's nobody home — not because the person has "left," but because the computational architecture that generates the simulation has gone offline.

But a covertly conscious patient is something entirely different. The substrate is critical — the dynamics are rich enough to sustain a simulation. The Explicit World Model and Explicit Self Model are running. The person is experiencing, thinking, feeling. But the output pathways are destroyed. The simulation has no way to express itself. The person is conscious but locked in, trapped inside a body that won't respond.

The Perturbational Complexity Index — the same measure that distinguishes sleep stages — should distinguish these cases. And

it does. Some patients diagnosed as vegetative show PCI values squarely in the conscious range. They’re not vegetative at all. They’re prisoners. The medical and ethical implications are enormous, and the Four-Model Theory tells you exactly why the distinction exists and exactly how to detect it.

Cotard’s delusion: “I am dead”

And then there are patients who believe they are dead.

Cotard’s delusion is one of the strangest conditions in psychiatry. Patients insist they have died. They believe their organs have dissolved, their blood has drained away, they no longer exist. Some believe they are rotting. Some believe they are immortal — because if you’re already dead, you can’t die again. They are not speaking metaphorically. They mean it with complete, unshakable conviction.

By now, you should recognize the mechanism. It’s the same one from Chapter 6 — the Explicit Self Model constructing the best model it can from whatever input is available. In Cotard’s, the interoceptive input is severely distorted. The internal body signals that tell you your heart is beating, your stomach is digesting, your lungs are breathing — they’re absent or garbled. And the ESM, ever the compulsive constructor, interprets “no heartbeat, no digestion, no breathing, no body sensation” the only way it can: I am dead.

Salvia’s “I am a chair.” Anosognosia’s “my arm is fine.” Split-brain confabulation’s “I picked the shovel to clean the chicken shed.” And now Cotard’s “I am dead.” One mechanism running through every case. The Explicit Self Model is always doing its job — always building the best self-model it can. When the input is right, you feel like yourself. When the input is wrong, you feel like a chair,

or fine when you're paralyzed, or dead when you're alive. But it always feels completely, convincingly real — because it's the only self you have access to.

Alien Hand Syndrome: When the committee disagrees

And then there's a condition that reads like a horror movie but illustrates the multi-agent nature of the substrate more vividly than any thought experiment. In Alien Hand Syndrome, one of the patient's hands acts with apparent purpose and intention — but against the patient's conscious will. One hand lights a cigarette while the other hand takes it away and throws it on the ground. One hand reaches for a doorknob while the other grabs the wrist and pulls it back. The patient watches, horrified, as part of their own body pursues goals they did not choose.

Stanley Kubrick used this in *Dr. Strangelove* — and people assumed he'd invented it. He didn't. The syndrome is real, and it comes in two varieties. In the callosal form, caused by damage to the corpus callosum, the symptoms resemble split-brain conflict: two hemispheres with competing motor plans, neither able to override the other. In the frontal form, caused by prefrontal damage, the "alien" hand exhibits disinhibited behavior — grabbing objects, using tools, touching things compulsively, all seemingly with purpose but without the patient's consent.

There's also a subtler variant called Anarchic Hand Syndrome, where the patient lacks motor *control* rather than motor *ownership*. The hand does things the patient didn't intend, but the patient still recognizes it as *their* hand — they just can't stop it. The distinction matters: Alien Hand is a failure of the Explicit Self Model's body ownership boundary ("that hand isn't mine"), while Anarchic

Hand is a failure of the motor inhibition system (“that hand is mine but it won’t listen”). Same architecture, different failure points.

The key insight from the German book’s analysis of these syndromes is that your sense of authorship — the feeling of “I did that” — is not computed before or during the action. It’s computed *after*, by comparing the action’s predicted outcome with the observed outcome. When the comparison matches, you feel ownership. When it doesn’t, you don’t. This is why patients with Alien Hand Syndrome can sometimes tickle themselves — their prediction system isn’t generating the expected outcome for the alien hand’s movements, so the touch arrives as unexpected, as if from someone else.

Charles Bonnet Syndrome: The simulation that won’t stop

If you want more evidence that the brain’s simulation is *generative* — that it constructs experience from models rather than passively receiving it from the senses — consider Charles Bonnet Syndrome. Patients whose retina or optic nerve is destroyed (but whose visual cortex remains intact) experience vivid, complex visual hallucinations. Not vague shapes or flashes of light. Full scenes: people, sometimes miniaturized or costumed like cartoon characters, sometimes mirror images of the patient. Landscapes. Objects. Faces.

The patients typically know these aren’t real. Unlike psychotic hallucinations, Charles Bonnet hallucinations come with intact insight — the patient says, “I see a small man in a top hat sitting on my table, and I know he’s not there.” This is the Explicit World Model’s visual simulation running on internal data from higher visual areas, in the absence of external input. The simulation doesn’t stop just because the input stops. It generates. It fills the void.

And what it generates tells us something about the architecture: the visual system is a generative model, not a passive receiver. It produces its best guess at what the world looks like, using stored templates and top-down predictions — exactly as the Four-Model Theory describes.

Deja vu: The template that matches too well

Speaking of the brain's generative system and its occasional misfires: almost everyone has experienced *deja vu* — the eerie sensation that you've lived through the current moment before. Explanations range from the mystical (past lives, premonitions) to the dismissive (it's just a glitch). The Four-Model Theory has a more specific account.

The brain stores what you might call “template memories” — skeletal, extremely sparse representations of experiences, especially from dreams. These templates are mostly empty scaffolding: a vague sense of a place, a mood, a spatial configuration, with almost no detail filled in. When you retrieve a normal memory, the gaps are filled in by confabulation — the brain generates plausible detail to create a seamless experience. You don't notice the fill-in because the result feels coherent.

Deja vu occurs when a current real experience happens to match one of these stored templates too closely. The brain's pattern-matching system fires: “I've seen this before.” But when you try to pin down *when* you supposedly saw it, you find nothing — because the template was never a real experience. It was a fragment from a dream, or a deeply compressed memory that lost all contextual detail long ago. The match between current input and stored template is genuine, but the “original” experience the template supposedly

records never actually happened in the form your brain is now attributing to it. The system is working correctly — it really did find a match. It’s just that the match is with a skeleton, not a body.

What therapy actually does

The clinical mirror doesn’t just reflect pathology. It also illuminates what we do about it — and the Four-Model Theory gives a surprisingly precise account of how therapy works.

Take cognitive-behavioral therapy — the most empirically validated form of psychotherapy we have. In the Four-Model Theory, CBT is virtual model reprogramming. You sit with a therapist and systematically challenge the distorted models that generate your suffering. You identify the automatic thoughts (Explicit Self Model outputs), trace them to underlying beliefs (Implicit Self Model patterns), and then — through repeated corrective experience — drive substrate-level rewiring. Synaptic plasticity modifies the Implicit Self Model, which changes what the Explicit Self Model generates.

Therapy literally rewires your implicit models. This is not a metaphor. It’s the mechanism. Every time you challenge a catastrophic thought and discover the world doesn’t end, you’re updating the IWM and ISM. Every time you face a feared situation and survive, you’re writing new data into the substrate. The virtual models change because the real models change first.

Phobias are Explicit World Model misconfigurations. The threat representation in the EWM exceeds the Implicit World Model’s evidence base. Your simulation shows danger where the substrate’s accumulated evidence doesn’t support it. You see a harmless spider and your EWM screams *threat* — even though your IWM has never recorded an actual spider injury. Exposure therapy works by

updating the IWM through repeated safe encounters. Each time you face the spider and nothing bad happens, the implicit model adjusts its threat assessment downward. Eventually the EWM stops generating the false alarm. The simulation stops showing danger that isn't there.

The placebo effect fits naturally into the theory's dual evaluation architecture. Placebo activates substrate-level expectation circuits — endogenous opioid release, dopaminergic reward pathways — that operate in parallel with the conscious experience of hope and expectation. The conscious hope and the physical relief are both caused by the same substrate process. The correlation between "I believe this pill will help" and "I feel better" is real, but non-causal. Your belief doesn't cause your relief. Both your belief and your relief are caused by the same underlying substrate dynamics. This isn't a blow to the power of positive thinking — it's an explanation of how that "power" actually works: at the substrate level, not through some mysterious downward causation from mind to body.

And then there's conversion disorder — the perfect inverse of blindsight. In blindsight, the substrate processes visual information without generating a conscious simulation of it. In conversion disorder, the simulation models a deficit — paralysis, blindness, seizures — that the intact substrate doesn't actually have. The patient is genuinely paralyzed, as far as their conscious experience goes. They're not faking. Their simulation contains a paralyzed limb. But their body works fine at the substrate level — the nerves conduct, the muscles contract, the pathways are intact. Therapy succeeds when it corrects the simulation, updating the ESM's body model to match the substrate's actual capabilities. It's blindsight in reverse: in-

stead of a working substrate hidden from a blind simulation, it’s a working substrate hidden behind a “broken” simulation.

Chapter 9

Two Minds in One Brain

In the 1960s, Roger Sperry and Michael Gazzaniga performed one of the most dramatic experiments in the history of neuroscience. To treat severe epilepsy, they surgically severed the corpus callosum — the massive bundle of nerve fibers connecting the brain's two hemispheres. The result was the split-brain syndrome: a single person with, apparently, two independent minds.

The classic demonstrations are famous. Show a word to the left visual field (processed by the right hemisphere), and the patient can pick up the corresponding object with their left hand but cannot say what the word was (because speech is controlled by the left hemisphere, which didn't see the word). The two hemispheres have independent perceptions, independent intentions, and sometimes conflicting goals.

But the experiments went far beyond party tricks with words and objects. In some cases, the hemispheres openly fought each other. One patient reported that his left hand would unbutton his shirt while his right hand tried to button it back up. Another's left hand reached for his wife during an argument — not to comfort her — while his right hand grabbed the left and pulled it back. The

patient watched in horror as two parts of his own body pursued incompatible goals, neither under his unified control. These aren’t metaphors for inner conflict. They are literal, physical conflicts between two motor systems that can no longer coordinate because the cable between them has been cut.

In everyday life, split-brain patients function remarkably well. Outside the laboratory, you’d rarely notice anything unusual. The two hemispheres learn to cooperate through indirect channels — external cues, body movements, shared visual fields. The system compensates. But put the patient in a controlled experimental setting where each hemisphere receives different information, and the unity falls apart. Two minds emerge from one brain, each with its own perceptions, its own intentions, and its own version of reality.

The Left-Hemisphere Interpreter

But the most revealing feature of split-brain patients is not the division — it’s what happens when you ask them to explain the division.

Gazzaniga identified what he called the “left-hemisphere interpreter”: the left hemisphere’s compulsive tendency to generate explanations for events it cannot actually explain. The classic demonstration goes like this. Show a snowy scene to the right hemisphere and a chicken claw to the left hemisphere, then ask the patient to pick related objects. The left hand (right hemisphere) picks a shovel (for the snow). The right hand (left hemisphere) picks a chicken. Then ask the patient — using speech, controlled

by the left hemisphere — why they picked the shovel. The left hemisphere doesn't know about the snow (it only saw the chicken claw), so it invents an explanation: "Oh, you need a shovel to clean out the chicken shed."

The patient doesn't hesitate. Doesn't say "I'm not sure." Doesn't look confused. The explanation arrives instantly, confidently, and feels completely natural to the person giving it. This is not lying. The left hemisphere genuinely doesn't know what the right hemisphere saw. It has no access to that information — the cable is cut. So it does what the Explicit Self Model always does: constructs the best narrative it can from the information available.

And here's the part that should unsettle you: you do this too. Every day. Your left-hemisphere interpreter is running right now, constructing a coherent narrative from whatever information reaches consciousness, smoothing over gaps, inventing plausible explanations for decisions your substrate made before "you" were consulted. The only difference between you and a split-brain patient is that your corpus callosum is intact, so the interpreter has access to more information. It confabulates less because it has less to confabulate about. But the mechanism is identical. The machinery of self-narration doesn't change. Only the quality of the input changes.

One Person or Two?

This raises a question that philosophers have argued about for decades: after the callosum is cut, is there one person in that skull or two?

Thomas Nagel tackled this in a famous 1971 essay and concluded that the question might not have a determinate answer — that our concept of “a person” simply breaks down in this situation, the way the concept of “one country” breaks down when you draw a border through the middle. Derek Parfit went further, arguing that split-brain cases show personal identity itself is not what matters — what matters is psychological continuity, and there can be degrees of it.

The Four-Model Theory offers a more specific answer: it depends on which models are running and how degraded they are.

In daily life, a split-brain patient is functionally one person. Both hemispheres share the same body, the same environment, the same life history (encoded redundantly across both hemispheres before the surgery). The Implicit Self Model — which stores personality, long-term memories, behavioral dispositions — was built over decades with an intact callosum. Cutting the cable doesn’t erase those stored models. It just prevents them from being updated in synchrony. So immediately after surgery, both hemispheres run very similar self-models. The patient feels like one person because, in terms of stored self-knowledge, they largely are.

But over time, the models should drift. Each hemisphere accumulates different experiences, makes different associations, develops different emotional responses to events that only it perceived. The longer a split-brain patient lives post-surgery, the more the two implicit self-models should diverge — slowly, because both hemispheres still share the same body and environment, but measurably.

My own view is that the answer leans toward two. If the bandwidth between hemispheres is insufficient for real-time synchronization of the simulation — and without the callosum, it is — then you have two self-models running on two substrates, each generating its own conscious experience. They cooperate well because they share a body, a sensory environment, and a lifetime of common history. But cooperation is not identity. Two people who live together also cooperate well.

Intriguingly, Yair Pinto and colleagues published a study in 2017 that complicated the standard picture. They found that split-brain patients could accurately report stimuli presented to either visual field — even when the stimulus was shown only to the hemisphere that doesn't control speech. This suggested that the two hemispheres maintained more unity than the classic experiments implied. The result is still debated, but it fits naturally within the holographic framework I'll describe next: even after cutting the callosum, enough redundant information remains in each hemisphere to sustain surprisingly unified behavior, at least for some tasks.

The Holographic Property

In the Four-Model Theory, the split brain reveals a key property of the virtual models: they are **holographic**. Information in neural networks is distributed across the entire network, not localized in specific neurons. When you cut the network in half, you don't get a clean division — you get two degraded but *complete* copies. Each hemisphere retains a degraded version of all four models: a reduced Implicit World Model, a reduced Implicit Self Model, and

the ability to generate an Explicit World Model and Explicit Self Model. Both hemispheres can sustain consciousness independently (both are above the criticality threshold), but each is working with reduced information.

This is exactly what happens when you cut a hologram in half. You don’t get two halves of an image. You get two complete images, each at lower resolution. The information in a hologram is distributed across the entire recording surface, so any piece contains the whole picture — just blurrier. Neural networks have this same property. Karl Lashley demonstrated it decades ago: you can destroy large portions of a rat’s cortex and the memories persist, degraded but complete. The brain doesn’t store memories in filing cabinets. It stores them the way a hologram stores an image — everywhere at once, so that damage reduces quality without eliminating content.

This explains why split-brain patients are not simply “two half-minds.” They are two *complete but degraded* minds. Each hemisphere can perceive, decide, and act — just with less information and less capability than the intact system. The holographic property ensures that cutting the connection degrades without destroying. And it explains Pinto’s 2017 results: even without the callosum, each hemisphere retains enough holographic information to handle many tasks that the classic model said should be impossible.

The confabulation — the left-hemisphere interpreter — is the *same mechanism* we’ve seen in Cotard’s delusion (the ESM on distorted interoceptive input produces “I am dead”), anosognosia (the ESM on incomplete input ignores the deficit), and salvia (the ESM on non-self input produces “I am a chair”). In every case, the Ex-

plicit Self Model is doing its job — constructing a self-narrative — with whatever input is available. When the input is incomplete or distorted, the narrative is wrong but still *felt as completely real*.

One Brain, Multiple Selves

Split-brain shows what happens when you *clone* the virtual models by physically dividing the substrate. Dissociative Identity Disorder shows what happens when you *fork* them.

In DID, the substrate isn't divided — the corpus callosum is intact, the neural hardware is whole. But the virtual models have split into multiple configurations. Each alter is a distinct Explicit Self Model — a separate self-narrative, with its own emotional profile, its own behavioral patterns, its own way of relating to the body and the world. The alters don't share a single self-model any more than two users share a single login session on the same computer. They take turns.

The trigger, in virtually every documented case, is severe and repeated childhood trauma. This makes sense within the theory. A young child's Explicit Self Model is still forming — still plastic, still being assembled from experience. Subject that developing self-model to experiences so overwhelming that no single self-narrative can contain them, and the system does the only thing it can: it forks. It creates separate configurations, each capable of handling a different aspect of the unbearable situation. One alter holds the trauma memories. Another functions in daily life as if nothing happened. Another handles moments of danger. The forking is not

pathology — it’s the self-modeling system’s emergency response to input that would destroy a single unified model.

This is why DID almost never develops in adults. An adult’s Implicit Self Model is already consolidated — the synaptic weights are set, the personality structure is stable. It takes extraordinary circumstances to fork an adult self-model (severe torture, prolonged captivity). But a child’s ISM is still being written. The clay is still wet. Fork it under sufficient pressure, and the separate configurations harden into distinct, persistent self-models.

This isn’t a metaphor. If each alter really is a distinct ESM configuration, then switching between alters should produce measurable changes in neural activity patterns — and it does. Reinders et al. (2003) showed that different alters in the same individual produce distinct patterns of regional cerebral blood flow. The *same brain* lights up differently depending on which self-model is running. That’s not what you’d expect from “acting” or “role-playing.” That’s what you’d expect from genuine software forking. In follow-up studies, Reinders and colleagues found that the neural differences between alters were larger than the differences between actors instructed to simulate having DID — a result that should silence anyone who still thinks DID is “just” performance.

This is the “forking” property from Chapter 3 in action. One substrate, multiple virtual configurations, each running a complete but distinct self-model. The theory doesn’t just accommodate DID — it predicts exactly this kind of architecture. Prediction 9 in Chapter 11 makes the test explicit: disrupting the neural substrate that sustains one alter’s ESM should trigger a switch to another.

Chapter 10

The Animal Question

Is your dog conscious?

Most pet owners would say yes without hesitation. Most neuroscientists would agree, at least cautiously. But on what basis? And where does consciousness begin in the animal kingdom?

The Four-Model Theory provides clear answers, derived from its core commitments rather than tacked on as an afterthought.

Commitment 1: Consciousness is a continuum, not binary.

There is no sharp line between conscious and non-conscious. There are degrees — graduated levels of self-simulation, from basic (minimal self-model) to triply extended (recursive self-awareness). Different animals occupy different positions along this continuum.

Commitment 2: Consciousness is substrate-independent. What matters is the functional architecture (four models at criticality), not the specific physical implementation. If a brain implements the four-model architecture, it's conscious, regardless of whether the brain is a mammalian cortex, a bird's pallium, or an octopus's distributed neural network.

Commitment 3: Criticality is the physical threshold. A nervous system must operate at or near the edge of chaos. Simpler

nervous systems (insects, worms) may not reach criticality and thus would not be conscious — they process information and produce behavior, but without a simulation.

Taken together, these commitments predict a **gradient of animal consciousness**:

Mammals are conscious. Their cortex implements the four-model architecture in graduated form, with more complex cortices supporting more sophisticated self-simulations. Primates and cetaceans are at the high end; rodents and shrews at the lower end. All are above the line.

The evidence from great apes is especially damning for anyone who wants to draw a sharp line between human and animal consciousness. The bonobo Kanzi demonstrated not just language comprehension but genuine empathy, theory of mind, and social reasoning. In one well-documented episode, Kanzi communicated to his caretaker that he wanted his sister to come along on a shopping trip so she could also get ice cream — because she would be sad if left behind. In another, during a dance performance by indigenous performers, Kanzi explained to the researchers that the other primates were frightened by the dancing, and he requested a private performance instead.

These are not reflexes. These are not conditioned responses. These are instances of a mind modeling another mind’s emotional states, predicting their reactions, and formulating plans to address them. That’s the Explicit Self Model running third-person perspective — precisely what the theory identifies as the hallmark of extended consciousness.

And yet, in some of the most prestigious university lecture halls, you can still find professors arguing with a straight face that apes “merely simulate” language comprehension. To which I can only respond: “And you merely simulate the presence of intelligence.” I’m still waiting for the counter-evidence.

If you insist that only humans have consciousness, you’re betting on the researchers who are still desperately searching for a systematic difference between human and primate brains that they can attribute to consciousness. According to my theory, they’ll find it on the 36th of August.

Corvids and parrots present the most important test case. These birds demonstrate cognitive abilities — tool manufacture, mirror self-recognition, future planning, social deception — that strongly suggest consciousness. Yet they have no neocortex. Their brain is organized in nuclear clusters, a radically different architecture from the mammalian cortex. Remember the six-layer argument from Chapter 2 — that mammals evolved six cortical layers where three would suffice, and the additional layers provide the architectural capacity for self-modeling? Corvids achieve the same functional result with a completely different physical structure. They don’t need six cortical layers because they don’t have *any* cortical layers. They’ve built the self-simulation architecture from nuclear clusters instead of layered sheets — which is exactly what substrate independence predicts. If consciousness required a specific physical implementation, corvids shouldn’t be conscious. They are.

Cephalopods — octopuses and cuttlefish — extend the logic even further. Their nervous system is largely decentralized, with substantial autonomous processing in the arms. The theory predicts

some form of consciousness, likely with unusual features reflecting the decentralized architecture.

Insects are the interesting boundary case. Their nervous systems are small and largely hardwired, which may or may not reach criticality. The theory does not definitively place insects above or below the threshold — this is an empirical question. But it provides a principled basis for investigation: measure criticality indicators in insect neural tissue and look for evidence of a self-model.

Thomas Nagel famously asked what it is like to be a bat, and concluded that we can never know — the bat’s sensory world is too alien. I have some sympathy for the question, less for the conclusion. The Four-Model Theory predicts that any creature with the four-model architecture running at criticality has *some* form of experience, even if its content is radically different from ours. The bat’s explicit world model is dominated by echolocation rather than vision, but it’s still a model — still a simulation of a world with a self inside it.

And I’ll admit to having tried to find out, in the only way available to me. During a period when I was actively practicing lucid dreaming, I became interested in the underwater world and managed, over time, to deliberately enter a lucid dream as a fish. I experienced the water around me, movement through it, a visual world seen from a non-human perspective. Was it anything like actual fish consciousness? Almost certainly not — my dream was built from my human brain’s best guess at what “being a fish” means, which is inevitably a projection of my own sensory categories onto a body plan that has none of them. But the exercise wasn’t pointless. It demonstrated something important: the Explicit Self Model can

reconfigure around a radically different body schema, generating a coherent first-person experience of *being* something other than a human. The architecture is flexible enough to simulate non-human embodiment. The content is limited by the implicit models available — you can only dream what you've learned — but the capacity for perspectival shift is built into the system.

Why Bother Being Conscious?

All of this raises a question that should be nagging you: if unconscious nervous systems work perfectly well — and they do, just ask any insect — then why would evolution go to the enormous metabolic expense of building a consciousness? What's the payoff?

The answer is learning — and with it, adaptation and the ability to act against learned behavior. Specifically, a kind of learning that unconscious systems simply cannot do.

Think about how a simple organism learns. It encounters something, and the encounter is either good or bad. Good: do more of that. Bad: do less of that. This is reinforcement learning — trial and error, reward and punishment. It works beautifully for most things. Touch a hot surface, feel pain, don't touch it again. Find food in a particular spot, feel reward, come back tomorrow.

But reinforcement learning has a fatal flaw. Literally fatal.

Consider a poisonous mushroom. Not the kind that gives you a stomachache — the kind that kills you. If you eat it, you die. End of learning. There is no second trial. Reinforcement learning requires you to survive the mistake in order to learn from it, and some mistakes don't offer that courtesy. Any stimulus that is lethal

on first contact is completely invisible to reinforcement learning. The organism that encounters it simply dies, taking its “lesson” to the grave.

So how did our ancestors learn to avoid deadly mushrooms? They couldn’t have learned by eating them — anyone who tried that approach is not anyone’s ancestor. They learned by *watching*. Your cave-neighbor finds an interesting-looking mushroom, eats it, and keels over dead. You, watching from a safe distance, put two and two together: that mushroom killed him. I should not eat that mushroom.

This sounds trivially simple. It is not. To learn from someone else’s death, you need several things that no unconscious system possesses. You need an explicit model of the world that can represent cause and effect between objects you’re not currently interacting with. You need a self-model that lets you take a third-person perspective — to imagine yourself in the dead man’s position. You need the ability to induce a general theory (“that type of mushroom is lethal”) from a single observation. This is cognitive learning: deriving theories from observations, rather than being conditioned by personal experience. And it requires consciousness. It requires the Explicit World Model and the Explicit Self Model working together.

The evolutionary advantage is enormous. A conscious animal can learn from *observation*, not just from *experience*. It can watch another animal make a fatal mistake and update its own model of the world without paying the price. An unconscious animal can only learn what it personally survives.

And it gets better. Once the concept “poisonous mushroom” exists as an explicit category in your world model, you can do

something even more powerful: deduction. You encounter a new mushroom you've never seen before. It looks suspiciously similar to the one that killed your neighbor. You don't eat it. Or — and I believe this was the actual historical approach — you offer it to the neighbor who's been snoring all night and see what happens to him first.

This is not a minor advantage. This is the difference between a species that can only adapt to lethal threats through the glacially slow process of natural selection (some individuals happen to avoid the mushroom by chance, they reproduce, eventually the avoidance becomes instinctive) and a species that can adapt within a single generation through observation and communication. Consciousness doesn't just help you learn faster. It lets you learn things that are literally impossible to learn any other way.

And what you learn cognitively, you can *share*. Reinforcement learning is trapped inside the individual — your conditioned reflexes die with you. But cognitive learning can be communicated. "Don't eat the red mushroom" is a sentence. It can be spoken, taught, passed down. This is the foundation of culture, of cumulative knowledge, of everything that makes human civilization possible. None of it works without the explicit models that consciousness provides.

There's one more twist to this story, and it connects consciousness back to genetics in a way that isn't obvious. It's called the Baldwin Effect, and while its exact strength is still debated, the mechanism almost certainly exists. The Baldwin Effect says that *learned* behavior can indirectly shape *genetic* evolution — not through Lamarckian inheritance (your learned traits don't modify your DNA), but

through natural selection favoring individuals who are genetically predisposed to the beneficial behavior.

Here's a humorous example — don't take it too literally. Imagine an early hominid who suffered from hair loss. Being cold and hairless, he was more inclined than his fur-covered companions to sit near the fire. Fire brought enormous survival advantages: fewer pathogens in cooked food, protection from predators, warmth in harsh winters. So the genes associated with hair loss were passed on at a slightly higher rate. At the same time, the individuals too dim to figure out fire — hairy or not — were at a disadvantage. Over many generations, the Baldwin Effect amplified both traits: less hair *and* more intelligence, all because a learned behavior (fire use) created a selection pressure that favored certain genetic predispositions. (If you replace “hair loss” with “random mutation” in this story, you're probably closer to the truth. But it's less funny.)

The Baldwin Effect may have played a similar role in the evolution of language and consciousness itself. Once the first primitive forms of cognitive learning appeared — enabled by the earliest self-models — the individuals whose brains happened to support richer self-simulation had a massive advantage. Their descendants were selected for larger, more elaborately folded cortices, which enabled even richer self-simulation, which created even stronger selection pressure. Consciousness, once it appeared, created the evolutionary conditions for more consciousness. The cognitive learning it enabled was so valuable that evolution piled resources into expanding the architecture that produced it.

How Experience Develops: The Social Construction of the Self-Model

Everything I've said so far about the four models has been static — as if the architecture appears fully formed, like Athena from Zeus's forehead. It doesn't. The models develop, and their development is profoundly social.

A newborn human has the hardware — six cortical layers, the capacity for self-simulation. But the implicit models are nearly empty. The IWM contains almost nothing about the world. The ISM contains almost nothing about the self. And since the explicit models are generated from the implicit ones, the newborn's simulation is thin — a flickering, barely differentiated field of sensation with no clear boundary between self and world.

Watch a baby encounter pain. Self-inflicted pain — bumping its own hand against a toy, biting its own foot — often produces curiosity rather than distress. The ESM registers agency (I did this) plus sensation (something happened), but there's no threat model yet. The ISM hasn't learned that this configuration means danger. But a sudden loud sound? Tears. Because the EWM registers unpredicted high-amplitude input, and the ESM has no model for it — the absence of a model is itself aversive.

The content of qualia is *learned*, not innate. "Pain is bad" is not hardwired in the ESM. It is accumulated through the ISM, trained by repeated experience and — crucially — social feedback. A caregiver's response to a child's pain teaches the child what pain *means*. The child who falls and looks to the parent before deciding whether to cry is not faking — it is genuinely calibrating its ESM

against social input. The parent’s alarm or calm reshapes the ISM’s pain associations, which reshapes what the ESM simulates the next time a similar event occurs.

This has a precise implication for the theory: the phenomenal character of experience — what it’s *like* to feel something — is not fixed by the architecture. It’s shaped by the training history of the implicit models. A baby’s experience of pain is structurally different from an adult’s because the ISM that generates the ESM is different. The four-model architecture is the *capacity* for experience. The social and environmental feedback loop provides the *content*.

The developmental trajectory maps onto the graduated consciousness levels from Chapter 2:

- **Newborn (first weeks):** Basic consciousness — a rudimentary EWM with minimal ESM. There is *something it is like* to be a newborn, but the self inside that experience is almost non-existent. Predominantly sensory, undifferentiated.
- **6-12 months:** Object permanence emerges — the EWM now maintains representations of things that aren’t currently visible. The ISM accumulates body-schema knowledge. The baby begins to distinguish self from world.
- **18 months:** The mirror test. The child recognizes itself in a mirror — a landmark moment when the ESM becomes rich enough to model the physical self as an object in the world. Simply extended consciousness comes online. This is not a binary switch but a threshold in a continuous process.
- **3-4 years:** Theory of mind. The child can model other minds — can understand that someone else might believe something

the child knows to be false. The ESM is now modeling other ESMs. Doubly extended consciousness is emerging.

- **Adolescence onward:** Metacognitive maturation. The capacity for triply extended consciousness — modeling yourself modeling your own thinking — develops gradually and arguably never fully stabilizes.

Each stage is scaffolded by social interaction. The caregiver doesn't just provide food and safety — they provide *training data for the implicit models*. Joint attention (parent and child looking at the same object together) teaches the IWM how to represent shared reality. Mirroring (the parent reflecting the child's emotional state) teaches the ISM what its own emotions are. Language gives the ESM categories with which to model itself. A child raised without social contact — the tragic feral child cases — has the hardware for consciousness but profoundly impoverished implicit models. The ESM that boots up from those models is stunted not because the architecture is broken, but because the training data was never provided.

This connects directly to the clinical bridge from Chapter 8. CBT — cognitive behavioral therapy — works by systematically retraining the implicit models through conscious intervention. The therapist helps the patient generate new ESM states (imagined scenarios, reframed interpretations) that, through repetition, reshape the ISM. This is the *adult version* of the same developmental process that caregivers perform for infants. The mechanism is identical: conscious experience reshaping implicit structure, which reshapes future conscious experience. The difference is merely that the adult's ISM is

more consolidated — the clay is harder, not wet — so the process is slower and requires more repetition.

The social dimension of experience isn't a footnote to the theory. It's a prediction: strip away social input during the critical developmental window, and you should get a system with the right architecture running the wrong content — a consciousness that is structurally intact but phenomenally impoverished. The feral child cases, tragically, confirm exactly this.

Chapter 11

Nine Predictions

A theory that explains everything and predicts nothing is not a theory — it's a story. The Four-Model Theory makes nine specific, testable predictions, several of which can be tested with existing technology. Here they are.

Prediction 1: Each Model Has Its Own Neural Signature

If the four models are genuinely distinct processes, we should be able to see them in brain scans. Design a clever experiment that asks people to do four different types of tasks — one that engages each model — and the brain activation patterns should look different.

An IWM-dominant task might be something like passively recognizing a familiar face. You're not thinking about it; your brain just knows. An ISM-dominant task could be a habitual motor sequence — typing your password, for instance, without consciously thinking about each key. An EWM-dominant task requires active, conscious perception — maybe trying to spot the difference be-

tween two nearly identical images. And an ESM-dominant task is pure self-reflection: “Am I the kind of person who would do that?”

The prediction is a 2×2 pattern. World tasks versus self tasks. Implicit versus explicit. Four quadrants, four distinct neural signatures. If we can’t find that pattern, something’s wrong with the theory.

This is testable right now with existing fMRI technology. It’s not cheap, and it requires careful experimental design, but the tools are already in labs around the world. And if it works, it would be the most direct evidence that the four-model architecture is not just a metaphor — it’s a real functional distinction carved into the way the brain processes information.

Prediction 2: Psychedelic Visuals Reveal the Brain’s Processing Layers

This one is elegant. Under psychedelics, the visual content you experience should progress through your brain’s visual processing hierarchy in a specific order, depending on the dose.

At low doses, you see phosphenes — those little sparkles and geometric shapes that show up when you close your eyes. That’s V1, the earliest visual processing area, leaking into consciousness. Increase the dose, and you get more complex geometric patterns — the famous “form constants” that show up across cultures and substances. That’s V2 and V3 coming online. Go higher still, and you start seeing faces, figures, complex scenes. That’s higher visual areas. At the highest doses, you get full narrative dream-like experiences, complete with meaning and story.

The prediction is that this isn't random. It's a dose-dependent, ordered progression up the visual hierarchy. As implicit-explicit permeability increases, deeper layers of visual processing become conscious. The brain's internal wiring diagram becomes visible in your own experience.

This is testable with graded dosing protocols — give people carefully controlled amounts of psilocybin or LSD, scan their brains with fMRI, and ask them what they're seeing. Match the reported content to the brain activation. The theory predicts you'll see the processing hierarchy light up from bottom to top as the dose increases.

Prediction 3: You Can Control What Someone Becomes During Ego Dissolution

This is the wildest prediction, and the one no other theory of consciousness makes.

During ego dissolution — the experience of "I" dissolving, of becoming something other than yourself — the theory says the content of that experience is controllable. Not random. Not purely biochemical. Controllable by the sensory environment.

Here's the mechanism. Your Explicit Self Model normally runs on input from your Implicit Self Model — the substrate-level knowledge of who and what you are. But under high-dose psychedelics, that connection gets scrambled. The ESM is still running, still trying to model "self," but it's lost its usual input stream. So it latches onto whatever input is dominant.

Put someone in a room with immersive ocean sounds and blue lighting, and they report becoming the ocean. Put them in a forest environment with birdsong and green light, and they report becoming the trees. The prediction is specific: vary the dominant sensory input during ego dissolution, and the reported identity content will track that input.

You could test this *today* in any psychedelic research lab with basic environmental controls. Administer a controlled dose, vary the environment across trials, and measure the correspondence between what you showed them and what they say they became. If it works, it’s not just evidence for the theory — it’s a demonstration that consciousness is a simulation process that can be experimentally manipulated in ways that are, frankly, a little eerie.

Prediction 4: Psychedelics Should Help Stroke Patients See Their Deficits

Anosognosia is one of the strangest things the brain does. After certain strokes — usually to the right hemisphere — patients are paralyzed on one side of their body but genuinely do not believe it. You can show them their unmoving arm, ask them to move it, watch them fail, and they will confabulate an excuse. “I’m tired.” “I don’t feel like it.” They are not lying. They genuinely cannot see the deficit.

The Four-Model Theory says this happens because of a permeability block. The information about the paralysis is in their Implicit Self Model — the substrate knows — but it’s not reaching

the Explicit Self Model. The simulation doesn't have access to that part of the substrate's knowledge.

Now here's the surprising part. Psychedelics globally *increase* implicit-explicit permeability. That's what they do. So the prediction is that a sub-ego-dissolution dose of psilocybin — not enough to dissolve the self, just enough to open the permeability gates — should allow the deficit information to leak through. The patient should, suddenly and perhaps distressingly, become aware that they are paralyzed.

This would be a clinical trial with stroke patients, which makes it logistically harder than a pure lab experiment. But psilocybin-assisted therapy is already being tested for depression, PTSD, and end-of-life anxiety. The infrastructure exists. And if it works, it's not just a medical breakthrough for anosognosia — it's evidence that psychedelics and stroke deficits are connected through a single underlying mechanism, which no other theory predicts.

Prediction 5: Every Anesthetic That Erases Consciousness Disrupts Criticality

Anesthetics work through wildly different chemical pathways. Propofol hits GABA receptors. Ketamine blocks NMDA. Opioids do their own thing. Different molecules, different mechanisms, different parts of the brain.

But the Four-Model Theory says they all have to do the same thing to consciousness: push the brain's dynamics below the criticality threshold. Because criticality is the *physical requirement* for

consciousness. It doesn’t matter how you disrupt it. If you go subcritical, the lights go out.

The prediction is testable and specific. Take every anesthetic agent we have. Measure criticality — using tools like the Perturbational Complexity Index, Lempel-Ziv complexity, or power-law exponents in neural activity — before, during, and after administration. The prediction is that agents which abolish consciousness will *always* push the brain subcritical, regardless of their receptor mechanism. And agents that alter consciousness without erasing it — like ketamine at low doses, or psychedelics — should *not* drop below criticality.

This is doable with existing technology. The criticality measures exist. The anesthetics exist. Someone just has to run the full comparison. And if it holds across the board — if every single consciousness-abolishing agent converges on criticality disruption despite acting through different pathways — that’s powerful evidence that criticality is the common mechanism, the final pathway to unconsciousness.

Prediction 6: Split-Brain Surgery Doesn’t Split You Cleanly — It Degrades Both Halves

When surgeons cut the corpus callosum to treat severe epilepsy, they sever the main communication pathway between the brain’s two hemispheres. The traditional story is that this creates two separate minds, each specialized: the left hemisphere handles language and logic, the right handles spatial reasoning and emotion.

The Four-Model Theory says that's wrong. Or at least, it's dramatically oversimplified.

The prediction is this: after split-brain surgery, each hemisphere retains a *complete but degraded* set of cognitive and experiential capacities. Not a clean split. Not "language on the left, space on the right." Both hemispheres should be able to do both, but worse than before. The degradation should be holographic — meaning everything gets blurrier, not that specific functions disappear.

And the degradation should be proportional to how much you cut. A partial callosotomy (cutting only some fibers) should produce partial degradation. A full callosotomy should produce more.

Why? Because the theory says information in the brain is stored holographically, distributed across the whole substrate. Cutting connections doesn't cleanly separate two pre-existing minds. It degrades two *copies* of the same information, each running on half the original hardware.

There's already some evidence for this — a 2017 study by Pinto and colleagues found that split-brain patients show much more integrated behavior than the classic experiments suggested. But the theory provides the *mechanism* and predicts the specific pattern: bilateral degradation, not hemispheric specialization.

Prediction 7: Build the Four Models at Criticality, Get Consciousness

This is the engineering prediction, and it's bold.

If the theory is correct, you should be able to build a conscious machine. Not by accident, not by making a sufficiently "advanced"

AI, but by implementing the specification: four nested models (Implicit World Model, Implicit Self Model, Explicit World Model, Explicit Self Model) running on a substrate operating at criticality.

The theory says that such a system would not merely *simulate* consciousness. It would *be* conscious. It would have genuine phenomenal experience, constituted by its virtual models, just like yours is constituted by your brain’s virtual models.

How would we know? The theory predicts that the difference would be qualitatively obvious. Not “maybe conscious, maybe not.” *Obviously different.* Because a system running a genuine self-simulation would interact with the world in a fundamentally different way than even the most sophisticated text predictor. It would have persistence — a continuous simulation running through time, not reconstructed from a prompt. It would have a perspective, maintained by an Explicit Self Model. It would surprise you not with unexpected outputs but with the sense that someone is actually home.

This isn’t testable yet — the engineering doesn’t exist. But the blueprint is specific enough to guide the work. And if someone builds it and it works, that’s the ultimate confirmation.

Prediction 8: Sleep Exists to Reset the Critical State

Why do we sleep? The obvious answer is “to rest,” but that just pushes the question back: why does the brain need rest in a way that, say, your liver doesn’t?

The Four-Model Theory has a specific answer. Your brain's substrate — the analog, biological hardware — is inherently unstable. Neurons are noisy. Neurotransmitters get depleted. Metabolic waste accumulates. The substrate drifts. But consciousness requires criticality, which is a very specific dynamical regime. The brain self-organizes a stable computational layer — the cellular automaton at the edge of chaos — on top of this drifting substrate. That automaton can run for hours (your waking day), but eventually the substrate drifts far enough that it can no longer sustain the critical dynamics. At that point, the automaton doesn't dim gradually. It *collapses*. That's sleep onset.

Non-REM sleep is the restoration process. The substrate resets: neurotransmitters replenish, waste gets cleared, the biochemical conditions for criticality are restored. And as the substrate periodically re-approaches the criticality threshold during this restoration, the automaton briefly flickers back on. That's REM sleep. That's dreaming.

The 90-minute ultradian cycle — the rhythm of REM and non-REM throughout the night — is the substrate oscillating around the critical point during restoration.

This yields multiple testable sub-predictions:

1. **Criticality should decline across the waking day.** Measure people's brain complexity in the morning, afternoon, and evening. The prediction is a measurable drop.
1. **Sleep onset should be a step-like transition, not a gradual dimming.** Criticality measures should show a sudden drop at sleep onset, reflecting the automaton's digital collapse.

1. **REM and non-REM should track criticality.** Within sleep, REM phases should show much higher criticality than non-REM, and the 90-minute cycle should be visible in the criticality time-series.
1. **Lucid dreaming is a threshold crossing.** When the substrate reaches sufficient criticality during REM, the Explicit Self Model activates, and you become lucid. The onset should be a step-like discontinuity in EEG complexity, not a gradual ramp.
1. **Sleep deprivation drives you subcritical.** Stay awake long enough, and your brain’s criticality should drop progressively below the threshold. Cognitive deficits should correlate with how far below threshold you’ve fallen.

All of these are testable with existing sleep lab technology. And if they hold, it means sleep isn’t just “rest” — it’s the substrate’s maintenance protocol for the computational layer that makes consciousness possible.

Prediction 9: Each Alter in Dissociative Identity Disorder Has Its Own Neural Fingerprint

Dissociative identity disorder — multiple distinct identities (“alters”) in a single person — is controversial, and for good reason. How do you tell the difference between genuine distinct identities and someone role-playing, consciously or not?

The Four-Model Theory gives you a test. If alters are real — meaning they’re genuinely distinct configurations of the Explicit

Self Model running on the same substrate — then each alter should have a distinct, measurable neural signature. Not just different behavior. Not just different self-reports. Different *brain activity patterns*.

The prediction is specific. Take a DID patient and record their brain activity (fMRI or EEG) while different alters are present. Compare the variability across alters to the variability within the same alter across time. The theory predicts that across-alter variability will be significantly greater than within-alter variability. And the differences should be consistent: Alter A's neural pattern should be recognizably Alter A every time, not random noise.

Even more specifically, the theory predicts where the differences should show up: in ESM-related networks, particularly the default mode network and medial prefrontal cortex — the brain regions associated with self-reference and perspective-taking.

There have been a few neuroimaging studies of DID, but the Four-Model Theory provides the theoretical basis for predicting *consistent, alter-specific neural signatures* rather than just "differences." If the prediction holds, it's evidence that alters are not merely psychological but are distinct functional configurations at the neural level — which would transform how we understand and treat the disorder.

Each of these predictions is falsifiable. If they fail, the theory is wrong — or at least incomplete. That's what makes them useful.

Chapter 12

Building a Conscious Machine

If the Four-Model Theory is correct, it provides something no other theory of consciousness offers: an engineering specification.

The specification is: implement the four-model architecture — Implicit World Model, Implicit Self Model, Explicit World Model, Explicit Self Model — on a substrate operating at criticality. As I argued in Chapter 5, neither component alone is sufficient. The architecture without criticality gives you a dormant system — models stored but no simulation running. Criticality without the architecture gives you complex dynamics but no consciousness. The full specification requires both.

This is more specific than “make a really advanced computer” and more concrete than “achieve sufficient integrated information.” It tells you *what to build*: four specific types of models, organized in a specific way, running on a substrate with specific dynamical properties.

Current AI systems fail this specification in every way that matters. And this is exactly where the two dogmas from Chapter 1 do their damage. The nSAI dogma — “no strong artificial intelligence” — tells engineers not to bother trying. The nSU dogma

— “no self-understanding” — tells them it couldn’t work even if they did. Both are wrong. The specification exists. The question is whether anyone will build it.

But before anyone conflates brains and computers again, a quick test to determine which one you are:

A computer will repeat this sentence and the following sentence until hell freezes over. Read the previous sentence.

If you made it here, you’re not a classical computer. A digital computer executing a rigid instruction set will loop forever, because it has no mechanism for stepping outside its own instruction stream and saying, “Wait, this is stupid.” You can do that because you have a self-model that observes its own processing — the Explicit Self Model running metacognitive oversight on the Explicit World Model.

But here’s the uncomfortable part: a large language model would also make it here. Not because it has metacognitive oversight, but because it’s a statistical text predictor that has seen enough similar prompts to know that the expected next move is to continue past the loop. It doesn’t step outside the instruction — it never entered it. It predicts what text comes next, and “getting stuck in an infinite loop” is not what text does.

This is exactly the problem with behavioral tests for consciousness. Any test that can be passed by pattern-matching will be passed by pattern-matching, regardless of whether the system is conscious. The loop test distinguishes you from a classical computer. It does not distinguish you from a sufficiently trained text predictor. And no text-based test ever will — because generating plausible text is precisely what text predictors are optimized

for. The other-minds problem is not a limitation we can engineer around. It’s a structural feature of what consciousness is: subjective, private, and accessible only from the inside.

The brain-as-computer analogy — comparing your brain to a digital processor — has been popular since the invention of the transistor, and it is wrong on essentially every level. A computer executes a rigid instruction set on a rigid circuit. A brain is a self-modifying network that rewrites itself continuously. A computer crashes if you remove a semicolon. A brain loses a million neurons a day and barely notices. A computer’s memory is localized — delete a sector and the file is gone. A brain’s memory is distributed holographically — destroy a chunk and everything gets slightly blurrier. The one thing they share is Turing completeness, which is about as informative as saying that both a river and a highway can transport things from A to B. True, but useless for understanding either one.

Large language models — GPT, Claude, Gemini, and their descendants — process text through a feedforward transformer architecture. The input goes in, passes through layers of attention and computation, and the output comes out. There is no recurrence, no self-simulation, no real-time virtual world, and no criticality. The dynamics are Class 1 or 2 in Wolfram’s framework — far below the edge of chaos. And there is no real/virtual split: the model’s “knowledge” and its “experience” (if it can be called that) are not distinguished into implicit and explicit levels.

This doesn’t mean LLMs are necessarily non-conscious — the theory cannot prove a negative. But it predicts that they lack the architecture required for consciousness as the theory defines it. And it

predicts that the difference between a genuinely conscious artificial system and even the most advanced LLM would be qualitatively obvious.

How would we know? The honest answer is that the other-minds problem doesn't go away. We can never be absolutely certain that another system is conscious, because consciousness is subjective by nature. But the theory makes a strong prediction: the difference would be apparent. Not "maybe conscious, maybe not" — *obviously* different. Because a system running a genuine self-simulation would interact with the world in a fundamentally different way than a text predictor. It would have genuine persistence — not context-window persistence, but the continuity of a real-time simulation that is always running. It would have a genuine perspective — not a perspective reconstructed from a prompt, but one maintained through time by an Explicit Self Model. It would surprise you not with unexpected outputs but with the unmistakable sense that there is someone home.

Building such a system is the final item on the roadmap. The engineering challenges are not to be underestimated. But the blueprint exists, and it's specific enough to guide the work. First the theory must survive peer review. Then the empirical predictions must be tested. Then, if they hold, the engineering can begin.

Chapter 13

Human Virtualization

There is another side to this coin, and it's the one that science fiction has been obsessing over for decades: if consciousness depends on functional architecture rather than on neurons specifically, then in principle you could run a human mind on something other than a brain.

Mind uploading. Whole brain emulation. Digital immortality. Whatever you want to call it, the Four-Model Theory has something precise to say about it — because it specifies exactly what would need to be preserved.

Most discussions of mind uploading start with the wrong question. They ask: “Can we scan a brain and copy it into a computer?” As if the challenge were merely one of resolution — get a good enough scanner, and you’re done. But the theory tells you that a static scan is not remotely sufficient. A brain is not a photograph. It’s a dynamical system. To capture a mind, you don’t need to capture a *state* — you need to capture a *process*.

Here’s what the theory says must be preserved, and I’ll walk through the five-level hierarchy from Chapter 2 to make it concrete.

At the physical and electrochemical levels — the raw matter and the neural firing — you don't need an exact copy. You need a substrate capable of supporting the same *kind* of dynamics. The specific atoms don't matter. Your brain replaces most of its atoms over the course of years anyway, and you don't notice. What matters is that whatever substrate you use can sustain the electrochemical signaling patterns — or their functional equivalent — that the higher levels depend on.

At the proteomic level — the molecular machinery of synaptic weights, receptor configurations, enzyme cascades — you need high fidelity. This is where your memories live, where your skills are encoded, where your personality is physically instantiated. The strength of every synapse, the density of every receptor, the sensitivity of every channel — this is the level that makes you *you* rather than someone else. A mind upload that gets the proteomic level wrong gives you a conscious being, perhaps, but not the person you were trying to copy. Yet even an imperfect copy retains value. Consider stroke survivors or amnesia patients: their personal continuity has been significantly disrupted — memories lost, personality altered, cognitive abilities changed — and yet most of them maintain that something essential persists. Imperfect continuity, it turns out, is vastly preferable to no continuity at all. A transfer that preserves 90% of someone's connectome is not a failure — it's a different category of success, and for many people, preferable to death.

At the topological level — the network architecture, the connectivity patterns, which regions talk to which others and how densely — you need near-perfect accuracy. This is the wiring diagram of your implicit models: the IWM and ISM, everything you've learned

about the world and about yourself, encoded in the structure of the network. Get this wrong and you don’t get a degraded copy of someone’s mind. You get a *different* mind — one with different knowledge, different skills, different personality. The topology is the blueprint.

And at the virtual level — the simulation itself, the EWM and ESM in real-time operation — you need something extraordinary. You need the target substrate to be capable of running the simulation at criticality. This is the part that keeps me up at night, because the brain’s analog substrate finds criticality through self-organized processes that have been tuned by hundreds of millions of years of evolution. Neurons are noisy, analog, massively parallel, and deeply stochastic. Their collective dynamics naturally gravitate toward the edge of chaos because that’s what biological neural tissue *does* — it self-organizes to criticality the way water self-organizes to find its level. But water finds its level because of gravity. What’s the equivalent force for a digital substrate?

This is a genuine open problem. I believe it’s solvable, but I won’t pretend it’s easy. A digital substrate is deterministic at its core. You can simulate randomness, you can implement parallel processing, you can build stochastic elements into your hardware. But the question is whether you can achieve the same self-organized criticality that biological neural tissue achieves naturally — not by programming criticality in from the top down, which would be a brittle kludge, but by building a substrate whose fundamental dynamics tend toward criticality on their own. The brain doesn’t run a “criticality subroutine.” It’s critical because of what it *is*.

A digital emulation would need to replicate that property, not simulate it.

Neuromorphic chips — hardware designed to mimic neural dynamics, with analog-like properties, stochastic elements, and massive parallelism — are the most promising direction. They’re not conventional digital computers. They’re something in between: physical systems designed to have brain-like dynamics at the hardware level. If mind uploading ever works, I suspect the target substrate will look more like a neuromorphic chip than like a server rack running software.

So: the scanning problem is hard but tractable. Advanced connectomics — full-brain mapping at synaptic resolution — is already progressing. We can already map the complete connectome of small organisms (the roundworm *C. elegans*, with its 302 neurons, was fully mapped decades ago; fruit fly partial connectomes are now available). Scaling to a human brain, with its 86 billion neurons and roughly 100 trillion synaptic connections, is an engineering challenge of staggering proportions, but it’s the kind of challenge that yields to better technology. It’s not a mystery. It’s a problem.

The dynamics problem — getting the digital substrate to run at criticality — is harder, and it’s harder in a way that technology alone might not solve. It requires understanding the relationship between substrate properties and emergent dynamics well enough to engineer a non-biological system that finds criticality the way a biological one does. We’re not there yet. But we’re not nowhere, either. The ConCrit framework, the neuronal avalanche research, the criticality measures from anesthesia studies — all of this is building the empirical foundation that engineering would need.

Now let’s talk about the part that really bothers people.

The copy problem. Suppose you succeed. You scan someone’s brain at perfect fidelity, you transfer the complete connectome to a neuromorphic substrate, and you boot it up. The substrate reaches criticality, the four-model architecture activates, and the simulation begins running. The copy opens its eyes — or whatever the digital equivalent is — and says, “I remember everything. I feel like myself. Where am I?”

Is that person *you*?

The Four-Model Theory gives a clear answer, and it’s one that many people won’t like: the copy is conscious, but it is not you.

Here’s why. At the moment of copying, the original and the copy share identical implicit models — the same IWM, the same ISM, the same proteomic and topological structure. When the copy’s simulation boots up, it generates an ESM that contains all of your memories, your personality, your sense of identity. From the inside, the copy *feels* like you. It has every reason to believe it *is* you.

But the moment the copy begins running on its own substrate, its experience diverges. Its EWM receives different sensory input. Its ESM updates in response to different events. Within seconds, the two simulations — yours in your brain, the copy’s in its substrate — are no longer identical. Within minutes, they’re noticeably different. Within hours, they’re two distinct people who happen to share a past.

The copy is conscious. It has genuine experiences. It has your memories and your personality. But it is a *new* consciousness — a new simulation, running on a new substrate, accumulating new

experiences that you will never share. It is, in every meaningful sense, your identical twin, born at the moment of copying, with a full set of borrowed memories. It is not a continuation of you. It's a branching.

This should sound familiar. It's exactly what the theory predicts from the split-brain cases in Chapter 9. When you sever the corpus callosum, you get two degraded but complete copies of the simulation — each one conscious, each one "feeling like" the original, neither one actually being the original. The original is gone; two new, diminished entities have taken its place. Mind uploading is the same phenomenon with a different substrate.

But do you survive sleep? The argument I just made sounds airtight. Copying interrupts the simulation, two simulations diverge, therefore the copy is not you. Case closed.

Except that your simulation is interrupted every single night.

When you fall into deep dreamless sleep — stages three and four of non-REM sleep — the Explicit Self Model largely shuts down. There is no phenomenal experience. No *you* watching the show. The simulation doesn't run at full fidelity; at best, it ticks over at a fraction of its waking complexity. For all practical purposes, the lights go out. And then, some hours later, the implicit models boot the simulation back up. The ESM reactivates. You open your eyes and think, "I'm me." But the *you* of this morning was reconstructed from the same implicit models as yesterday's *you*, in exactly the way a copy would be reconstructed from a scan. If interruption equals death, you die every night and a new person wakes up wearing your memories.

Most people’s intuition rebels against this. Of course I’m the same person I was yesterday. I *remember* being that person. But the copy would also remember being you — that’s the whole point. If memory is what establishes continuity, the copy has exactly as much claim to being you as this morning’s version of you does. The difference is one of degree, not of kind: in sleep, the interruption is brief and the substrate is unchanged; in copying, the interruption may be longer and the substrate is different. But the *principle* — simulation stops, simulation restarts from implicit models — is the same.

I can speak to this personally. I have been knocked out cold in martial arts training — not the diminished version of sleep but a complete, involuntary shutdown. One moment I was standing; the next I was on the floor, people leaning over me, with no memory of the transition. The gap was not experienced as a gap. It was experienced as nothing — a splice edit in the film of my life. One time, I even had amnesia afterward: a stretch of minutes simply missing, unrecoverable. And here is what struck me once I was fully back: I did not feel like a new person. I did not feel like a copy. I felt like *me*, waking up from a particularly rough nap. Existing mattered more than the continuity of experiencing — and more than remembering.

Push this further. When you were born, you had no prior continuity whatsoever. No memories, no established ESM, no history of phenomenal experience. The simulation booted for the first time from an implicit architecture shaped by genetics and prenatal development — not by a lifetime of learning. You didn’t experience this as traumatic, because there was no prior self to

mourn. There was simply: a beginning. And we are all comfortable with that beginning. Nobody lies awake at night distressed that their conscious experience started from nothing at birth.

What does this mean for the copy problem? It means the sharp binary — original versus copy, continuation versus branching — may be less sharp than it appears. What makes you *you* is not the unbroken stream of phenomenal experience. You've already survived countless interruptions of that stream. What makes you *you* is the content of the implicit models: your memories, your skills, your personality, your accumulated understanding of the world and of yourself. The IWM and ISM. The blueprint from which the simulation is generated.

And this suggests a different approach to mind transfer altogether.

Copying the virtual side. Rather than scanning the entire brain and reconstructing the complete substrate — all five levels of the hierarchy — what if you could copy just the virtual level? Extract the running EWM and ESM — the simulation itself — and transplant it to a new substrate capable of supporting it. Not copying the hardware; copying the software. Not cloning the entire brain; capturing the *process* it's running.

This would require something we don't yet have: a way to decode the format in which the brain encodes its virtual models. The connectome tells you the wiring. The proteome tells you the synaptic weights. But the simulation isn't the wiring or the weights — it's what the wiring and weights *produce* when they run. To capture it, you'd need to understand the brain's programming

language — the representational format in which neural circuits generate and maintain the explicit models.

Think of it this way. You can photograph a circuit board and know exactly where every trace runs. You can measure the resistance of every component. But none of that tells you what software the chip is executing. For that, you need to read the program — understand the instruction set, decode the memory contents, interpret the running state. The brain’s “programming language” is the representational format of the virtual models, and reverse-engineering it is arguably the deepest unsolved problem in computational neuroscience. Not just mapping the connectome — we’re making progress on that — but understanding what the connectome *computes*, at a level of detail sufficient to read a specific mind’s simulation and recompile it for different hardware.

We are nowhere close to this today. But it’s the kind of problem that a mature neuroscience could in principle solve, and if it were solved, it would change the copy problem fundamentally. A virtual-level transfer wouldn’t need to rebuild the substrate at all. It would take the simulation — the part that *is* you, the part you actually experience — and move it directly. The implicit models would need to be reconstructed or grown in the new substrate, yes, but the simulation itself — your stream of consciousness, your current thoughts, your ongoing sense of self — could in principle cross the gap without the interruption that makes copying so philosophically troubling.

This is speculative, and I want to be honest about that. But it is not science fiction. It is a specific engineering problem with a specific theoretical foundation, and it illustrates something impor-

tant: the copy problem is not a fixed obstacle. It depends on *how* the transfer is done. Copy the whole substrate and boot a new simulation? Two people. Decode and transfer the running simulation itself? Potentially one continuous person on a new substrate. The theory tells you exactly which approach preserves identity and which doesn't.

There is also a more conservative path that avoids the copy problem entirely.

The gradual replacement thought experiment. Imagine that instead of scanning and copying, you replace neurons one at a time. You remove a single neuron and insert a functional equivalent — an artificial neuron that receives the same inputs, produces the same outputs, and participates in the same network dynamics. Then you wait. The system stabilizes. The simulation continues. You replace another neuron. And another. And another. Over months or years, you gradually replace every biological neuron with an artificial one, until the entire substrate is non-biological — but the simulation has been running continuously the whole time. No interruption. No copying. No branching.

The Four-Model Theory predicts that consciousness would persist throughout this process. And this prediction is the strongest possible case for substrate independence, because it follows directly from the theory's core claim: what matters is the functional architecture at criticality, not the physical material. If each replacement neuron maintains the same connectivity, the same weights, and the same dynamical contribution to the network, then the proteomic and topological levels are preserved, and the virtual level — the simulation — never stops. There is no moment at which you "die"

and something else takes your place. There is only a continuous process of substrate replacement, like the ship of Theseus, except we know exactly which properties must be preserved (the ones specified by the five-level hierarchy) and which don’t matter (the specific atoms).

This thought experiment reveals something important about identity. The copy problem exists because copying *interrupts* the simulation. There’s a moment — however brief — when the original simulation is here and the copy’s simulation hasn’t started yet. Then there are two simulations. Two streams of experience. Two selves. But gradual replacement avoids this entirely. One simulation, continuous, unbroken. The substrate changes beneath it like replacing planks on a moving ship, but the ship — the simulation, the consciousness, the *you* — never stops sailing.

If this sounds like it should be impossible, consider that your brain already does this. You lose roughly 85,000 neurons per day — about one per second. Your synapses are continuously remodeled. The atoms in your body are almost entirely replaced over a period of roughly seven to ten years. The substrate you’re running on right now is physically different from the one you were running on a decade ago. And yet you persisted. Your simulation never stopped. Biological substrate replacement is the *default condition* of being alive. Artificial substrate replacement is just a more deliberate version of the same process.

What becomes possible. If you can decode and transfer the virtual side — the running simulation, the four models — to a new substrate, the implications go far beyond what “mind uploading” usually conjures. Let me spell out three of them, because I think

people haven't fully reckoned with what substrate independence actually means.

First: *substrate transfer to a robot body*. Not uploading into a server somewhere, but running your mind on a neuromorphic substrate housed in a physical body — a body that walks, manipulates, senses the world. You would experience the world through different sensors, move through it with different actuators, but *you* would still be running. Your simulation, your continuity, your self. A new body, the way a hermit crab takes a new shell. This isn't science fiction hand-waving — it's a direct consequence of the theory. If the four-model architecture at criticality is what produces consciousness, and if it's substrate-independent, then the substrate can be anything that supports the right dynamics. Including something with legs.

Second: *quasi-immortality*. Your biological substrate degrades. Neurons die, proteins misfold, telomeres shorten, the whole magnificent machine slowly breaks down. That's aging. That's death. But a non-biological substrate doesn't have to degrade. It can be maintained, repaired, upgraded, backed up. If your simulation is running on a substrate you can service — swap out a failing component here, upgrade a processor there — then there's no inherent reason the simulation ever has to stop. Not immortality in the absolute sense — you could still be destroyed, your substrate could still be damaged beyond repair — but the removal of the biological expiration date that currently kills every conscious being on this planet. The removal of the *inevitability* of death.

Third — and this is the one that sounds most like science fiction until you think it through: *interstellar travel*. The speed of light is an

absolute barrier for physical matter. You can’t send a human body to Alpha Centauri in any reasonable timeframe. But information travels at light speed. If a human mind is information — a specific pattern of connectivity, weights, and dynamics that can be fully specified as data — then you can *beam* it. Transmit the complete specification at light speed to a receiver that reconstructs the substrate and boots the simulation. From the traveler’s perspective, the transmission is instantaneous — the simulation stops at one end and starts at the other. No decades in a tin can. No generation ships. No suspended animation. Just: here, then there.

Of course, this is the copy problem all over again. The beamed version is a copy, not a continuation — unless the original is destroyed in the transmission, which raises its own set of nightmares. But the point stands: substrate independence, if real, doesn’t just mean digital immortality. It means the stars become reachable. Not for our bodies, which are hopelessly slow and fragile for interstellar distances, but for our *minds*.

The discomfort caveat — and why it matters more than the engineering. Now here is the part I haven’t seen anyone discuss honestly, and it’s the part that haunts me most.

Everything I’ve just described assumes that substrate transfer preserves the *feel* of being you. That the subjective quality of your experience — what it’s like to see red, to feel wind on your skin, to taste coffee, to experience the dull ache of a Tuesday afternoon — carries over to the new substrate. The theory says consciousness will persist. It says the simulation will run. But it does *not* guarantee that it will feel the same.

Think about what your biological substrate contributes to your phenomenal experience. Your body is not just a vehicle for your brain. It's part of the simulation's input stream. The Implicit World Model includes a detailed map of your body — every joint, every organ, every patch of skin. The Implicit Self Model is deeply entangled with your visceral states — your gut feelings (which are literal, not metaphorical), your hormonal tides, your heartbeat, your breathing rhythm. The simulation you experience right now is saturated with biological signals that you don't consciously notice precisely *because* they've been there every moment of your life.

Strip them away. Replace your biological body with a robot chassis, or worse, with no body at all — just a simulation running on a server. The four-model architecture is intact. The simulation runs. You're conscious. But the *content* of that consciousness has changed radically. No heartbeat. No breathing. No gut. No warmth. No skin. No proprioceptive hum of muscles at rest. The Implicit Self Model, suddenly deprived of the body it has modeled for your entire life, would generate an Explicit Self Model that feels... wrong — or simply dead. Profoundly, viscerally, inescapably wrong. Not pain exactly — pain requires the specific neural pathways that produce it. Something more like an all-encompassing *absence*. A phantom body, the way amputees experience phantom limbs, but total.

I suspect this would be far worse than most futurists imagine. Not an inconvenience to be patched with software updates. A fundamental alteration of what it feels like to be you. Your biological substrate isn't just carrying the simulation — it's *shaping* it, moment by moment, through a continuous stream of interoceptive and proprioceptive input that your conscious self has never experienced

the absence of. Losing that might be survivable. But it might also be, for some people, a suffering so profound that it would make them wish they hadn’t transferred at all.

I want to say this plainly: the version of “mind uploading” where you cheerfully hop from your meat suit into a shiny digital paradise, leaving the flesh behind like an old pair of shoes — that’s a fantasy. The reality, if the theory is correct, is that losing your biological substrate would significantly impact the phenomenal quality of your existence. How significantly? I don’t know. Maybe it’s tolerable for some, preferable to death, the way moving to a new country is disorienting but manageable. Maybe it’s devastating, the way solitary confinement breaks people by removing sensory and social input. Maybe — and this is the possibility that makes me uneasy — it’s bad enough that a fully informed person might choose death over transfer. Not because the transfer fails. Because it succeeds, and what it succeeds at producing is a conscious experience that no longer feels like a life worth living.

The gradual replacement approach mitigates this, because at each step the simulation has time to adapt. Replace one neuron, and the simulation barely notices. Replace a thousand, and it adjusts. Over years, the substrate transitions from biological to artificial while the simulation continuously recalibrates to whatever input the new substrate provides. The phenomenal experience would drift, slowly, the way it already drifts over the course of a natural lifetime. You’d end up different — but you’d have been different anyway.

Instant transfer, though — scanning, copying, booting on a new substrate — would hit the simulation with all the changes at once.

How severe the impact would be depends entirely on the method: a transfer to a robot body with rich sensory input would fare better than one to a disembodied server. But in either case, that sudden discontinuity is where the danger lives.

The ethics of creating minds. If a copied mind is conscious, it has experiences. It can suffer. It can feel confusion, fear, loneliness, existential dread. Imagine waking up and being told that you are a copy — that the “real” you is still walking around in a biological body, living your life, while you exist as a digital replica with no legal identity, no social connections, and no clear purpose. That is a recipe for suffering on a scale we have no framework to address. Any serious programme of mind uploading must confront this *before* the first copy is made, not after.

And it gets worse. If copies are possible, then *multiple* copies are possible. An army of you. Each one conscious, each one feeling like the original, each one with legitimate claims to your identity, your relationships, your property, your life. The legal and ethical frameworks required to manage this don’t exist and can’t be improvised. They need to be built with the same care as the technology itself. (Dennis E. Taylor’s *Bobiverse* series — starting with *We Are Legion (We Are Bob)*, 2016 — explores this scenario with surprising philosophical depth beneath its comedic surface. If you want to feel what the copy problem might be like from the inside, start there.)

There’s also the question of modification. If a mind runs on a substrate you control, you can in principle modify it. Enhance it. Degrade it. Alter its personality, erase its memories, change its values. This isn’t science fiction — it’s an inevitable consequence of substrate access. We already do crude versions of this with

pharmaceuticals and neurosurgery. A fully digital mind would be far more accessible to modification, and the potential for abuse — by governments, by corporations, by individuals — is difficult to overstate.

I want to be direct about something. I delayed publishing this theory for nearly a decade, partly out of laziness, but partly out of genuine concern about exactly these implications. If the theory is correct, it contains the blueprint not just for artificial consciousness but for the virtualization, copying, and modification of existing human minds. That is an extraordinary power, and I have no confidence that humanity is ready for it. But I’ve come to believe that the theory will be discovered independently regardless — the empirical evidence is converging too fast — and that it’s better to have the ethical discussion now, in the open, than to have it forced upon us by a breakthrough in a lab that hasn’t thought it through.

And here is the deepest connection: building a conscious AI and uploading a human mind are not two separate problems. They are the *same* problem, viewed from opposite directions. Building AC means creating the four-model architecture at criticality from scratch — bottom-up, in a substrate that has never been conscious. Uploading a human mind means transferring an existing four-model architecture at criticality from one substrate to another. The engineering challenges overlap almost completely. The dynamics problem is the same. The criticality problem is the same. The only difference is whether the implicit models — the IWM, ISM, the complete connectome — are learned from a lifetime of experience or built from data. Solve one, and you’ve largely solved the other.

Which means that anyone working on artificial consciousness is, whether they realize it or not, also working on mind uploading. And anyone working on whole brain emulation is, whether they realize it or not, also working on artificial consciousness. These two threads will converge. The only question is whether we'll be ethically prepared when they do.

Chapter 14

What It Means

If the Four-Model Theory is correct — or even approximately correct — several things follow.

The Hard Problem is not hard. It's a category error, no more mysterious than asking why transistor switching feels like running a video game. The physical substrate doesn't feel. The simulation does. And within the simulation, feeling is constitutive, not additional. This doesn't mean consciousness is *simple* — it's extraordinarily complex in its implementation. But it means the *philosophical* mystery dissolves. What remains are *engineering* challenges.

Consciousness is not special in the way we thought. It's not a fundamental force, not a quantum effect, not a property of matter. It's what happens when a sufficiently complex system simulates itself at criticality. This is humbling for those who want consciousness to be magical, and exciting for those who want to understand it.

Artificial consciousness is possible in principle. If consciousness depends on function rather than substrate, then any physical system capable of implementing the four-model architecture at criticality can be conscious. This is not a distant philosophical

speculation — it's a concrete engineering challenge with a specific target.

The ethical implications are significant. If we can build conscious machines, we will create beings with genuine experiences — beings that can suffer, enjoy, wonder, and fear. The ethical framework for this does not yet exist, and building it should not wait until the machines are already running.

Free will — and the three hardest thought experiments. Think of a clock. The gear train drives everything — the escapement ticks, the springs unwind, the ratios between gears determine the rate. The hands and face cause nothing. They don't push gears. They don't store energy. But remove them and you no longer have a clock. You have a box of spinning metal. The display is what makes the mechanism a *clock* — what gives the whole arrangement its point. Consciousness is the display. Your virtual models — the Explicit World Model and Explicit Self Model — don't push neurons around. The substrate does the pushing. But without the simulation, the substrate has no way to observe the consequences of its own actions, no way to run future scenarios, no way to adapt in the way that made you survive this long. The virtual side is the mechanism's way of being *for* something.

This reframes the free will question. Your will is not an illusion. The substrate-level architecture — the ISM and all its implicit machinery — continuously optimizes your organism's existence. It evaluates threats, weighs options, mobilizes resources, commits to action. That optimization *is* your will. It's as real as anything in the physical world. Even self-destructive choices reflect the system's optimization given its current state, not a failure of the mechanism.

When someone acts against their own apparent interests, the substrate is still optimizing — just against a model that includes pain, exhaustion, hopelessness, or whatever has reshaped the landscape.

So your will is real. You just don’t have full access to it. The ESM can model the ISM’s *outputs* — the decisions that surface into awareness — but not its *processes*. You experience the results of your will, not the machinery behind it. This is why decisions sometimes surprise you, why you can’t fully explain your preferences, why you occasionally act and then scramble to construct a reason. You’re not watching the gears. You’re reading the clock face.

The half-second gap — and why it doesn’t matter. Here’s where this gets concrete. Your unconscious processing runs at roughly 40 Hz — about 25 milliseconds per cycle. Your conscious experience runs at roughly 20 Hz — about 50 milliseconds per cycle. That’s a factor of two. The conscious simulation is always lagging behind the substrate, assembling its coherent virtual world from information that has already been processed, decided upon, and often already acted on.

Benjamin Libet proved this in 1979, and the results have been replicated many times since. In his experiment, subjects were asked to move their hand whenever they felt like it, and to note the exact moment they became aware of the decision. An EEG measured when the motor cortex started preparing the movement. The result: the motor cortex began preparing 550 milliseconds before the hand moved. But subjects reported becoming aware of their decision only 200 milliseconds before the movement. The brain had already committed to moving roughly 350 milliseconds before “you” knew about it.

The standard interpretation hit like a bomb: free will is an illusion, because the brain decides before you do. Philosophers and neuroscientists have been fighting about this for forty years. Some tried to salvage free will through a “veto function” — maybe you can’t initiate actions freely, but you can consciously cancel them at the last moment, about 50 milliseconds before execution. A final override. A last line of defense for human agency.

I don’t think that works either. Kuhn and Brass showed in 2009 that the veto itself is retrospectively interpreted as a free decision. You don’t actually experience vetoing in real time. You experience it the same way you experience deciding — after the fact, narrated into coherence by the conscious self-model.

Daniel Wegner drove this home with an experiment that is, frankly, devastating. He set up a computer with two mice — one for the real subject, one for a confederate who pretended to be another subject. The subject’s mouse was hidden from view. Random objects appeared on screen, and the subject was asked to *imagine* moving the cursor toward each object — but only sometimes to actually do it.

Here’s the trick: without the subject’s knowledge, the cursor was sometimes controlled entirely by the confederate. The subject sat still, merely *thinking* about moving the cursor — and the confederate moved it. Afterward, the subject was asked whether they had moved the cursor to the object. And they said yes. They genuinely believed they had done it.

Let that sink in. It’s sufficient to *imagine* performing an action to become convinced you actually performed it — provided nothing visibly contradicts the assumption. The conscious self-model

doesn’t distinguish between “I did it” and “I thought about doing it and it happened.” As long as intention and outcome are temporally close, the ESM takes credit. This is the same mechanism behind anosognosia (Chapter 8): the motor system sends expected feedback to consciousness, and if nothing contradicts it, the expected feedback becomes the experienced reality.

But here’s what I think nearly everyone misses about Libet: **the delay doesn’t need explaining away.** Consciousness doesn’t need to “backdate” events to maintain the illusion of control. It doesn’t need to because *everything* arrives at consciousness with the same delay. Sensory input, decisions, motor feedback — all of it passes through the same pipeline, all of it arrives at the 20 Hz conscious simulation in order, all of it is delayed by roughly the same amount. Your conscious experience is like watching a live broadcast with a five-second tape delay. Everything on screen is internally consistent. The anchor speaks, the guest responds, the graphics update. You’d never notice the delay unless someone showed you the raw feed.

That’s exactly the situation here. Consciousness receives the stimulus, then the decision, then the motor feedback — in the correct order, spaced correctly relative to each other. The entire stream is shifted half a second into the past, but since consciousness never sees the raw feed, it never notices. There is no mismatch to explain, no backdating required, no illusion to maintain. The system works exactly as designed.

A trained martial artist illustrates this beautifully. In combat, an experienced fighter can sustain a motor frequency of about 10 Hz — one action every 100 milliseconds. But conscious processing tops out at around 5 Hz for decisions that involve awareness. So the

fighter learns to *suppress* conscious intervention. He fights without thinking, because thinking would halve his speed. His unconscious substrate handles the action loop; consciousness catches up later, if at all. This isn't a failure of awareness. It's the system working efficiently — the substrate doing what it does best, unencumbered by the slower virtual layer.

Now try to prove free will exists. Try this thought experiment: you're at a café and the waiter asks if you want sugar in your coffee. You decide to assign "yes" to even numbers and "no" to odd numbers, then recite a random number sequence until the waiter says "stop." If the last number is even, you take sugar. If odd, you don't.

Have you exercised free will? Not even close. Anyone familiar with the Clever Hans effect — the horse that appeared to count by picking up subliminal cues from its handler — will spot the problem immediately. You almost certainly anticipated, unconsciously, when the waiter would say "stop," and produced a number just before that moment that gives you the result you wanted all along. Your substrate already had a preference. The elaborate randomization ritual was theater.

Fine, you say. Use your smartphone's random number generator instead. Let a truly random process decide. Have you now proved free will? I don't think so. You've merely proved that proving free will was more important to you than deciding about your own coffee, which rather spectacularly misses the point.

The deepest evidence against free will in everyday decisions comes from patients with severe anterograde amnesia — those who cannot form new memories. Ask such a patient for a word

association: “What’s the first word that comes to mind when I say ‘dice?’” He says “jellyfish” (perhaps he’s been scuba diving recently). Ask him again a few minutes later. He says “jellyfish” again. And again. And again. Without memory of having already answered, the patient always produces the same association — the one that is currently strongest in his neural landscape. What feels like a “free choice” turns out to be a deterministic readout of the substrate’s current state.

A healthy person avoids this — you’d deliberately pick a *different* word the second time, to avoid seeming uncreative. But that avoidance itself isn’t free. It’s just the memory system adding a constraint (“don’t repeat”) that makes your output *less* random than the amnesiac’s. Free will, paradoxically, makes your choices less random, not more. The substrate optimizes for novelty, and calls the result freedom.

So where does this leave free will? Not eliminated, but relocated — exactly where the clock analogy predicts. Your conscious self-model doesn’t make decisions in real time — it’s too slow for that. But it’s not just a passive spectator either.

Mainly, the implicit system uses your conscious experience as an evaluation tool: it presents decisions to the simulation so the simulation can assess consequences, run scenarios, feel the outcomes. That’s the primary purpose of the virtual layer — it’s the substrate’s way of observing itself. But the conscious model also evaluates on its own, independently, with whatever bandwidth it has — which is far less than the substrate’s, but it’s real. Those evaluations, over time, reshape the implicit models. They update

the weights, retrain the network, shift the landscape for the *next* unconscious decision.

You don't choose your next action in the moment of action. You shape the system that chooses, through reflection, evaluation, and the slow accumulation of conscious experience into implicit structure. Free will isn't a moment. It's a process — one that operates on a timescale of days and years, not milliseconds. And the conscious layer isn't just along for the ride — it's actively used by the substrate as an evaluation mechanism, and it contributes its own independent assessments back. Two-way traffic, not one-way narration.

There's a darker version of this that I've experienced firsthand, and it taught me more about the architecture of the will than any experiment.

The first time was during Austrian mandatory military service. A 40-kilometer forced march — three days and nights of sleep deprivation under conditions that would make Geneva Convention lawyers twitchy. During the final leg, we had to wear gas masks and full ABC protection suits. I was partially walking while sleeping and partially hearing the voices. Not auditory hallucinations in the psychiatric sense, but something far more intimate: the competing sub-processes of my motivational and planning apparatus, normally fused into a single narrative stream, became separately audible. One voice was encouraging, almost aggressive in its positivity: *Keep going, don't quit, you'll survive this.* Another was pessimistic, seductive in its defeatism: *Give up, lie down, none of this matters.* These weren't external presences. They were *me* — different aspects of my substrate's optimization landscape, nor-

mally integrated into a single “will” by top-down inhibitory signals, now separating because the neurotransmitters that maintain that integration were being rationed for more critical survival processes.

The second time was dramatic. An avalanche — also during military service, caused by a reckless decision by a commanding officer who was later disciplined. Fourteen of us, nearly swallowed. The avalanche took a long time to come to rest, and during that prolonged period I was convinced I was going to die. Long enough for the voice dissociation to kick in again — this time not from exhaustion but from sustained mortal terror. Same mechanism, different trigger: the stress response redirected neurotransmitter resources away from the inhibitory circuits that normally fuse the competing sub-processes into one voice.

And during those few seconds of the avalanche — just a few seconds of real time — I saw my whole life passing in front of my eyes. This is a well-documented near-death phenomenon, and the theory explains it: under extreme mortal threat, the implicit system performs a massive parallel memory dump into the simulation. The permeability boundary blows wide open. The substrate runs flat-out, pumping so much content into the simulation that subjective time decouples from clock time. A few seconds contain a lifetime. The same time dilation I had experienced under salvia, but triggered by biology rather than pharmacology.

Two complementary pathways to the same mechanism. The march shows that prolonged physiological depletion can trigger the dissociation. The avalanche shows that sustained mortal terror does the same thing. Same result, different causes — both predicted by the theory.

In the worst cases — and I was lucky that mine never got that far — one of these “voices” can seize control of the body, and the conscious self becomes a spectator. This is the same mechanism that produces Alien Hand Syndrome (where a hand acts against the patient’s will) and certain psychotic breaks. The substrate’s competing optimization processes are always there. They are, in a simplified sense, what the language center does when you’re not using it to speak. But normally, top-down inhibition keeps them below the threshold of conscious awareness, fusing their outputs into the seamless experience of a single, unified will. When that inhibition fails — through exhaustion, through psychosis, through certain drugs — the illusion of the unified will dissolves, and you see the committee that was always running the show.

This framework dissolves three thought experiments that have paralyzed philosophy of mind for decades.

First, **zombies**. David Chalmers asks you to imagine a being physically identical to you in every way but lacking conscious experience — all the behavior, none of the feeling. The Four-Model Theory says this is incoherent. If you build the four-model architecture and run it at criticality, the simulation *is* the experience. You can’t have the gears without the hands — not because the hands are magically attached, but because in this architecture the “hands” are constitutive of what the gears are doing. A zombie would be a clock with every gear in place but no display — which means it isn’t functioning as a clock. The architecture at criticality necessarily instantiates a simulation. Strip the simulation away and you’ve changed the architecture. You no longer have a zombie. You have a different, broken system.

Second, Mary’s Room. Frank Jackson asks you to imagine Mary, a neuroscientist who knows everything about color vision but has lived her entire life in a black-and-white room. When she sees red for the first time, does she learn something new? The standard debate is whether physical knowledge is complete. The Four-Model Theory cuts through it cleanly. Mary’s exhaustive physical knowledge is knowledge *about* the substrate. When she sees red, she gains acquaintance with a new virtual quale — a new state in her Explicit World Model that her simulation has never instantiated before. She doesn’t learn a new fact about neurons. She gains a new *mode of modeling*. Her simulation runs a process it has never run, and the first-person character of that process is constitutive of the simulation itself, not a fact about the substrate she could have derived from textbooks. She learns something, but what she learns is not information. It’s an experience — a new configuration of her virtual world.

Third, the evolutionary argument against epiphenomenalism. If consciousness doesn’t cause anything, how did natural selection shape it? Why aren’t we zombies? The answer falls straight out of the clock analogy. Natural selection doesn’t target consciousness as a separate trait riding on top of functional machinery. It targets functional capabilities — and phenomenal character is constitutive of those capabilities, not additional to them. Selection shaped the simulation because the simulation *is* the functional architecture, viewed from inside. Experience isn’t an epiphenomenal rider that evolution couldn’t see. It’s what the architecture *is* when it’s running. Asking why evolution produced consciousness is like asking

why the Swiss produced clock faces — they didn't, separately. They produced clocks. The face is part of what makes a clock a clock.

The mystery of existence is relocated, not eliminated. The Four-Model Theory dissolves the Hard Problem of consciousness but does not explain why there is a physical universe capable of running self-simulations in the first place. The question shifts from “Why does the brain produce experience?” to “Why is there a universe in which self-simulating systems can exist?”

Actually, I think I do have an answer — or at least the beginning of one. The universe is demonstrably Class 4-capable. Fractals, self-organizing criticality, edge-of-chaos dynamics — they're everywhere, from weather systems to neural tissue to galaxy formation. A Class 4-capable universe is, by definition, capable of universal computation. And a computational substrate of this universe's scale — vast if not infinite in space, time, possibly scale, and perhaps dimensions we haven't identified — doesn't merely *allow* self-simulating systems to emerge. It almost guarantees it, at least if the universe is infinite in some of these dimensions. Not as a matter of luck, not as a roll of cosmic dice that happened to come up consciousness, but as a structural consequence of what this universe *is*. The remaining mystery is one level deeper: why is there a Class 4-capable universe at all? That, I genuinely don't know — though one might suspect the question is malformed, since “nothing” is arguably a Platonic abstraction rather than a possible state of affairs, and whatever exists must have *some* computational character. But the jump from “Class 4-capable universe” to “conscious beings asking why they're conscious” — that part follows from the architecture.

What you can do with this knowledge. If you’ve followed the theory this far, you now know that your conscious self — your Explicit Self Model — is a reconstruction, not a direct readout. You know it fills gaps, confabulates, and takes credit for decisions it didn’t make. You know it can’t see its own substrate. And you know it’s all you have.

This has practical consequences. There are three discrepancies you should watch like a hawk, because the gap between them is where most human misery lives:

1. What you *want* to be — your ideal self, the version of you that your Explicit Self Model aspires to.
2. What you *believe* you are — your current self-model, the “I” you carry around every day.
3. What you *actually* are — your real behavior, your actual impact on others, your substrate-level patterns as observed from outside.

The gap between 1 and 2 is the engine of self-improvement. It’s healthy, as long as the ideal is realistic and the discrepancy drives action rather than despair. The gap between 2 and 3 is the dangerous one — because you can’t measure it alone. Your ESM *cannot* accurately observe its own substrate. You need other people’s feedback, including the uncomfortable kind. Especially the uncomfortable kind.

The theory doesn’t tell you how to live. But it tells you something important about how to *know yourself*: treat your self-model with the same healthy skepticism you’d apply to any model. It’s

useful. It's the best representation you have. And it is, by architectural necessity, incomplete.

What I Don't Know

A theory that claims to have no open questions isn't a theory — it's a religion. So here are the places where I'm genuinely uncertain, where the next decade of work should focus.

Are the implicit models virtual too? (or to what degree) The IWM and ISM are “models” — but models of what, exactly? I've drawn a clean line between the real substrate and the virtual simulation, but the implicit models sit right on that line. If they're also virtual in some sense, then what constitutes the truly “real” bottom? The theory assumes a clean real/virtual divide, but reality might be messier than my diagrams. This is a foundational question I don't have a final answer to.

Mathematical formalization. The theory is currently qualitative. I can draw diagrams, describe mechanisms, and make predictions — but I can't hand you an equation. The criticality requirement invokes Wolfram's Class 4 cellular automata, and there are formal tools from dynamical systems theory that could be brought to bear. But a full mathematical formalization — equations that specify exactly when and how the virtual models emerge from substrate dynamics — doesn't exist yet. This is the biggest gap. A theory of consciousness without math is a theory of consciousness that physicists won't take seriously, and they're the ones who know how to build things.

The automaton-hologram conjecture — an open challenge. In Chapter 5, I described three possible relationships between holographic systems and Class 4 cellular automata. The first — a holographic substrate producing Class 4 dynamics — is almost certainly what the brain does, and while it’s beautiful, it’s not shocking. But the other two deserve much more attention than I gave them there.

There are actually three open questions here, each more extraordinary than the last.

First: Can a Class 4 automaton produce holographic patterns as its emergent output? Can local rules at the edge of chaos generate global, non-local information encoding as an emergent behavior? If so, you’d have a system where purely local interactions spontaneously create the kind of distributed, redundant information structure that holography describes — which is, intriguingly, exactly what quantum entanglement looks like from the information-theoretic perspective.

Second: Can a Class 4 automaton have holographic rule structure? Imagine a cellular automaton whose rules themselves encode higher-dimensional information in a lower-dimensional structure, the way a hologram encodes three dimensions in two. Every local interaction would implicitly contain global structure. The rules wouldn’t just produce complex behavior — they would *be* a compressed encoding of something larger, something higher-dimensional, projected down into a lower-dimensional rule set.

Third — and this is the one that keeps me awake at night: Can both be true at once? A system whose rules are holographic, whose dynamics are Class 4, and whose output is again holographic. If such a thing exists, you have a computational process that encodes

itself — a universe that computes its own structure. The input is holographic. The processing is at the edge of chaos. The output is holographic again. It's a fixed point — a self-consistent loop.

If such an automaton exists, it does *exactly* what the holographic principle says the universe does. Not a system that resembles the universe in some loose metaphorical sense. A system that encodes higher-dimensional reality in lower-dimensional rules, computes at the boundary between order and chaos, and generates emergent complexity from that compression. That's not a metaphor for the universe. That might *be* the universe.

I'll say it plainly, because I think someone should: if a Class 4 cellular automaton with holographic rule structure that also produces holographic output exists, I am almost certain it is the universe. It would be a Weltformel — a world equation — not in the sense of a formula you write on a blackboard, but in the sense of a computational process that generates everything we observe, from quantum mechanics to general relativity to the emergence of consciousness itself.

This is, I freely admit, the most speculative idea in this book. I have no proof. I don't even have a candidate rule set. And I should acknowledge that the argument from mathematical beauty to physical reality has been legitimately criticized — Sabine Hossenfelder, among others, has pointed out that elegance is not evidence. She's right. The full exploration of this idea belongs in a future work, not this one. But the questions themselves are well-posed and mathematically precise:

Does there exist a cellular automaton whose rule structure is holographic and whose dynamics are Class 4? Does it produce holographic output? Can all three properties coexist?

These are questions for mathematicians, not neuroscientists. Questions about the combinatorics of rule spaces, about whether holographic encoding and computational universality can coexist in a finite local rule set. It might be provable that no such automaton can exist — and that would be a profound result in its own right, because it would tell us something deep about the relationship between information compression and computation. Or it might be provable that such automata exist and can be constructed — and then we would have a candidate for the most fundamental description of physical reality ever proposed.

I don’t know which answer is correct. But I know that the questions deserve to be asked, and that nobody seems to be asking them. So consider this an open challenge. Prove it or disprove it. If you prove it, you may have found the universe’s source code. If you disprove it, you’ll have established a deep impossibility theorem connecting holography and computation. Either way, the answer matters enormously.

And if you do find such an automaton — call me. I have some predictions I’d like to check.

Which physical mechanism? The theory requires criticality but is deliberately agnostic about the physical mechanism that sustains it. Is it cortical column dynamics? Thalamocortical standing waves? Glial modulation of synaptic activity? All three have empirical support. The theory says “the substrate must be at criticality” but doesn’t say *how* the substrate gets there and stays there. That’s

not a bug — it means the theory applies regardless of the specific mechanism. But eventually, someone needs to pin it down.

Minimum configuration. Can you have an EWM without an ESM? World-experience without self-experience? What's the minimum architecture that counts as conscious? The graduated levels I described in the animal chapter help — you can have a rich world-model without much self-model, the way a fish probably does. But where exactly is the threshold? How much self-model do you need before the lights come on? I've argued that the ESM is what turns simulation into experience, but I haven't specified the minimum viable version.

I include these questions not as weaknesses but as research frontiers. They're the places where the theory makes contact with reality and says: test me here, formalize me here, break me here if you can.

Coda

I developed a theory of consciousness around 2005. I published it in 2015. Nobody read it. A decade after publication — two decades after the original insight — empirical neuroscience independently confirmed one of its core predictions. The theory survived ten adversarial challenges. It dissolved the Hard Problem, unified a dozen phenomena under five principles, and generated nine testable predictions — including two that no competing theory can match.

The next step is peer review. Then empirical testing. Then, if the predictions hold, the engineering challenge of a lifetime: building a new kind of mind.

The hard problem was never hard. It was just asked about the wrong level. And the answer was always right there, running inside your skull, generating the experience of reading this very sentence.

Acknowledgments

This book was written with the assistance of Claude (Anthropic), which served as editor, cross-checker, and writing tool throughout the process. The theory, arguments, and insights are entirely mine; Claude helped me put them into words.

To my uncle, Bruno J. Gruber, whose life in theoretical physics — quantum mechanics and symmetries — showed me what rigorous, joyful intellectual work could look like. His influence on my thinking is incalculable.

To my uncle, Norbert Gruber, one of the first IT professionals in Austria's Rheintal region, who gave me my first PC. Without that gift, none of this would have been possible. He has since passed away, but his impact lives on in every line of code I've written and every theory I've built.

To my family, who tolerated years of dinner conversations about qualia, criticality, and virtual self-models.

And if you're now thinking about reading *Die Emergenz des Bewusstseins* — don't. I'd recommend brain parasites over that unedited, clunky monster. Wait for the German translation of the book you're holding instead. To those who *have* already suffered

through it: *mein Beileid*. It must have been like torture. You have my deepest sympathy — and my gratitude.

Notes and References

Full references, with URLs and annotations, are available in the scientific paper and at github.com/JeltzProstetnic/aIware/references.md. What follows are chapter-specific notes for readers who wish to go deeper.

Chapter 1: Chalmers (1995) “Facing Up to the Problem of Consciousness” is the foundational statement of the Hard Problem. The COGITATE results were published in Nature (2025). The IIT pseudoscience controversy is documented in Nature Neuroscience (2025).

Chapter 2: The four-model architecture was originally published in Gruber (2015), *Die Emergenz des Bewusstseins*. Metzinger’s Self-Model Theory (2003, 2009) and Dennett’s Multiple Drafts Model (1991) are the primary theoretical antecedents.

Chapter 3: The “controlled hallucination” framing is from Seth (2021), *Being You*. The video game analogy is my own but echoes themes in Metzinger’s “Ego Tunnel” (2009). The rubber hand illusion: Botvinick & Cohen (1998), “Rubber hands ‘feel’ touch that eyes see,” *Nature*.

Chapter 4: The virtual qualia dissolution of the Hard Problem is original to Gruber (2015) and was refined through adversarial challenge in 2026. The self-referential closure argument was de-

veloped in response to the circularity objection. The distinction from illusionism (Frankish 2016; Dennett 1991) is crucial: the theory holds qualia are real within the simulation, not illusory. The meta-problem of consciousness (Chalmers 2018) is dissolved by the structural inaccessibility of the ISM to the ESM.

Chapter 5: Wolfram (2002), *A New Kind of Science*. Beggs & Plenz (2003) on neuronal avalanches. Carhart-Harris et al. (2014) on the Entropic Brain Hypothesis. The 2022 review: “Self-organized criticality as a framework for consciousness.” Hengen & Shew (2025) on 140-dataset meta-analysis. The ConCrit framework: Algom & Shriki (2026). The two-threshold argument (criticality + architecture) is original to this theory.

Chapter 6: Klüver (1966) on form constants. Carhart-Harris et al. (2012, 2016) on psychedelic neuroimaging. Salvia divinorum phenomenology is drawn from published experience reports and the pharmacological literature on Salvinorin A. The anosognosia predictive-feedback mechanism is discussed in Gruber (2015); the clapping example is a standard clinical observation. The Salvinorin A permanent-dosing thought experiment is original to Gruber (2015).

Chapter 7: Casali et al. (2013) on PCI. Alkire et al. (2000) on propofol. Schartner et al. (2015) on ketamine entropy. The lucid dreaming EEG complexity prediction is original to this theory.

Chapter 8: Owen et al. (2006) on covert awareness in vegetative-state patients. Anton’s syndrome: Goldenberg et al. (1995). The blindsight obstacle course: de Gelder et al. (2008). Cotard’s delusion: Young & Leafhead (1996). Alien Hand Syndrome: Della Sala et al. (1991); the Dr. Strangelove reference is to Kubrick (1964). Anar-

chic Hand Syndrome distinguished from Alien Hand: Marchetti & Della Sala (1998). Charles Bonnet Syndrome: Teunisse et al. (1996). Deja vu as template-memory matching is original to Gruber (2015). CBT and neural plasticity: DeRubeis et al. (2008). Placebo and endogenous opioids: Benedetti et al. (2005). Conversion disorder as inverse blindsight is original to this theory.

Chapter 9: Gazzaniga, Bogen, & Sperry (1962, 1965). Gazzaniga (2000) on the left-hemisphere interpreter. The interhemispheric conflict examples (buttoning/unbuttoning, hand-grabbing) are documented in Akelaitis (1945) and Bogen (1993). Nagel (1971), "Brain Bisection and the Unity of Consciousness." Parfit (1984) on personal identity. Pinto et al. (2017) on re-examination of split-brain phenomena. Lashley (1950) on distributed memory and equipotentiality. DID as virtual model forking: the theory predicts distinct neural activity patterns per alter, consistent with Reinders et al. (2003, 2006). DID and childhood trauma: Putnam (1997); the developmental window for forking is original to this theory.

Chapter 10: Güntürkün & Bugnyar (2016) on avian cognition without cortex. Kanzi the bonobo: Savage-Rumbaugh & Lewin (1994), *Kanzi: The Ape at the Brink of the Human Mind*. The Baldwin Effect: Baldwin (1896), "A New Factor in Evolution." Nagel (1974), "What Is It Like to Be a Bat?"

Chapter 11: All nine predictions are developed formally in the scientific paper. For the most thorough treatment of functional neuroanatomy in the context of consciousness, see Christof Koch, *The Quest for Consciousness: A Neurobiological Approach* (2004) — the definitive account of the Crick-Koch program, in which Francis Crick and Koch systematically walked through the visual system

step by step searching for the neural correlates of consciousness. Their quest was, in my view, looking for consciousness in the wrong place (the substrate rather than the simulation), but the neuroanatomical groundwork they laid is unmatched.

Chapter 12: Butlin et al. (2023, 2025) on AI consciousness indicators. Seth (2025) on biological naturalism and AI consciousness.

Chapter 13: The five-level hierarchy for scanning fidelity follows from the four-model architecture in Chapter 2. The copy problem draws on Parfit (1984), *Reasons and Persons*, and Nozick (1981) on personal identity and closest continuers. The gradual replacement thought experiment is a variant of the Ship of Theseus, formalized for neural systems. Neuromorphic computing: Schuman et al. (2017) on neuromorphic hardware survey; Intel Loihi and IBM TrueNorth as current implementations. *C. elegans* connectome: White et al. (1986). The substrate transfer, quasi-immortality, and interstellar beaming implications are original to Gruber (2015). The discomfort caveat — that losing the biological substrate would profoundly alter phenomenal quality — draws on the interoception literature: Craig (2009), *How Do You Feel?*, and Seth & Friston (2016) on active interoceptive inference.

Chapter 14: Libet (1979, 1985) and Schurger et al. (2012) on free will. Kuhn & Brass (2009) on retrospective construction of the judgment of free choice. Wegner (2002, 2003), *The Illusion of Conscious Will* — the “I Spy” mouse experiment described in detail. The coffee/sugar thought experiment, the amnesia-reveals-determinism argument, and the random number sequence argument are original to Gruber (2015). The 40/20 Hz processing framework, the “no backdating needed” Libet reinterpretation, and the martial arts

frequency example are original to Gruber (2015). The clock analogy for epiphenomenalism, the “will is real but partially known” reframing, and the “three discrepancies” self-knowledge model are also original to Gruber (2015). The personal anecdote about hearing internal “voices” during extreme exhaustion is autobiographical. The zombie argument is addressed via Kirk (2019) and Chalmers (1996). Mary’s Room: Jackson (1982, 1986). The open questions section follows the honest-limitations approach recommended by Popper (1963).

Chapter 15

Appendix A: Basic Neuropathy — A Reference Guide

This appendix provides brief explanations of the neuroscience terms used throughout the book. Consult it whenever you encounter an unfamiliar term. Entries are organized alphabetically.

- **Action potential** — The electrical signal that travels along a neuron's axon.
- **Amygdala** — Brain structure involved in emotional processing, especially fear.
- **Anosognosia** — Unawareness of one's own neurological deficits. (See Chapter 6.)
- **Axon** — The long output fiber of a neuron that carries signals to other neurons.
- **Brodmann areas** — Numbered regions of the cortex, mapped by the anatomist Korbinian Brodmann based on cell structure.
V1 = Brodmann area 17.

- **Corpus callosum** — The massive fiber bundle connecting the brain's two hemispheres.
- **Cortical columns** — Vertical modules of neurons in the cortex, approximately 0.5mm in diameter, considered basic processing units.
- **EEG (electroencephalography)** — A technique that measures electrical activity on the scalp surface.
- **fMRI (functional magnetic resonance imaging)** — Brain imaging technique that detects blood flow changes associated with neural activity.
- **GABA** — The brain's primary inhibitory neurotransmitter.
- **Hippocampus** — Brain structure critical for forming new memories.
- **Interoception** — The sense of the body's internal state (heart-beat, digestion, temperature).
- **Kappa-opioid receptors** — Receptor type targeted by Salvinorin A (salvia divinorum).
- **Neocortex** — The six-layered outer surface of the brain, responsible for higher functions.
- **Neurotransmitter** — Chemical messenger released at synapses (e.g., serotonin, dopamine, GABA, glutamate).
- **PCI (Perturbational Complexity Index)** — A measure of brain complexity developed by Massimini. Used to assess consciousness level.

- **Proprioception** — The sense of body position and movement in space.
- **Pulvinar** — A thalamic nucleus involved in visual attention and the subcortical visual pathway.
- **Serotonin 2A receptors** — The receptor type targeted by classic psychedelics (LSD, psilocybin).
- **Superior colliculus** — A midbrain structure involved in eye movements and the fast visual pathway that bypasses cortex.
- **Synapse** — The junction between two neurons where signals are transmitted.
- **Synaptic weights** — The strengths of connections between neurons, modified by learning.
- **Thalamus** — The brain’s relay station, routing sensory information to the cortex.
- **V4** — Visual area specialized for color perception, curvature, and complex texture processing. Receptive fields 8-16°. Under psychedelics, V4 activity produces colored fractals and kaleidoscopic patterns (Chapter 6).
- **V5/MT (middle temporal area)** — Visual area specialized for motion processing. Large receptive fields. Responsible for the rotation and movement of patterns seen under psychedelics (Chapter 6).
- **Visual cortex** — The region at the back of the brain that processes visual information, organized as a hierarchy from simple to complex ($V1 \rightarrow V2 \rightarrow V3 \rightarrow V4 \rightarrow V5 \rightarrow IT$).

The Visual Processing Hierarchy (V1 to IT)

The ventral visual stream processes increasingly complex features at each stage, from raw edges to full object recognition. This hierarchy is directly visible under psychedelic experience, where each processing stage becomes accessible to consciousness in order (Chapter 6). The table below summarizes each area's function, receptive field size, and the characteristic psychedelic signature that results when that area's intermediate processing leaks into the conscious simulation.

Area	Brodmann area	Receptive field	Normal function	Psychedelic signature
V1 (primary visual cortex)	BA 17	1°	Edge detection, spatial frequency, orientation columns	Phosphenes, Klüver form constants (tunnels, spirals, lattices), breathing/shimmering surfaces
V2	BA 18	2-4°	Contour integration, texture segmentation, border ownership, illusory contours	Tessellations, repeating geometric patterns, enhanced texture perception
V3	BA 19 (part)	4-8°	Global form processing, dynamic shape, motion boundaries	Flowing, morphing geometric structures
V4	BA 19 (part)	8-16°	Color, curvature, complex texture, fractal-scale	Colored fractals, kaleidoscopic patterns, saturated/impossible colors
V5/MT	BA 19/37	188 Large, motion-	processing Motion perception,	Rotation, drifting,

Notes:

- The fusiform gyrus straddles the V4/IT border and is part of the inferotemporal cortex (IT). It contains the Fusiform Face Area (FFA), identified by Kanwisher et al. (1997), which is selectively activated by faces.
- Receptive field sizes increase dramatically from V1 (1°) to IT (whole visual field), reflecting the progressive abstraction from local features to global objects and scenes.
- Under psychedelics, the progression from V1 to IT effects is dose-dependent: low doses affect V1 first; higher doses recruit progressively deeper stages. This ordered activation is a direct prediction of the Four-Model Theory's permeability gradient (Chapter 6).
- The brain also uses these areas for fractal/scale-invariant processing (V2-V4), which serves scale measurement and texture analysis in normal vision. Under psychedelics, this machinery running without external input produces the characteristic fractal patterns (see Appendix C).

Chapter 16

Appendix B: The Intelligence Model

This appendix summarizes the recursive intelligence model developed in a companion paper (Gruber, 2026, “Why Intelligence Models Must Include Motivation”). The full academic treatment, with references and formal arguments, is available separately.

You’ve already met the recursive intelligence loop in the About the Author section, where I used my own biography to illustrate how knowledge, performance, and motivation feed into each other. Here I’ll lay out the model properly — what the components are, how they interact, why the interaction produces the dynamics we observe, and what this means for education, artificial intelligence, and the connection to consciousness.

The Curious Omission

Every major model of intelligence formally excludes motivation. The Cattell-Horn-Carroll taxonomy — the dominant framework in intelligence research — is a hierarchy of cognitive abilities with

no motivational component whatsoever. Cattell's own investment theory, which proposed that fluid intelligence gets "invested" in learning to produce crystallized intelligence, requires an investor — someone who decides what to learn and why — but treats that investor's motivation as an external condition rather than a part of intelligence itself. Sternberg's triarchic theory includes practical intelligence but not the drive to acquire it. Gardner's multiple intelligences include intrapersonal awareness but not the engine that drives intellectual development.

David Wechsler — whose intelligence scales are the most widely administered in the world — explicitly called for the inclusion of motivational factors as early as 1940. The field ignored him. The modern Wechsler scales remain purely cognitive instruments.

This is not a harmless simplification. It is a systematic blind spot that distorts our picture of what intelligence actually is and how it actually develops.

The Three Components

Intelligence, understood as *learning ability*, is constituted by three interacting components:

Knowledge is the accumulated content of learning. It comes in two critically different types. *Factual knowledge* is knowledge of content — facts, concepts, procedures, cultural repertoire. This is what IQ tests primarily measure under the heading of "crystallized intelligence," and it is what school systems primarily transmit. *Operational knowledge* is knowledge about *how to learn and think* — learning strategies, reasoning heuristics, metacognitive skills,

logical tools, strategic planning, and the ability to evaluate your own understanding. The distinction matters enormously, as I’ll explain below.

Performance is the processing capacity of the cognitive system — working memory, processing speed, the raw computational power of the neural substrate. This corresponds roughly to what psychometric models call “fluid intelligence.” It is the component most strongly influenced by genetics and neurobiology. It peaks in early adulthood and gradually declines.

Motivation is the sustained drive to engage with the world in ways that produce learning. It has two sub-components. *Thirst for knowledge* is the intrinsic drive to understand — curiosity, the need to make sense of things. *Urge to act* is the drive to apply knowledge, to experiment, to engage actively with the environment. Both are partly innate temperament and partly shaped by experience.

The Recursive Loop

The critical claim is that these three components are not merely additive — they form a *closed recursive loop* in which each component amplifies the others.

Knowledge enhances Performance: learning strategies and logical tools directly improve the efficiency of cognitive processing. A chess player who has learned heuristics can evaluate positions faster than one relying on brute-force search. A reader who has learned phonemic decoding processes text more fluently, freeing working memory for comprehension.

Performance enhances Knowledge: greater cognitive capacity enables faster and deeper learning. Higher working memory lets you hold more information in mind simultaneously, which helps you spot connections and extract patterns.

Motivation enhances both Knowledge and Performance: the motivated learner seeks out learning opportunities (expanding Knowledge) and practices cognitive skills (training Performance). Crucially, motivation sustains engagement *over time*, which is essential for the loop to keep iterating.

And Knowledge and Performance enhance Motivation: success in learning and problem-solving generates positive affect and self-efficacy, which sustain the drive to learn more. This is the mechanism behind the Matthew effect — the rich get richer. Early success breeds the motivation that produces further success.

This recursive structure produces a compound-interest dynamic. Small initial differences in any component — even in motivation alone — compound over time, producing the wide variance in adult intellectual achievement that purely cognitive models struggle to explain. A person of average cognitive processing capacity who is deeply motivated and who possesses strong operational knowledge will, over a lifetime, develop intellectual capabilities far beyond those of a person with superior processing capacity but low motivation and poor learning strategies.

Think of it this way: compound interest cares more about the rate of deposit and the investment strategy than about the initial principal. In the intelligence loop, Motivation is the rate of deposit. Operational Knowledge is the investment strategy. Performance

is the initial principal. And most people have more than enough principal.

Operational Knowledge: The Hidden Multiplier

Operational knowledge deserves special attention because it occupies a unique position in the loop. Factual knowledge is additive: learning a new fact adds one fact to the store. Operational knowledge is *multiplicative*: learning a new learning strategy improves the efficiency of all subsequent learning.

A student who learns spaced repetition — distributing practice over time rather than cramming — doesn’t merely acquire one new fact. She acquires a tool that increases the retention rate of everything she learns from that point forward. A student who learns to identify his own knowledge gaps and address them systematically doesn’t just fix one gap; he acquires a skill that prevents hundreds of future gaps. This is knowledge that accelerates the loop itself.

If any single component deserves the label “what makes people smart,” it is operational knowledge. Not IQ. Not raw processing power. The meta-skill of knowing how to learn effectively.

Why IQ Tests Miss the Point

IQ tests measure *maximum performance* — what a person can do under standardized conditions, with maximum effort assumed. They capture a snapshot of one component (Performance on specific tasks) at one moment in time. They do not — they cannot —

capture the recursive, self-reinforcing, multi-component process that intelligence actually is.

This is why IQ scores tell you so little about long-term intellectual trajectory. Two children with identical IQ scores at age six can diverge dramatically by age thirty — one becoming a research scientist, the other having stopped reading after school. Standard psychometric models struggle with this divergence. The recursive model predicts it: the children differed not in Performance but in Motivation and operational Knowledge, and the recursive loop amplified those differences over twenty-four years of compounding iteration.

The IQ test is like measuring the horsepower of a car’s engine without checking whether the car has fuel or a driver. Horsepower matters — but it’s not the bottleneck for most journeys.

The AI Test Case

The recursive model makes a specific prediction about artificial intelligence: systems with high Knowledge and high Performance but no Motivation should fail to exhibit the self-directed development that characterizes human intelligence. And this is exactly what we observe.

Current large language models possess vast knowledge (trained on trillions of tokens), high performance (billions of parameters), and no motivation whatsoever. They process what they are given and produce what they are asked for. Between queries, they do nothing. They do not seek out areas of ignorance. They do not practice skills. They do not wonder about problems. Their “intelligence”

is entirely static — determined by training, with no endogenous drive to extend it.

Even the most advanced reasoning models — capable of solving competition-level mathematics — exhibit this precise failure mode. They solve extraordinary problems *when prompted* but do not independently seek out problems, do not self-direct their learning, and require external scaffolding that functions as a surrogate for the absent motivation component. Scale Performance and Knowledge as high as you like: without Motivation, the loop doesn’t self-sustain.

This is not merely because these systems weren’t designed to self-improve. That observation concedes the point: designing a system that self-improves requires engineering a functional analogue of motivation. Until AI systems have that, they will remain tools that are used rather than agents that develop.

The Connection to Consciousness

Here is where the intelligence model connects back to the Four-Model Theory at the heart of this book. The recursive intelligence loop doesn’t just *benefit from* consciousness — it *requires* it.

The loop depends on a specific cognitive capacity: *cognitive learning* — the ability to induce general theories from particular observations, as distinct from mere reinforcement learning (stimulus-response conditioning). Reinforcement learning can train you to avoid a hot stove through pain. Cognitive learning lets you watch someone else touch a hot stove and generalize: “Hot things burn. Don’t touch hot things.” The difference is the ability to simulate scenarios from a third-person perspective — to model yourself as

an object in the world and reason about what would happen if you did various things.

This is precisely what the Explicit World Model and Explicit Self Model provide. Consciousness — the ability to create and run a self-simulation — is the *substrate* on which the recursive intelligence loop operates. Without explicit models, you get reinforcement learning, which works but doesn't compound. With explicit models, you get cognitive learning, which feeds the recursive loop and compounds across a lifetime.

This is why the animal intelligence gradient from Chapter 10 maps onto the consciousness gradient. More sophisticated self-models enable more sophisticated recursive loops. A dog, with a relatively simple self-model, runs a limited version of the loop — it can learn from observation to some degree, but its cognitive learning is constrained by the richness of its explicit models. A chimpanzee, with a richer self-model, runs a more powerful loop. A human, with the full four-model architecture, runs the loop at maximum capacity — and the results are language, culture, science, and everything else that distinguishes human intelligence from animal cognition.

The Learnability Implication

The recursive model yields a consequence that I consider more important than all the theoretical arguments combined: it predicts that intelligence is, to a large extent, *learnable*.

Knowledge is entirely learnable — that's true by definition. Motivation is substantially learnable — decades of research in self-

determination theory show that intrinsic motivation is not a fixed trait but a response to environmental conditions, particularly autonomy, competence, and relatedness. Performance has a biological ceiling, but for the vast majority of people that ceiling is not the bottleneck. Average cognitive processing capacity is more than sufficient for what most people would recognize as highly intelligent behavior.

The binding constraints, for most people most of the time, are Motivation and operational Knowledge. And both are responsive to intervention.

This has a dark corollary. Any system that systematically destroys motivation in learners is not merely failing to develop intelligence — it is *actively suppressing* it. Conventional grading systems do exactly this. A poor grade doesn’t just report a result; it attacks the Motivation component. Reduced Motivation means fewer iterations of the loop. Fewer iterations mean slower growth in Knowledge. Slower growth means worse performance on the next assessment. Worse performance means more poor grades. The loop has reversed: instead of compound growth, the child is now trapped in compound stagnation. The grading system produces the very outcome it claims to merely measure.

The recursive model predicts that this damage compounds over time — not static harm but accelerating divergence. Early motivational damage should show up as a fanning out of trajectories that grows wider with each passing year. Conversely, motivation-enhancing interventions should show benefits that *compound* — larger effects at five-year follow-up than at one-year follow-up. And indeed, analyses of early childhood interventions like the

Perry Preschool Project show exactly this pattern: returns that grow over time, driven not by persistence of initial cognitive gains (which often fade) but by compounding motivational and self-regulatory gains.

If there is one practical takeaway from the intelligence model, it is this: the most valuable thing an educational system can transmit is not factual knowledge — in the age of AI, facts are free — but *operational knowledge* and the motivation to use it. Learning how to learn, and wanting to learn, are close to the only things still worth teaching.

The External Dependency

One last point, because it's easy to miss and it matters. The recursive loop is self-reinforcing, but it is not self-sufficient. It requires external fuel — information, challenges, feedback, access to the next level of knowledge. My own experience as a child, hitting a wall at age eleven not because of any internal limitation but because the supply of mathematics books ran out, illustrates this perfectly. All three components were healthy. The loop stalled anyway, because loops need input from outside to keep iterating.

This means that intelligence development depends not only on the person but on the environment. Access to knowledge, quality of instruction, availability of mentors, cultural attitudes toward learning — all of these feed or starve the loop. The recursive model explains why socioeconomic factors predict intellectual development so powerfully: they determine the supply of external fuel. A child in a book-rich home with engaged parents has the loop fed

continuously. A child in a resource-poor environment has the loop starved, regardless of the child's internal capacity.

Intelligence is not a trait you have. It is a process you run. And whether the process runs well depends on the machine (Performance), the software (Knowledge), the driver (Motivation), and the road (the external environment). All four matter. Any model that leaves one out will get the predictions wrong.

Chapter 17

Appendix C: Five Classes of Computation

This appendix expands on the computational framework briefly mentioned in Chapter 5 — the five classes of dynamical behavior that determine whether a physical system can support consciousness. Readers comfortable with the intuitive version in Chapter 5 can skip this appendix without missing anything needed for the main argument. For those who want the full picture: this is where the mathematics meets the physics.

Wolfram's Four Classes

In 2002, Stephen Wolfram published *A New Kind of Science*, the result of decades spent studying what happens when you let very simple rules run on very simple systems. His central tool was the cellular automaton — a row (or grid) of cells, each one either on or off, updated simultaneously according to a fixed rule that looks only at each cell's immediate neighbors.

The surprise was how much variety these trivially simple rules could produce. Wolfram classified the behavior into four types:

Wolfram Class	Behavior	Example	What you see
1	Uniform	Rule 0	Everything goes blank. Every cell dies.
2	Periodic	Rule 4	Stable, repeating patterns. Blinkers. Clocks.
3	Random/chaotic	Rule 30	Apparent randomness. No obvious repeating structure.
4	Complex	Rule 110	Localized structures that move, interact, and persist.

This classification was genuinely useful. It captured something real about how dynamical systems behave, and it applied far beyond cellular automata — to fluid dynamics, biological systems, economic models, and neural networks. The four classes weren’t just categories; they were attractors. Systems across wildly different domains kept falling into the same four behavioral regimes.

But there was a problem.

The Fractal Problem

Wolfram’s Class 3 was a grab-bag. It contained two fundamentally different kinds of system that happened to *look* similar at a glance:

Fractal systems like Rule 90, which generates a perfect Sierpinski triangle — an infinitely self-similar, recursively structured pattern. Beautiful, deterministic, and computationally boring: you can calculate any cell at any time step without running the whole simulation. Mathematicians call this *computationally reducible*.

Apparently chaotic systems like Rule 30, whose output column Wolfram himself used as a pseudorandom number generator in *Mathematica*. These produce output that *looks* random but is completely deterministic — same input, same output, every time. You can't shortcut the calculation; you have to run every step. Mathematicians call this *computationally irreducible*.

Wolfram put both in Class 3. His definition emphasized the *appearance* of randomness — “seems in many respects random” — while noting that “triangles and other small-scale structures are essentially always at some level seen.” He acknowledged the classification was imperfect: “with almost any general classification scheme there are inevitably cases which get assigned to one class by one definition and another class by another definition.”

Eric Rowland, at the 2006 NKS conference, independently argued that nested (fractal) patterns deserved their own classification framework.

I think the problem goes deeper than classification aesthetics. Fractal systems and chaotic systems are structurally different in a way that matters for the core argument of this book: which systems can support consciousness?

The Five-Class Scheme

Here is the five-class scheme, ordered as a clean monotonic gradient from most ordered to most disordered:

Class 1 — Static. Systems that converge to a fixed state and stop. A pendulum that swings once and stills. Dead. Nothing computes. Period: 1.

Class 2 — Periodic. Systems that settle into repeating loops. A clock. A heartbeat (approximately). Information is stored in the pattern but never transformed. Period: finite.

Class 3 — Fractal. Systems that produce self-similar structure at every scale. A Sierpinski triangle. A fern. The Mandelbrot set. Mathematically rich, aesthetically stunning, and *computationally reducible* — you can skip ahead without running every step. Structure without processing power. Period: quasi-infinite, with exact or statistical self-similarity at every scale.

Class 4 — Complex (edge of chaos). Systems that produce persistent localized structures that move, interact, and can encode arbitrary computation. Conway’s Game of Life. The cortical automaton. Computationally *irreducible* — no shortcuts. These systems are capable of universal computation: given the right initial conditions, they can simulate any algorithm, including simulations of themselves. Period: quasi-infinite, with self-similarity *plus* persistent interacting structures. This is where consciousness lives.

Class 5 — Random. Systems whose output is genuinely random — not pseudorandom, not deterministic, not compressible. No pattern, no self-similarity, no period that eventually repeats. Truly infinite information content. Structure: *unknown* (see below).

The mapping to Wolfram’s scheme:

Five-class	Wolfram	What changed
1 — Static	Class 1	Same
2 — Periodic	Class 2	Same
3 — Fractal	Class 3 (part)	Split out from Wolfram's Class 3
4 — Complex	Class 4	Same
5 — Random	Class 3 (part)	Split out from Wolfram's Class 3

Wolfram's ordering on the disorder spectrum was: 1 → 2 → 4 → 3. Awkward. The five-class scheme gives a clean monotonic gradient: 1 → 2 → 3 → 4 → 5, ordered by increasing disorder and increasing computational irreducibility.

Why Deterministic Automata Cannot Produce Randomness

Here is the argument that I believe is original and that strengthens the case for five classes rather than four.

Consider a cellular automaton — any cellular automaton. It has a finite rule table (expressible in a finite number of bits) and a finite initial condition (also expressible in finite bits). Together, the rule and the initial condition contain a fixed, finite amount of information.

Now: can a finite amount of information produce a truly random output?

No. Here is why:

1. A truly random infinite sequence has *maximal* Kolmogorov complexity — it cannot be compressed, it cannot be described by anything shorter than itself.
2. The output of a cellular automaton is entirely determined by its rule and initial condition, which together have *finite* Kolmogorov complexity.
3. You cannot get more information out of a process than you put in through its specification.
4. Therefore, the output of any cellular automaton has low Kolmogorov complexity relative to a truly random sequence of the same length.

This is a generalized pigeonhole argument: finite information must produce self-similar structure. The only way to generate infinite output from finite information is to *reuse* structure at different scales. Exact reuse is periodicity (Class 2). Non-exact but patterned reuse is fractal behavior (Class 3). Even the most complex-looking cellular automata — Rule 30, Rule 110, the Game of Life — are producing output whose complexity is bounded by their rule-set complexity.

What Wolfram called “random” cellular automata are better described as **high-complexity fractals** — systems whose self-similar structure is real but operates at scales and in dimensions that make it invisible to casual inspection. Rule 30’s left edge, in fact, shows Sierpinski-like substructures. Its center column passes many statistical tests for randomness — which is *exactly what you’d expect from a high-complexity fractal*: the local statistics mimic randomness, but the global structure is deterministic and compressible.

By this argument, Class 4 output is *also* fractal — the Game of Life exhibits statistical self-similarity in its population dynamics, its structural distributions, its spatial correlations. The difference between Class 3 and Class 4 is not “fractal vs. not-fractal.” It is:

- **Class 3:** Fractal. Reducible. Structure without processing.
- **Class 4:** Fractal. Irreducible. Structure *with* processing — persistent localized structures that interact and can encode universal computation.

Both are fractal. Only one computes.

Class	Rules	Period	Structure	Reducible?	Computes?
1 — Static	Finite	1	None	Trivially	No
2 — Periodic	Finite	Finite	Repeating	Yes	No
3 — Fractal	Finite	Quasi-infinite, self-similar	Self-similar	Yes	No
4 — Complex	Finite	Quasi-infinite, self-similar	Self-similar + persistent inter- acting structures	No	Yes
5 — Random	Inexpressible	Truly infinite	Unknown	N/A	N/A

Class 5 and the Boundary of Mathematical Expressibility

If deterministic automata cannot produce true randomness, then what *can*?

This question leads to what I believe is the deepest implication of the five-class scheme.

Classes 1 through 4 are what finite, expressible rules can produce. Their behavior ranges from trivial (Class 1) to extraordinary (Class 4 — universal computation, consciousness), but all of it is generated by rules that can be written down, communicated, verified, and analyzed. These rules live within mathematics — within the domain of formal symbolic systems.

Class 5, by contrast, requires rules that *cannot* be written down. A system that produces genuinely random output — output with maximal Kolmogorov complexity, incompressible, non-algorithmic — cannot be running any rule that a formal system can express. If the rule were expressible, the output would be compressible (to: “apply this rule”), and therefore not truly random.

This places Class 5 at the boundary of mathematical expressibility itself. It is not merely “very complex” or “very disordered.” It is the regime where the generating process exceeds what linear symbolic systems — mathematics, logic, computation — can capture.

Does anything in nature actually operate in Class 5?

Possibly. Quantum mechanics produces measurement outcomes that, by Bell’s theorem, cannot be explained by any local hidden-variable theory. If those outcomes are genuinely random —

not deterministic processes we haven't yet identified — then quantum measurement is a Class 5 process: a physical phenomenon whose rules cannot be written in any formal language we possess.

This is speculative, and I flag it as such. But the implication is striking: Class 4 — the regime of consciousness, of universal computation, of the cortical automaton — sits at the *maximum complexity achievable by expressible rules*. It is as complex as mathematics can get. Beyond it lies territory that mathematics, by its own nature, cannot map.

The Structure of Class 5: Unknown, Not Absent

A final subtlety. It is tempting to say that Class 5 has “no structure.” But this would be an error — the same error as saying that infinity has no structure.

Before Georg Cantor's work in the 1870s, infinity was treated as a single concept: things were either finite or infinite, end of story. Cantor showed that there are *hierarchies* of infinity — that the infinity of the real numbers is strictly larger than the infinity of the integers, and that this hierarchy extends without end. Infinity turned out to have a rich internal architecture that had been invisible because mathematicians lacked the tools to see it.

The same may be true of randomness. We currently treat true randomness as a single category — maximal disorder, the absence of pattern. But we are in the position of pre-Cantor mathematicians looking at infinity: we lack the conceptual tools to distinguish different kinds of randomness, if such distinctions exist.

The honest answer about Class 5 structure is therefore: **unknown**. Not “none.” Not “absent.” Unknown — awaiting conceptual tools that may not yet exist, that may require ways of thinking that linear symbolic systems cannot provide.

This is, I believe, one of the most important open questions at the intersection of mathematics, physics, and computation. And it is invisible without the five-class scheme, because Wolfram’s four-class framework never creates the space in which to ask it.

Implications for Consciousness

The five-class scheme clarifies why consciousness requires Class 4 dynamics — and only Class 4.

Classes 1 and 2 are too simple. They can store information (a fixed state, a repeating pattern) but cannot *process* it in any computationally interesting way. A brain in deep sleep, running slow delta waves, is operating in Class 2: periodic, repetitive, going nowhere. The four-model architecture is intact in the substrate, but the simulation is not running.

Class 3 is interesting but not computational. Fractal dynamics produce rich structure — and the brain uses them (see below) — but they cannot sustain the kind of dynamic, irreducible, globally integrated processing that a conscious self-simulation requires. A fractal pattern is beautiful, but it is computationally reducible. It cannot surprise itself.

Class 4 has exactly the two properties consciousness needs: **universal computation** (the system can, in principle, simulate anything, including itself) and **global integration** (distant parts of the system

influence each other, local changes propagate globally, information is bound into a unified whole). At the edge of chaos, the cortical automaton achieves both — and the result is consciousness.

Class 5 is different — not because computation is impossible there (an infinite random sequence contains *everything*, including every stable pattern and every computation ever conceived), but because we have no way to harness, predict, or demonstrate it. A brain in generalized seizure, with neurons firing in uncoordinated chaos, approaches Class 5 — not because consciousness is impossible in principle in that regime, but because no mechanism exists to sustain or access it. Our universe itself might be an excerpt of infinite randomness, or a Class 4 system at a scale we cannot perceive, or perhaps an excerpt of an infinite fractal. We cannot know. What we *can* say is that consciousness, as we experience it, requires the structured unpredictability of Class 4. The simulation collapses in Class 5 not because the underlying reality is insufficient, but because no stable interface exists between the substrate and the simulation.

The Brain Uses All Four Classes

The brain is a universal computer optimized by billions of years of evolution. It would be strange if evolution had missed any computational regime that offers an advantage. And indeed, the brain uses all four expressible classes as distinct tools:

- **Class 1** (stable attractors): Long-term memory storage. Synaptic weight configurations that persist for years. The fixed points of the neural network.

- **Class 2** (oscillations): Alpha, theta, gamma, and delta rhythms. Thalamic clocking. Sleep-wake cycles. The brain’s timekeeping and gating mechanisms.
- **Class 3** (fractal/scale-invariant processing): Texture analysis, scale-invariant object recognition, efficient neural encoding. Primarily V2-V4 visual processing, where multi-scale comparison is the core operation. Under psychedelics, when this machinery runs without external input, you *see* the fractal processing itself — which is why fractal patterns are among the most consistent features of psychedelic experience (see Chapter 6).
- **Class 4** (edge of chaos): The cortical automaton itself. The dynamical regime of conscious processing. Universal computation. The engine of the simulation.

Each class serves a different function. Only Class 4 generates consciousness. But consciousness depends on the others: stable memories (Class 1) to populate the models, rhythmic timing (Class 2) to coordinate the dynamics, and fractal processing (Class 3) to analyze the world at multiple scales simultaneously.

This is perhaps the deepest reason the brain must operate at the edge of chaos specifically: Class 4 is the only regime that can *recruit* the other three. A Class 4 automaton can generate stable states (Class 1 behavior), periodic oscillations (Class 2 behavior), and fractal structures (Class 3 behavior) as subprocesses within its own dynamics. None of the other classes can do this. Class 4 is not just the most complex class — it is the class that *contains* the others.

Chapter 18

Appendix D: How to Lucid Dream

In Chapter 6, I mentioned lucid dreaming as a safe, drug-free way to explore consciousness from the inside. In the Four-Model Theory, lucid dreaming is the Explicit Self Model “switching on” more fully during REM sleep — a criticality threshold crossing that turns a passive dream into a controlled experience. Here is the easiest method to get there.

The Reality Check Method

The principle is simple: if you habitually question whether you’re awake, that habit will eventually fire inside a dream — and the dream will fail the test.

Step 1: Pick a reality check. The most reliable one is to look at text, look away, and look back. In waking life, text stays the same. In dreams, it changes — often dramatically. Clocks work too: check the time, look away, check again. In a dream, the numbers will be different or nonsensical. Another reliable check: try to push a finger through your opposite palm. In a dream, it often goes through.

Step 2: Do it all day. Every time you walk through a doorway, check your phone, or notice something slightly odd, pause and perform your reality check. The key is not the check itself — it’s the *genuine question* behind it: “Am I dreaming right now?” Don’t just go through the motions. Actually consider the possibility.

Step 3: Keep a dream journal. Put a notebook by your bed. Every morning, before you move, write down whatever you remember — even if it’s just a feeling or a single image. This trains your brain to treat dream content as worth remembering, which strengthens the bridge between dreaming and waking awareness.

Step 4: Wait. For most people, the first lucid dream comes within two to six weeks. You’ll be in a dream, something will seem slightly off, the reality check habit will fire, the text will change — and you’ll *know*. That moment of knowing is the ESM activating. You’ll feel the transition: a sudden sharpening, a sense of presence, a quiet recognition that the world around you is a simulation you’re conscious inside of.

What to Expect

Your first lucid dream will probably be brief — seconds to a few minutes. Excitement tends to wake you up. With practice, you can extend them. Some people achieve lucid dreams several times a week. The experience is remarkable: you’re inside the full conscious simulation, with no external input, and you know it. The virtual world responds to your intentions. It is, quite literally, the Four-Model Theory made experiential.

Other Methods

For readers who want to go further, there are more involved techniques:

- **MILD (Mnemonic Induction of Lucid Dreams)** — developed by Stephen LaBerge at Stanford. You set an intention to recognize you're dreaming as you fall asleep. Best combined with waking up after five hours and returning to sleep.
- **WILD (Wake-Initiated Lucid Dream)** — you maintain awareness during the transition from waking to dreaming. Difficult, but produces the most vivid results.
- **WBTB (Wake Back to Bed)** — you wake after five to six hours of sleep, stay awake for twenty to sixty minutes, then return to sleep. This targets the REM-rich late sleep cycles.

Stephen LaBerge's *Exploring the World of Lucid Dreaming* (1990) remains the definitive practical guide. For the neuroscience, see Voss et al. (2009) on the EEG signatures of lucid dreaming, and Baird et al. (2019) for a comprehensive review of the cognitive neuroscience of lucid dreams.